


The long-run interval is cheaper at every d :

1.4--3.4 $\times$ faster than the i.i.d. plug-in (asymptotically $d^{1.01}$ vs $d^{1.86}$)



- BGSN: long-run interval
- the same, with the i.i.d. plug-in variance
- ◆ un-thinned curvature ($p=1$)
- ▲ ASGD (first order)
- ✕ offline HAC (not streaming)