

Inference, Not Just Optimisation: Streaming Inversion-Free Quasi-Newton Estimation on Dependent Data Streams

Ihor Kendiukhov^{a,*}

^a*Institute of Medical Genetics and Applied Genomics, University of
Tübingen, Tübingen, Germany*

Abstract

Streaming quasi-Newton methods fit a model in one pass, as cheaply as stochastic gradient descent but with the accuracy of Newton’s method. Their error bars, however, assume that observations are independent, and real streams – traffic counts, air quality, sensor readings – are not. The resulting confidence intervals do not merely lose a little coverage: they converge to a wrong coverage and stay there. On the hourly air-quality stream studied here, a nominal 95% interval covers about 47%.

One change repairs the estimate and the interval together, within the same cost budget: read the stream in contiguous blocks. The block-averaged gradients such a method already computes are exactly the ingredients of a consistent estimator of the variance the intervals need, so the correct statistical object and the computational shortcut coincide – and the valid interval turns out cheaper than the invalid one it replaces. We prove consistency, a rate and a central limit theorem with the correct limit under geometric mixing, derive the exact coverage that dependence-blind intervals converge to, and report the one cost of blocking: a step-size condition it makes non-trivial, with a remedy. Across simulated designs and two real hourly streams, coverage rises from 0.46–0.83 to 0.78–0.95.

Keywords: streaming estimation, stochastic Newton, dependent data, long-run variance, confidence intervals, mixing

1. Introduction

Many estimation problems now arrive as a stream: observations appear one at a time or in small batches, they are too numerous to store, and a decision or a report is required at any moment. Stochastic approximation ([Robbins and Monro, 1951](#)) is the natural tool, and averaged stochastic gradient descent

*Corresponding author.

(Ruppert, 1988; Polyak and Juditsky, 1992) is asymptotically efficient with $O(d)$ work per observation. Its weakness is conditioning: when the curvature of the objective has eigenvalues spread over orders of magnitude, first-order methods crawl (Bottou et al., 2018; Leluc and Portier, 2023). Second-order methods fix the conditioning but appear to cost $O(d^3)$ per step.

A line of work on *inversion-free* stochastic Newton methods removes that obstacle. When the Hessian has the form $\nabla^2 F(\theta) = \mathbb{E}[\alpha(\theta; \xi) \Psi(\theta; \xi) \Psi(\theta; \xi)^\top]$ —as it does for linear, logistic, softmax, Poisson and ridge regression, and for the geometric median—the running curvature estimate is a sum of rank-one terms, so its inverse can be maintained directly by the Sherman–Morrison rank-one update (Sherman and Morrison, 1950) — also called the Riccati form, the name we use in the proofs — without ever forming or inverting a matrix (Bercu et al., 2020; Cénac et al., 2025; Boyer and Godichon-Baggioni, 2023; Godichon-Baggioni and Lu, 2024). Godichon-Baggioni and Werge (2025) push this to its natural conclusion: thinning the rank-one updates and taking mini-batches of size d gives an estimator whose *total* cost is $O(dN)$ —first-order cost—while remaining asymptotically efficient. A concurrent construction achieves the same budget by masking Hessian columns (Godichon-Baggioni et al., 2026).

All of this theory assumes the stream is independent and identically distributed. The streams that motivate it are not. Hourly traffic volume, pollutant concentrations, queue lengths, and user activity are serially dependent, and after any regression adjustment the *scores* remain serially dependent: in the two real streams we study, the lag-one residual autocorrelation after a fixed weather-and-calendar adjustment is 0.76 (traffic) and 0.92 (air quality). The consequence is not a small correction. Writing $g_i = \nabla_\theta f(\theta^*; \xi_i)$ for the score at the truth, $\Gamma_k = \text{Cov}(g_0, g_k)$, $H = \nabla^2 F(\theta^*)$ and

$$S = \sum_{k \in \mathbb{Z}} \Gamma_k \quad (\text{the long-run score covariance}), \quad (1)$$

any efficient streaming M -estimator satisfies $\sqrt{N}(\hat{\theta}_N - \theta^*) \rightsquigarrow \mathcal{N}(0, H^{-1} S H^{-1})$, whereas the independent-data theory certifies $\mathcal{N}(0, H^{-1} \Gamma_0 H^{-1})$. The ratio of the two variances, coordinatewise, is a factor κ_j that does not shrink with N . It is not a rounding error: we estimate it at about 2 for the traffic stream and about 20 for the air-quality stream, and at $\kappa_j = 3$ a nominal 95% interval covers 75% of the time *forever*, however long the stream runs.

This paper. We ask what has to change in a streaming inversion-free quasi-Newton method so that both its point estimate and its *uncertainty* are correct on a dependent stream, without leaving the $O(dN)$ budget. Our answer is

55 that one structural choice does both jobs. Consume the stream in *contiguous*
 56 *blocks* of length b and take one preconditioned step per block; place curvature
 57 updates on a *gap* of m observations. Then

- 58 • the b -averaged gradients the algorithm forms anyway are, up to a factor
 59 b , the non-overlapping batch means of the score process, so they *are* an
 60 estimator of S —the statistically correct object and the computational
 61 shortcut coincide; and
- 62 • gapping decorrelates the rank-one curvature terms, which is what makes
 63 the inversion-free recursion behave, at matched cost, as it would on
 64 independent data.

65 Neither device substitutes for the other: gapping alone leaves the variance
 66 wrong, and blocking alone leaves the curvature recursion aggregating depen-
 67 dent terms. We call the resulting estimator BGSN (blocked–gapped streaming
 68 Newton).

69 *Contributions..*

70 1. **Theory for the inversion-free curvature recursion under de-**
 71 **pendence** (Section 3). We prove almost sure convergence, the rate
 72 $\|\theta_T - \theta^*\|^2 = O(\log N/N)$, and a central limit theorem with the long-run
 73 sandwich $H^{-1}SH^{-1}$, for a rank-one Sherman–Morrison curvature recur-
 74 sion with gapped updates on a geometrically mixing stream (Theorems 3
 75 to 6). The independent-data analysis rests on the score being a martingale
 76 difference sequence with respect to the algorithm’s filtration; dependence
 77 destroys that, and existing dependent-data results are first-order, while
 78 existing streaming second-order inference assumes conditionally unbiased
 79 gradients. A by-product (Proposition 7) is that for this algorithm blocking
 80 is *free*: the leading term of the error is $-N^{-1}H^{-1}\sum_{i\leq N}g_i$ for *every* b ,
 81 so the asymptotic variance does not depend on the block length.

82 2. **A long-run covariance estimator from block gradients, and why**
 83 **it is fast** (Section 4). S is estimated at $O(d^2)$ per block and $O(d^2)$
 84 memory from the same block gradients. The key point is *what* is being
 85 blocked: batch means taken over the *iterate* path need blocks long enough
 86 for the iterates to mix, which forces $b \asymp N^{3/4}$ and caps the rate at $N^{-1/8}$
 87 (Zhu et al., 2023; Roy and Balasubramanian, 2023; Samsonov et al., 2025);
 88 block *gradients* inherit their dependence from the data, so $b \asymp \log N$
 89 suffices and Theorem 10 gives $\tilde{O}(N^{-1/2})$. This does not contradict the
 90 minimax rate $\Theta(N^{-(1-a)/2})$ for Hessian-free covariance estimation from
 91 an SGD trajectory with step exponent a (Ni and Huo, 2026), which is at

best $N^{-1/4}$ over the admissible range $a > 1/2$: a quasi-Newton method is not in that class, because it already maintains curvature and forms block gradients.

3. **An exponentially small block-length bias** (Theorem 10). What makes $b \asymp \log N$ admissible is that the two-scale (flat-top / lugsail) correction $\hat{S}^{\text{FT}} = \hat{S}^{\text{BM}} + b(\hat{\Lambda}_1 + \hat{\Lambda}_1^\top)$ has bias $O(\rho^b)$ rather than the $O(1/b)$ of plain batch means. (A flat-top lag window is one whose weights are flat near lag zero, so its bias picks up only the far tail of the autocovariances rather than shaving every lag.) The correction and this tail-sum property are due to the flat-top literature (Politis and Romano, 1995; Politis, 2011; Vats and Flegal, 2022; Singh et al., 2025), not to us (Section 4 is explicit about the boundary); what we add is that it is available at *block-gradient* resolution on a dependent stream, and a streaming (dyadic) implementation that needs no advance knowledge of N .

4. **Exactly how badly dependence-blind intervals fail** (Proposition 8). Their asymptotic coverage is $2\Phi(z_{\alpha/2}/\sqrt{\kappa_j}) - 1$; over a sweep in dependence strength the largest gap between prediction and measurement is 0.0085, and 0.0063 for κ up to 2.9, with no fitted constants (Figure 1). This turns a previously reported qualitative failure into a quantitative one.

5. **A stability condition that blocking creates, and its fix** (Proposition 1). Blocking the gradient makes the step’s contraction condition $\gamma_t \|\bar{H}^{-1} \hat{H}_t\|_{\text{op}} < 2$ non-trivial, where γ_t is the step size, $\bar{H}$ the running curvature estimate used as preconditioner, and $\hat{H}_t = b^{-1} \sum_{i \in B_t} \alpha_i \Psi_i \Psi_i^\top$ the Hessian of block t alone. On a dependent stream a block of b observations can span far fewer than b distinct covariate directions, so $\hat{H}_t$ is near-singular and the norm is large. We give the condition, a step-size shift t_0 read off the warm-up window that restores it at $O(c_w d^2)$ one-off cost, and a proof that the shift changes no asymptotic statement. Without it our own method diverges on a regime-switching design; with it that design is stable and its coverage tracks the infeasible oracle interval at every sample size (Tables 2 and 5). We report this because it is a cost of blocking, and it is ours.

6. **Curvature schedules** (Proposition 11). At a matched number of rank-one updates, deterministic gapping incurs excess curvature variance $O(\rho^m)$ whereas the Bernoulli thinning used as the cost knob in the independent-data method incurs $(\kappa_H - 1)/m$. Gapping is thus a strict, if modest, improvement on the existing cost knob, and it makes the per-block cost

130 deterministic.

131 **7. Evidence** (Section 5). Four dependent designs with exact population
 132 targets (three with closed-form S), plus hourly traffic and air-quality
 133 streams under two protocols, one of which has an exact ground truth on
 134 real covariates.

135 *What we do not claim..* The $O(dN)$ budget and the idea of thinning curvature
 136 updates are due to Godichon-Baggioni and Werge (2025); blocked or growing
 137 mini-batches as a dependence-breaking device are due to Godichon-Baggioni
 138 et al. (2023); the long-run sandwich limit under mixing or Markovian sampling
 139 is known for first-order stochastic approximation (Liu et al., 2023; Li et al.,
 140 2023; Samsonov et al., 2025; Li et al., 2026); recursive online estimation of
 141 the *instantaneous* asymptotic variance under independent data is due to
 142 Godichon-Baggioni (2019) and Chen et al. (2020); Zhu et al. (2023); online
 143 estimation of a *long-run* covariance under dependent sampling is known for
 144 first-order methods (Roy and Balasubramanian, 2023; Samsonov et al., 2025);
 145 online covariance estimation for streaming second-order methods is known
 146 under independent data (Kuang et al., 2026; Wang et al., 2026); the two-scale
 147 correction is known (Politis and Romano, 1995; Vats and Flegal, 2022; Singh
 148 et al., 2025); the semiparametric efficiency bound under Markovian sampling
 149 is due to Li et al. (2023), and we attain rather than establish it. Section 6
 150 states each boundary precisely.

151 2. Setting and notation

152 We observe a strictly stationary stream $\xi_1, \xi_2, \dots$ on $\mathbb{R}^p$ and wish to
 153 estimate

$$\theta^* = \arg \min_{\theta \in \mathbb{R}^d} F(\theta), \quad F(\theta) = \mathbb{E}[f(\theta; \xi)], \quad (2)$$

154 where $f(\cdot; \xi)$ is a loss. Write $g_i(\theta) = \nabla_\theta f(\theta; \xi_i)$ and $g_i = g_i(\theta^*)$, so $\mathbb{E}[g_i] = 0$.
 155 Let $H = \nabla^2 F(\theta^*)$, $\Gamma_k = \mathbb{E}[g_0 g_k^\top]$ and let S be as in (1). Throughout, $\Phi(x)$ is
 156 the standard normal distribution function, $z_{\alpha/2} = \Phi^{-1}(1 - \alpha/2)$, $\|\cdot\|_{\text{op}}$ is the
 157 operator norm, and N counts observations while T counts blocks.

158 *Terminology..* Five terms recur and are used in their standard senses. The
 159 *score* is the gradient of the loss at one observation, $g_i(\theta)$; an *M-estimator* is
 160 a minimiser of an empirical average of losses. A *sandwich* covariance is one
 161 of the form $H^{-1}\Sigma H^{-1}$; it is *instantaneous* when $\Sigma = \Gamma_0$ and *long-run* when
 162 $\Sigma = S$. *Batch means* estimates a variance by averaging within consecutive

163 blocks and taking the sample covariance of the block averages. The strong
 164 (or α -) *mixing coefficient* of the stream at lag k is

$$\alpha_{\text{mix}}(k) = \sup \{ |\mathbb{P}(A \cap B) - \mathbb{P}(A)\mathbb{P}(B)| : A \in \sigma(\xi_i, i \leq 0), B \in \sigma(\xi_i, i \geq k) \}, \quad (3)$$

165 a measure of how nearly independent the distant past and the distant future
 166 are; it is zero for all $k \geq 1$ if and only if the stream is independent.

167 Nothing below requires the model to be correctly specified. All that is used
 168 is that θ^* solves $\nabla F(\theta^*) = 0$ for the *working* loss, so under misspecification
 169 θ^* is the pseudo-true parameter and $H^{-1}SH^{-1}$ is the corresponding long-run
 170 sandwich, in the usual sense of [White \(1982\)](#) and [Domowitz and White \(1982\)](#).
 171 This matters in practice, because omitted dynamics are one of the commonest
 172 ways for a score to become serially correlated (Remark 1).

173 **Remark 1** (Dependence in the covariates is not enough). If the conditional
 174 law of y given x is correctly specified and its innovations are independent
 175 over time, then (g_i) is a martingale difference sequence and $S = \Gamma_0$ *exactly*,
 176 however persistent x is. A genuine long-run-variance problem needs the *score*
 177 to be serially correlated: serially correlated errors, a latent dynamic factor,
 178 or omitted dynamics. All our designs therefore couple dependent covariates
 179 with a dependent error or latent process, and we report the resulting κ_j
 180 explicitly. Covariate dependence still matters for the effective sample size of
 181 the curvature estimate, which is what Proposition 11 is about.

182 2.1. The algorithm

183 *Blocks..* Partition the stream into contiguous blocks $B_t = \{(t-1)b+1, \dots, tb\}$,
 184 $t = 1, 2, \dots$, and write $\bar{g}_t(\theta) = b^{-1} \sum_{i \in B_t} g_i(\theta)$ for the block gradient and
 185 $\varepsilon_t = \bar{g}_t(\theta^*)$ for the block score.

186 *Curvature..* Let $\mathcal{C} = \{i \geq 1 : i \equiv 0 \bmod m\}$ be the *gapped* curvature slots
 187 and let $n_t = |\mathcal{C} \cap [1, tb]|$ count the rank-one updates made by the end of block
 188 t . Maintain an unnormalised accumulator A and a weight W :

$$A_n = A_{n-1} + \iota_n e_{k_n} e_{k_n}^\top + \alpha_{i_n} \Psi_{i_n} \Psi_{i_n}^\top, \quad W_n = W_{n-1} + 1, \quad (4)$$

189 where $\alpha_i = \alpha(\theta_{t-1}; \xi_i)$ and $\Psi_i = \Psi(\theta_{t-1}; \xi_i)$ for $i \in B_t$, e_k cycles through the
 190 canonical basis, and $\iota_n = c_\iota \varsigma_\star n^{-q_r}$ with $q_r \in (0, 1/4)$ is a vanishing ridge
 191 that keeps $\lambda_{\min}$ of the curvature estimate bounded below without any matrix
 192 operation. Its scale $\varsigma_\star$ is the mean curvature magnitude of the *first* block
 193 carrying curvature slots, $\varsigma_\star = (d|\mathcal{C} \cap B_1|)^{-1} \sum_{i \in \mathcal{C} \cap B_1} \alpha_i \|\Psi_i\|^2$, computed once
 194 and then frozen. Freezing matters and is not a detail: a scale that tracked

the *running* curvature would shrink exactly when the curvature estimate degenerated, and Lemma 2 would then presuppose its own conclusion. With $\varsigma_\star$ fixed, positive and finite almost surely under Assumption 3, the lemma holds unconditionally. The scale exists only to make c_ℓ dimensionless; setting $\varsigma_\star = 1$ recovers an unscaled ridge and Table 5(d) shows the results are the same. The curvature estimate and preconditioner are

$$\bar{H}_t = A_{n_t}/W_{n_t}, \quad \bar{H}_t^{-1} = W_{n_t}A_{n_t}^{-1}, \quad (5)$$

and A^{-1} is updated by two rank-one Sherman–Morrison steps per curvature slot, each $O(d^2)$; no matrix is ever inverted. Maintaining A^{-1} rather than $\bar{H}^{-1}$ is what keeps every update exactly rank one despite the changing normalisation W .

Parameter step. With $A_0 = \lambda_0 I$, $W_0 = 1$ (λ_0 a small positive constant in the units of the curvature; Table 5(e) varies it over eight orders of magnitude), and $\bar{H}_{t-1}$ built only from blocks $1, \dots, t-1$, the raw step is

$$\theta_t = \theta_{t-1} - \Pi_t \left[\frac{1}{t} \bar{H}_{t-1}^{-1} \bar{g}_t(\theta_{t-1}) \right], \quad \Pi_t[v] = v \cdot \min \left\{ 1, \frac{c_D (1 + \|\theta_{t-1}\|)}{\|v\|} \right\}, \quad (6)$$

where Π_t is a step-length safeguard, explained next, that is inactive whenever the step is of a reasonable size. The step $1/t$ is what makes the last iterate efficient without averaging; because it is preconditioned by $\bar{H}^{-1}$ there is no learning-rate constant to tune. A weighted averaged variant (step t^{-a} , $a \in (1/2, 1)$, with logarithmically weighted Polyak–Ruppert averaging) is available and reported alongside.

Why a blocked step needs a safeguard. Linearising (6) at $\theta^\star$, ignoring Π_t , gives the error recursion

$$\theta_t - \theta^\star = (I - \gamma_t \bar{H}_{t-1}^{-1} \hat{H}_t)(\theta_{t-1} - \theta^\star) - \gamma_t \bar{H}_{t-1}^{-1} \varepsilon_t + o(\|\theta_{t-1} - \theta^\star\|), \quad (7)$$

where $\hat{H}_t = b^{-1} \sum_{i \in B_t} \alpha_i \Psi_i \Psi_i^\top$ is the Hessian of block t alone. Since $\bar{H}^{-1} \hat{H}_t$ is a product of two positive semidefinite matrices, its eigenvalues are real and non-negative, so the spectral radius of the linear map in (7) is below one exactly when

$$\gamma_t \lambda_{\max}(\bar{H}_{t-1}^{-1} \hat{H}_t) < 2, \quad \text{for which} \quad \gamma_t \|\bar{H}_{t-1}^{-1} \hat{H}_t\|_{\text{op}} < 2 \quad (8)$$

is sufficient, since $\lambda_{\max} \leq \|\cdot\|_{\text{op}}$. We work with the operator norm because it is what we can compute and because it is the conservative choice; and we say

222 “spectral radius” rather than “contraction” deliberately, because for a non-
 223 symmetric map a spectral radius below one bounds the *asymptotic* growth of
 224 powers, not the norm of a single step. What (8) rules out is the regime we
 225 actually observe, in which $\gamma_t \lambda_{\max} \gg 2$ and a single step multiplies the error
 226 by a large factor. On an *independent* stream with $b \geq d$ the block Hessian
 227 has full rank and is close to H , so $\|\bar{H}^{-1} \hat{H}_t\|_{\text{op}} \approx 1$ and (8) holds already at
 228 $\gamma_1 = 1$. On a *dependent* stream it need not: a block of b observations from a
 229 persistent stream spans far fewer than b distinct covariate directions, $\hat{H}_t$ is
 230 close to singular, and the norm is inflated by roughly the ratio of the largest
 231 to the smallest curvature. We measured $\|H^{-1} \hat{H}_t\|_{\text{op}}$ at $b = 20$: median 20
 232 on our autoregressive design and 33 on our regime-switching design, with
 233 per-block maxima of 55 and 89. So (8) is violated at $\gamma_1 = 1$ on *both*, and
 234 the autoregressive designs are lucky rather than safe. Taking a full-length
 235 step under those numbers makes the first iterates overshoot by orders of
 236 magnitude, and with $\gamma_t = 1/t$ the excursion is undone only at rate $1/t$, so at
 237 realistic N it can survive to the end of the stream: without a safeguard our
 238 own method’s relative variance on the regime-switching design is 1.1×10^{14}
 239 times the efficient value, against 1.41 with the safeguard (Table 5(i)), and
 240 individual replicates diverge outright. This is not a defect of the base method.
 241 It is created by *blocking* the gradient, which is the one thing we add to it, so
 242 it is ours to fix and ours to report.

243 The fix in (6) bounds each step relative to the current iterate, $\|\Pi_t[v]\| \leq$
 244 $c_D(1 + \|\theta_{t-1}\|)$ with $c_D = 1$. It is scale-free in θ , costs $O(d)$, and — the point
 245 — it is *self-limiting*: it acts only on steps that are actually too long, and leaves
 246 every other step exactly as (6) prescribes.

247 **Proposition 1** (The safeguard is eventually inactive). *Let $\mathcal{A} = \{\Pi_t \text{ is active for only finitely many } t\}$.
 248 On $\mathcal{A}$ there is a finite $T_0(\omega)$ with $\theta_t = \theta_{t-1} - \gamma_t \bar{H}_{t-1}^{-1} \bar{g}_t(\theta_{t-1})$ for all $t > T_0$,
 249 so the safeguarded and unsafeguarded iterates coincide from T_0 on and every
 250 asymptotic statement about one holds for the other; in particular Theorems 4
 251 to 6 and Proposition 7 are unaffected. Moreover $\mathbb{P}(\mathcal{A}) = 1$ whenever $\theta_t \rightarrow \theta^*$
 252 and $\|\bar{H}_t^{-1}\|_{\text{op}} \bar{g}_{t+1} \rightarrow 0$ almost surely, because then $\|\gamma_t \bar{H}_{t-1}^{-1} \bar{g}_t\| \rightarrow 0$ while
 253 $c_D(1 + \|\theta_{t-1}\|) \rightarrow c_D(1 + \|\theta^*\|) > 0$.*

254 We state Proposition 1 conditionally on $\mathcal{A}$ because that is what we can prove,
 255 and we say plainly what is therefore *not* proved: we do not establish almost
 256 sure convergence for the safeguarded recursion from scratch. The obstruction
 257 is real and worth naming — the factor $\min\{1, c_D(1 + \|\theta_{t-1}\|)/\|v\|\}$ depends
 258 on the current block’s gradient, so it is $\mathcal{F}_t$ - rather than $\mathcal{F}_{t-1}$ -measurable and
 259 correlates with the noise, which breaks the martingale structure the proofs use.
 260 A complete treatment is available in the stochastic-approximation literature

on expanding truncations (Andrieu et al., 2005; Kushner and Yin, 2003), at the price of machinery we do not need here. What we do instead is measure $\mathcal{A}$: across every design and every sample size we report, the safeguard binds between 2 and 7 times, and the count does not grow with N (Table 5(i), Table 2), which is the empirical content of $\mathbb{P}(\mathcal{A}) = 1$. A unit test asserts it.

The alternative we rejected, and why.. (8) suggests a different remedy: shift the schedule to $\gamma_t = (t + t_0)^{-1}$ with $t_0 = \frac{1}{2} \max_{t \leq T_w} \|\bar{H}_{T_w}^{-1} \hat{H}_t\|_{\text{op}}$ measured on the warm-up blocks, which makes (8) hold at $t = 1$ by construction and has a cleaner proof (it is a finite deterministic reindexing; see Appendix C). We implemented it and it is worse, for a reason worth stating: it slows *every* step, whereas the overshoot needs correcting on only a handful. On the regime-switching design at $N = 200,000$ the shift leaves the relative variance at 3.08 against 1.41 for the safeguard, and on a stream short enough that t_0 is comparable to the number of steps — which is the case for our real-data segments — it freezes the estimator altogether. Both variants are in the code and both are reported (Table 5(i)); the safeguard is the default.

Curvature warm-up.. The first $\min\{\lceil c_w d/b \rceil, \lfloor \varpi_w N/b \rfloor\}$ blocks are spent on curvature updates only, with no parameter step and with every observation contributing a rank-one update. The first term is the rule; the second caps the warm-up at a fraction $\varpi_w = 0.1$ of the stream and binds only when $c_w d > \varpi_w N$, which is exactly the regime in which calling an $O(d)$ -observation warm-up negligible would be false. It matters: on a 4000-observation stream with $d = 11$, spending $c_w d = 2200$ observations on curvature leaves 45 blocked steps and a relative variance of 2.51, whereas capping at 10% leaves 90 steps and a relative variance of 1.28 (Table 5(j)). Our real-data segments are of exactly that length. This consumes $O(d)$ observations and $O(c_w d^3)$ flops— $O(1)$ in N —and changes no asymptotic statement. It is not cosmetic: a $1/t$ step taken while $\bar{H}$ is still uninformative leaves a component that decays like t^{-c} with $c < 1$, and at realistic N that component dominates the sampling error and destroys coverage (Table 5). We regard “do not start the efficient schedule before the curvature estimate is informative” as a practical finding of independent interest.

Cost.. Per block after the warm-up: $O(bd)$ for the gradient, $O(d^2)$ for the $\lceil b/m \rceil$ curvature updates and the preconditioned step, and $O(d^2)$ for the covariance accumulators. With $b = m = d$ this is $O(d^2)$ per block of d observations, i.e. $O(d)$ amortised per observation, matching Godichon-Baggioni and Werge (2025). The warm-up adds a one-off $O(c_w d^3)$: it performs $c_w d$ rank-one updates at $O(d^2)$ each, which is the unavoidable cost of producing a

299 $d \times d$ curvature estimate at all and is of the same order as a single matrix fac-
 300 torisation. The stability shift (C.1) adds a further one-off $O(c_w(d^2 + db)d/b)$,
 301 which is 4–5% of the total operation count at the settings of Table 6 and
 302 vanishes relative to the streaming pass as N grows. Total time is therefore

$$O(dN + c_w d^3), \quad \text{i.e. } O(dN) \text{ once } N \gtrsim c_w d^2, \quad (9)$$

303 and Table 6 reports the crossover explicitly rather than hiding it in the
 304 constant. Memory is $O(d^2)$: the running inverse and two $d \times d$ covariance
 305 accumulators. We state both plainly because it is easy to misread an $O(dN)$
 306 *time* claim as an $O(d)$ *memory* claim; it is not, and the same is true of the
 307 independent-data method we build on.

308 3. Theory

309 3.1. Assumptions

310 **Assumption 1** (Regularity). F is twice continuously differentiable with
 311 $\|\nabla^2 F(\theta)\|_{\text{op}} \leq L$ for all θ ; $H = \nabla^2 F(\theta^*) \succ 0$; and $\nabla^2 F$ is Lipschitz on a
 312 neighbourhood of θ^* , so that $\|\nabla F(\theta) - H(\theta - \theta^*)\| \leq L_\delta \|\theta - \theta^*\|^2$ for all θ .

313 **Assumption 2** (Expected smoothness and moments). There are $C, C' \geq$
 314 0 with $\mathbb{E}\|g(\theta)\|^2 \leq C + C'(F(\theta) - F(\theta^*))$ for all θ ; there is $\eta > 0$ with
 315 $\mathbb{E}\|g(\theta)\|^{2+2\eta} \leq C_\eta(1 + F(\theta) - F(\theta^*))^{1+\eta}$; and $\theta \mapsto \mathbb{E}[g(\theta)g(\theta)^\top]$ is continuous
 316 at θ^* .

317 **Assumption 3** (Curvature form). $\nabla^2 F(\theta) = \mathbb{E}[\alpha(\theta; \xi)\Psi(\theta; \xi)\Psi(\theta; \xi)^\top]$ with
 318 $\alpha \geq 0$; $\mathbb{E}\|\alpha(\theta; \xi)\Psi(\theta; \xi)\|_{\text{op}}^{\eta'} \leq C_{\eta'}$ for some $\eta' > 2$ uniformly on a neighbourhood
 319 of θ^* ; and $\theta \mapsto \alpha(\theta; \xi)\Psi(\theta; \xi)\Psi(\theta; \xi)^\top$ is Lipschitz in θ in L^2 near θ^* .

320 **Assumption 4** (Identifiability). θ^* is the only stationary point of F outside
 321 which the gradient points away from it: for every $\eta > 0$, $\inf_{\|\theta - \theta^*\| \geq \eta} \langle \nabla F(\theta), \theta - \theta^* \rangle > 0$.
 322

323 Assumption 4 holds whenever F is convex with a unique minimiser, which
 324 covers every objective we consider (least squares, logistic, ridge, softmax, the
 325 geometric median). It is what turns “the gradient step cannot keep moving”
 326 into “ $\theta_t \rightarrow \theta^*$ ”, and the independent-data analysis needs it too.

327 **Assumption 5** (Geometric mixing). $(\xi_i)_{i \geq 1}$ is strictly stationary and its
 328 mixing coefficient (3) decays geometrically: $\alpha_{\text{mix}}(k) \leq K c_{\text{mix}}^k$ for some $c_{\text{mix}} \in$
 329 $(0, 1)$, $K < \infty$. Moreover $\mathbb{E}\|g_1\|^{4+\nu'} < \infty$ for some $\nu' > 0$, $\mathbb{E}\|g_1\|^{2+\nu} < \infty$
 330 for some $\nu > 0$ (without loss of generality $\nu \leq 2\eta$), and $S \succ 0$. The fourth

moment with a margin is used only for the variance of the covariance estimator in Theorem 10(iii): the products $b\bar{g}_{tj}\bar{g}_{tk}$ whose autocovariances that bound needs are quadratic in the scores, and Davydov's inequality applied to them requires $2 + \delta$ moments of the products, i.e. $4 + 2\delta$ of the scores. Plain fourth moments would leave the argument one ϵ of integrability short, which is why we ask for the margin rather than quietly using it.

Assumptions 1 to 3 are those of the independent-data analysis (Godichon-Baggioni and Werge, 2025), with Assumption 3 strengthened from a moment bound to a local Lipschitz condition that we need because the curvature terms are now dependent. Assumption 5 is the only genuinely new assumption; a geometrically ergodic Markov chain (Meyn and Tweedie, 2009) satisfies it, as do stationary causal ARMA/GARCH-type recursions with geometrically decaying dependence in Wu's sense (Wu, 2005, 2011). Davydov's covariance inequality (Davydov, 1968; Rio, 1993) converts Assumption 5 into a geometric bound on the score autocovariances that we use repeatedly:

$$\|\Gamma_k\|_{\text{op}} \leq C_\Gamma \rho^{|k|}, \quad \rho = c_{\text{mix}}^{\nu/(2+\nu)} \in (0, 1). \quad (10)$$

In particular S in (1) converges absolutely.

3.2. The curvature recursion

The first result is where dependence bites. Under independence the terms accumulated in (4) form (conditionally on the iterates) an independent array and standard strong laws apply. Under Assumption 5 they are dependent, and the gapped design is what restores control: the subsequence indexed by $\mathcal{C}$ is itself stationary with mixing coefficients $\alpha_{\text{mix}}(jm) \leq Kc_{\text{mix}}^{jm}$.

Lemma 2 (The curvature estimate is never ill-conditioned). *Under Assumption 3, with $\iota_n = c_t \varsigma_\star n^{-q_r}$, $\varsigma_\star > 0$ fixed and $q_r \in (0, 1)$, there is a finite random $c_1 > 0$ with*

$$\lambda_{\min}(\bar{H}_t) \geq c_1 n_t^{-q_r} \quad \text{and hence} \quad \|\bar{H}_t^{-1}\|_{\text{op}} \leq c_1^{-1} n_t^{q_r} = O(t^{q_r}) \quad \text{for all } t, \text{ a.s.}$$

This is the whole purpose of the vanishing ridge, and it needs no assumption on the iterates: it makes $\sum_t \gamma_t^2 \|\bar{H}_{t-1}^{-1}\|_{\text{op}}^2 = \sum_t O(t^{2q_r-2}) < \infty$, the square-summability half of the Robbins–Monro condition, for every t rather than eventually, and without any eigenvalue computation.

Theorem 3 (Curvature rate). *Under Assumptions 1, 3 and 5, suppose $\theta_t \rightarrow \theta^\star$ almost surely with $\|\theta_t - \theta^\star\| = O(t^{-1/2} \log^c t)$ a.s. for some $c < \infty$. Then for every $\nu_H < q_r$,*

$$\|\bar{H}_t - H\|_{\text{op}} = O(t^{-\nu_H}) \quad \text{and} \quad \|\bar{H}_t^{-1} - H^{-1}\|_{\text{op}} = O(t^{-\nu_H}) \quad \text{a.s.}$$

Remark 2 (Order of the proofs). Theorem 3 presupposes a rate for the iterates while the results below use Theorem 3, so the logical order must be stated. It is a four-rung ladder, each rung using only the ones below it: (i) Lemma 2, from the ridge alone, with no assumption on the iterates; (ii) Theorem 4, $\theta_t \rightarrow \theta^*$ a.s., using only Lemma 2; (iii) a preliminary rate $\|\theta_t - \theta^*\| = O(t^{-1/2} \log^{1+\varpi} t)$ a.s. (Lemma 16, proved in Appendix A.4), from (A.11) with $\bar{H}_{t-1}^{-1}$ bounded by Lemma 2 and $\bar{H}_t \rightarrow H$ — consistency only, which follows from (ii); (iv) Theorem 3, and then Theorems 5, 6, 9 and 10. Appendix A is organised in this order, and rung (iii) is where the restriction $q_r < 1/4$ is used — everything else needs only $q_r < 1/2$.

3.3. Consistency, rate, and the central limit theorem

The independent-data proofs use that ε_t is a martingale difference with respect to $\mathcal{F}_{t-1} = \sigma(\xi_1, \dots, \xi_{(t-1)b})$. Under Assumption 5 it is not: $\mathbb{E}[\varepsilon_t | \mathcal{F}_{t-1}] \neq 0$, and no choice of block length repairs this, because the conditional bias of a block average is $O(1/b)$ and does not vanish at fixed b .

Our substitute is Gordin’s martingale–coboundary decomposition. Under Assumption 5 the conditional expectations $\mathbb{E}[\varepsilon_n | \mathcal{F}_0]$ decay geometrically in L^2 , so

$$\varepsilon_t = m_{t+1} + z_t - z_{t+1}, \quad \mathbb{E}[m_{t+1} | \mathcal{F}_t] = 0, \quad \mathbb{E}\|z_1\|^2 < \infty, \quad (11)$$

with (m_t) stationary and square integrable and $\mathbb{E}[m_1 m_1^\top] = S/b$ (Lemma 13). Every weighted sum $\sum_t \gamma_t \bar{H}_{t-1}^{-1} \varepsilon_t$ then splits into a square-integrable martingale, which converges almost surely because $\sum_t \gamma_t^2 \|\bar{H}_{t-1}^{-1}\|_{\text{op}}^2 = \sum_t O(t^{2q_r-2}) < \infty$ by Lemma 2, and a telescoping coboundary, which converges because one rank-one Riccati update moves the preconditioner by only $O(1/t)$ (Lemma 15). This is what makes the argument work with weights that are *random and adapted* — unavoidable here, since the weights contain $\bar{H}_{t-1}^{-1}$ — and it is why the long-run covariance S , rather than the instantaneous Γ_0 , is the object that appears: it is the second moment of the martingale part in (11). Appendix A gives the details.

Theorem 4 (Almost sure convergence). *Under Assumptions 1 to 5, the iterates (6) satisfy $\theta_t \rightarrow \theta^*$ almost surely, for every fixed block length $b \geq 1$ and gap $m \geq 1$.*

Theorem 5 (Rate). *Under the assumptions of Theorem 4, $\|\theta_t - \theta^*\|^2 = O(\log(N_t)/N_t)$ almost surely, where $N_t = bt$.*

396 **Theorem 6** (Central limit theorem with the long-run sandwich). *Under*
 397 *Assumptions 1 to 5, for each fixed block length $b \geq 1$,*

$$\sqrt{N_T}(\theta_T - \theta^*) \rightsquigarrow \mathcal{N}(0, H^{-1}SH^{-1}), \quad N_T = bT,$$

398 *with S as in (1). The same limit holds for the weighted averaged variant. If*
 399 *instead $b = b_N \rightarrow \infty$ with $b_N = o(\sqrt{N})$, the same limit holds provided the*
 400 *random constants Λ of Lemma 14 are bounded in probability uniformly in N*
 401 *along the sequence b_N .*

402 We separate the two cases because only the first is proved unconditionally, and
 403 the difference matters: Theorem 10 takes $b \asymp \log N$, so the interval we actually
 404 recommend lives in the second case. The proof of the fixed- b statement is
 405 almost sure and goes through Lemma 14, whose conclusion carries a finite
 406 random constant; for a triangular array in which b changes with N , “finite for
 407 each N ” is not enough, and the naive route around it — Cauchy–Schwarz on
 408 $\mathbb{E}\|H^{-1} - P_{t-1}\|_{\text{op}}\|\varepsilon_t\|$ — is exactly the bound that the proof of term (A) in
 409 Appendix A shows is too weak, by a factor $T^{1/2}$. Establishing the uniformity
 410 would require redoing Lemma 14 with explicit constants in b , which we have
 411 not done. What we can say is that the condition is about the remainder
 412 terms and not about the limit, that growing b makes the block scores *less*
 413 dependent rather than more, and that empirically coverage is flat in b over
 414 $b \in [10, 320]$ at $N = 200,000$ (Table 5(b)), which is the composition the two
 415 theorems need. Section 7 lists this as an open point rather than burying it.

416 The condition $b = o(\sqrt{N})$ is what makes the number of steps $T = N/b$
 417 large enough for the $1/t$ schedule to erase the initial condition: with $\gamma_t = 1/t$
 418 the effect of θ_0 decays like $1/T = b/N$, so its contribution to $\sqrt{N}(\theta_T - \theta^*)$ is
 419 of order $b/\sqrt{N}$. It is amply satisfied by the $b \asymp \log N$ of Theorem 10, and it
 420 is the reason Table 5(b) shows the point estimate degrading once b becomes
 421 a substantial fraction of $\sqrt{N}$.

422 **Proposition 7** (Blocking is free). *Under the assumptions of Theorem 6,*

$$\theta_T - \theta^* = -\frac{1}{N_T}H^{-1} \sum_{i \leq N_T} g_i + R_T, \quad \|R_T\| = o_p(N_T^{-1/2}),$$

423 *so the asymptotic variance in Theorem 6 is the same for every block length b .*

424 Proposition 7 matters for the design argument. It says that, for this
 425 algorithm, averaging the gradient over a block costs nothing statistically—the
 426 leading term is the same sample average of scores whatever b is—which is
 427 why we are free to choose b for the covariance estimator. It also answers

the objection that single-sample streaming updates are more sample-efficient than batched ones (Li et al., 2026): that is true for a step size $\gamma_t \asymp t^{-a}$ with $a < 1$, and false here.

3.4. The variance-inflation factor and the exact failure of dependence-blind intervals

Proposition 8 (Coverage of dependence-blind intervals). *Let $\kappa_j = [H^{-1}SH^{-1}]_{jj} / [H^{-1}\Gamma_0H^{-1}]_{jj}$ and let $\hat{\Gamma}_0 \rightarrow_p \Gamma_0$. Then the interval $\theta_{T,j} \pm z_{\alpha/2}\{[\bar{H}^{-1}\hat{\Gamma}_0\bar{H}^{-1}]_{jj}/N\}^{1/2}$ has asymptotic coverage*

$$2\Phi(z_{\alpha/2}/\sqrt{\kappa_j}) - 1.$$

The proof is one line, and the content is entirely in κ_j : dependence-blind intervals do not merely lose a little coverage, their coverage converges to a computable number below the nominal level. For the homogeneous autoregressive design of Section 5—covariates $x_t = \phi x_{t-1} + \dots$, errors $u_t = \psi u_{t-1} + \dots$, $x \perp u$ —one gets $\Gamma_k = \sigma_u^2(\phi\psi)^{|k|}\Sigma_x$ and $H = \Sigma_x$, hence

$$\kappa_j = \frac{1 + \phi\psi}{1 - \phi\psi} \quad \text{for every } j, \quad (12)$$

so at $\phi = \psi = 0.7$ a nominal 95% interval covers 74.9%. (12) is textbook geometric summation and we present it only because it makes Proposition 8 checkable with no free parameters (Figure 1); the qualitative failure has been observed before (Liu et al., 2023), whose simulated coverage “hovers around 85%” at $\kappa = 2$, against our prediction of 83.4%.

4. Estimating the long-run covariance from block gradients

The algorithm forms $\bar{g}_t(\theta_{t-1})$ once per block and then discards it. Retain it for two $O(d^2)$ accumulations. Over a retained window $\mathcal{T}$ of T' blocks (see below), define

$$\hat{S}^{\text{BM}} = \frac{b}{T'} \sum_{t \in \mathcal{T}} \bar{g}_t \bar{g}_t^\top, \quad \hat{\Lambda}_1 = \frac{1}{T' - 1} \sum_{t, t+1 \in \mathcal{T}} \bar{g}_t \bar{g}_{t+1}^\top, \quad \hat{S}^{\text{FT}} = \hat{S}^{\text{BM}} + b(\hat{\Lambda}_1 + \hat{\Lambda}_1^\top). \quad (13)$$

$\hat{S}^{\text{BM}}$ is the non-overlapping batch-means estimator at block resolution. $\hat{S}^{\text{FT}}$ is the two-scale (flat-top, or “lugsail”) correction. It has the same expectation as $2\hat{S}^{\text{BM}}(2b) - \hat{S}^{\text{BM}}(b)$ — Richardson extrapolation in $1/b$ — and coincides with it exactly if $\hat{\Lambda}_1$ is formed from disjoint rather than all adjacent block pairs; we use all adjacent pairs because it has the same expectation and a

smaller variance. Either way the expectation is that of a lag-window estimator with weights 1 for $|k| \leq b$ and $(2b - |k|)/b$ for $b < |k| < 2b$. **This device is not new:** it is the flat-top window of Politis and Romano (1995), refined by Politis (2011) and recast as “lugsail” by Vats and Flegal (2022), and Singh et al. (2025, Eqs. (15) and (17)) apply exactly these two forms to stochastic gradient descent under independent data. Nor is the $O(\rho^b)$ bias order a new mechanism: for a stationary sequence whose autocovariances decay geometrically it is the classical tail-sum property of a flat-top window. What we contribute is that it is available *at block-gradient resolution on a dependent stream*, where it is what makes $b \asymp \log N$ admissible and therefore what buys the rate, together with the streaming implementation below.

Confidence intervals. With $u_j = \bar{H}^{-1}e_j$ (one $O(d^2)$ product using the already-maintained inverse), report $\theta_{T,j} \pm z_{\alpha/2} \{u_j^\top \hat{S}^{\text{FT}} u_j / N\}^{1/2}$. Only the requested coordinates cost anything; the full $d \times d$ sandwich is never needed.

Retained window (a streaming detail that matters). Early blocks have θ_{t-1} far from θ^* and inflate (13). Discarding a fixed fraction requires knowing N in advance, which a streaming method does not. We instead keep two accumulators covering the current and previous dyadic windows $[2^k, 2^{k+1})$ and report from their union, so at any time at least half and at most three quarters of the blocks are retained and the discarded set is a prefix. This answers, in our setting, the open problem of an online implementation of equal-batch-size estimators raised by Singh et al. (2025, Remark 4).

Theorem 9 (L^2 rate). *Under the assumptions of Theorem 4, $\mathbb{E}\|\theta_t - \theta^*\|^2 = O(\log(N_t)/N_t)$, uniformly over $b \leq C \log N$.*

The L^2 statement is separate from the almost sure one because the drift analysis in Theorem 10 needs a moment bound and an almost sure rate with an unspecified random constant does not supply one. All remainder bounds in Appendix A are uniform over $b \leq C \log N$, which is the only regime in which we use them.

Theorem 10 (Long-run covariance estimation). *Let Assumptions 1 to 5 hold and let $b = b_N \rightarrow \infty$ with $b_N \leq C \log N$. Write $\Delta = \sum_{k \geq 1} k (\Gamma_k + \Gamma_k^\top)$. Then, entrywise,*

$$(i) \quad \mathbb{E}[\hat{S}^{\text{BM}}] = S - \Delta/b + O(\rho^b) + O(\tau_{b,N});$$

$$(ii) \quad \mathbb{E}[\hat{S}^{\text{FT}}] = S + O(\rho^b) + O(\tau_{b,N});$$

$$(iii) \quad \text{Var}[\hat{S}_{jk}^{\text{FT}}] = O(b/N); \text{ and consequently}$$

490 (iv) $\|\widehat{S}^{\text{FT}} - S\|_{\text{op}} = O_p(d\sqrt{b/N} + \rho^b + \tau_{b,N})$, which at fixed d and $b \asymp \log N$
 491 gives $\|\widehat{S}^{\text{FT}} - S\|_{\text{op}} = \widetilde{O}_p(N^{-1/2})$, whereas $\|\widehat{S}^{\text{BM}} - S\|_{\text{op}} = O_p(d\sqrt{b/N} +$
 492 $1/b)$ is minimised at $b \asymp (N/d^2)^{1/3}$ and gives $O_p((d^2/N)^{1/3})$.
 493 The factor d in (iv) is the cost of converting the entrywise bound (iii) into an
 494 operator norm; we make no attempt to remove it, and every rate statement in
 495 this paper is at fixed d . Here $\tau_{b,N} = (b+d)\log^2 N/N$ is the contribution of the
 496 drift $\theta_{t-1} \neq \theta^*$: the b comes from the parameter error and the d from the fact
 497 that the block curvature $\widehat{H}_t$ is a b -sample estimate of a $d \times d$ matrix, so that
 498 $\mathbb{E}\|\widehat{H}_t\delta\|_{\text{op}}^2 \lesssim (1+d/b)\|H\delta\|_{\text{op}}^2$. Both are dominated by the $\sqrt{b/N}$ standard
 499 deviation whenever $b+d = o(\sqrt{N})$ up to logarithms. Consequently $\sqrt{N}(\theta_{T,j} -$
 500 $\theta_j^*)/\{u_j^\top \widehat{S}^{\text{FT}} u_j\}^{1/2} \rightsquigarrow \mathcal{N}(0,1)$ and the intervals above have asymptotically
 501 exact coverage.

502 **Remark 3** (Why blocking gradients beats blocking iterates). Batch-means
 503 estimators for stochastic approximation are usually built from the *iterate*
 504 path (Chen et al., 2020; Zhu et al., 2023; Roy and Balasubramanian, 2023;
 505 Singh et al., 2025; Samsonov et al., 2025). Under dependent sampling this
 506 forces long blocks: the iterates θ_k move by $\gamma_k \asymp k^{-a}$, so a block must
 507 span $\sum_{k=t+1}^{t+b} \gamma_k$ of parameter movement before its average behaves like an
 508 independent draw, which requires $b \asymp N^a$ and yields rates of order $N^{-1/8}$
 509 (e.g. Samsonov et al., 2025, Prop. 2 and Cor. 2). Block *gradients* do not
 510 have this problem: their serial dependence is inherited from the data process,
 511 whose mixing time is a constant, and the iterate’s own movement enters
 512 only through the $O(b/N)$ drift term in Theorem 10. This is the substantive
 513 reason for the rate difference, and it is available only to an algorithm that
 514 forms block gradients and keeps a curvature estimate. It is also why the
 515 minimax rate $\Theta(N^{-(1-a)/2})$ for *Hessian-free* covariance estimation from an
 516 SGD trajectory (Ni and Huo, 2026) does not apply: a quasi-Newton method
 517 is outside that class.

518 **Remark 4** (Positive semidefiniteness). $\widehat{S}^{\text{FT}}$ is a difference of positive semidef-
 519 inite matrices and need not be positive semidefinite; the same caveat attaches
 520 to all flat-top and lugsail estimators (Singh et al., 2025, Remark 5). For
 521 marginal intervals only the scalars $u_j^\top \widehat{S}^{\text{FT}} u_j$ must be positive, and we fall
 522 back to $u_j^\top \widehat{S}^{\text{BM}} u_j$ if one is not. We report the frequency of that fallback in
 523 every experiment; it is 0.0% throughout.

524 **Proposition 11** (Gapping versus Bernoulli thinning at matched cost).
 525 Let $Y_i = \alpha(\theta^*; \xi_i) \Psi_i \Psi_i^\top$, $\Gamma_k^H = \text{Cov}(Y_0, Y_k)$, $S^H = \sum_k \Gamma_k^H$ and $\kappa_H =$
 526 $\|S^H\|_{\text{op}}/\|\Gamma_0^H\|_{\text{op}}$. From a stream of length N , take n rank-one curvature

updates either at a deterministic gap $m = N/n$ or by Bernoulli($p = n/N$) thinning of every index. Then, as $n \rightarrow \infty$,

$$n \operatorname{Var} [\hat{H}^{\text{gap}}] = \Gamma_0^H + O(\rho^m), \quad n \operatorname{Var} [\hat{H}^{\text{Bern}}] = \Gamma_0^H + \frac{1}{m} (S^H - \Gamma_0^H).$$

So the excess over the independent-data benchmark Γ_0^H/n is exponentially small in m for gapping and of order $(\kappa_H - 1)/m$ for Bernoulli thinning. Gapping also makes the per-block work deterministic. We verify both statements in Supplementary Fig. S2. We are explicit that this is a curvature-side improvement: $\bar{H}$ enters the limit distribution only through the rate ν_H of Theorem 3, so the end-to-end effect on $\hat{\theta}$ is second order, and we do not claim otherwise.

5. Experiments

Every number quoted below is generated from the saved results by `experiments/make_tables.py` and injected as a macro, so the prose cannot drift from the data. Monte Carlo standard errors are shown or stated everywhere.

Designs.. Four simulated dependent designs with $d = 20$, described in full in Appendix B.1: **AR-hom** (autoregressive covariates and autoregressive errors, closed-form S , the same variance inflation $\kappa = 2.92$ in every coordinate); **AR-het** (coordinatewise-varying persistence, closed-form S , κ_j ranging over $[1.00, 4.41]$, so a scalar correction cannot pass); **Markov** (covariates are a Markov-modulated Gaussian mixture with $d + 1$ regimes, hence non-Gaussian with discrete latent state and regime dependence, closed-form S , $\kappa_j \equiv 7.55$); **Logistic** (a correctly specified marginal logistic model with a serially dependent latent error, so θ^* is exactly the generating value while the score is autocorrelated; S from an independent long simulation, relative Monte Carlo error 7.5×10^{-3}). All covariate covariances are the ill-conditioned design of the base paper ($\operatorname{cond}(H)$ up to 400); Table 3 adds a well-conditioned counterpart, and two real hourly streams are described under Table 4.

Methods compared.. Our estimator BGSN; the same point estimator with plain batch means or with the i.i.d. plug-in variance (so that the interval, and only the interval, changes); the base independent-data streaming Newton method with its own curvature schedule ($m = 1$, Bernoulli $p = 1/d$) and its own prescribed plug-in variance; averaged SGD and streaming AdaGrad with tuned step sizes, each with the i.i.d. plug-in, with our long-run variance, or with random scaling (a studentisation that needs no variance estimate




Figure 1: Dependence-blind confidence intervals fail by a computable amount, and only estimating the long-run score covariance repairs them. (a) Coverage of nominal 95% intervals on the homogeneous autoregressive design as the dependence strength varies, against the parameter-free prediction of Proposition 8; error bars are ± 1.96 Monte Carlo standard errors. (b) Coverage of every method on that design at $\phi = \psi = 0.7$ ($\kappa = 2.92$), one row per method; the dotted line is the nominal level and the number at the right is the interval width relative to the infeasible oracle interval. Read it in three groups. The five rows at the top all account for the autocovariances of the score — the infeasible oracle, the offline reference, and the three streaming estimators that use our long-run covariance — and they are the rows that reach nominal, at width ≈ 1 . The next two use the instantaneous covariance the independent-data theory prescribes; they are narrower by the factor $\sqrt{\kappa} = 1.71$ and they miss, by the amount Proposition 8 predicts. The bottom four are first-order methods whose *point* estimates have not converged on this ill-conditioned design (their $\sqrt{N}$ error is 1.03 to 9.77 times the efficient value), so their coverage is dominated by that and not by which covariance they use; Table 3 is where they become informative about inference.

Table 1: Coverage of nominal 95% confidence intervals, interval width, and estimation error on four dependent designs ($d = 20$, $N = 200,000$, block length $b = 20$, $R = 300$ replications; $R = 60$ for the offline reference). *cov*: coverage, pooled over the d coordinates (Monte Carlo standard error ≤ 0.008 throughout). *w*: interval width divided by the width of the infeasible oracle interval. *err*: $\sqrt{N}$ RMSE divided by the asymptotic standard deviation, averaged over coordinates; 1.00 is efficient. κ is the variance inflation caused by dependence. Only methods that estimate the *long-run* score covariance reach nominal coverage; methods using the instantaneous covariance that the independent-data theory prescribes undercover by an amount predicted exactly by Proposition 8. **Read the first-order rows with care**: these designs are ill-conditioned (that is what streaming second-order methods are for), so those methods’ *point* estimates have not converged — their *err* column is several times 1 — and their coverage is dominated by that, not by which covariance they use. Table 3 repeats the AR-hom column on a well-conditioned design where every point estimator is efficient, which is where the first-order rows become informative about inference. The Markov column is the one design where *no* streaming method’s point estimate has converged at this N ; Table 2 resolves it by sweeping N against the oracle interval. ASGD is averaged stochastic gradient descent (Polyak–Ruppert averaging of the iterates); AdaGrad is the diagonally rescaled first-order method; *random scaling* is the variance-estimator-free studentisation of Lee et al. (2022).

method	AR-hom			AR-het			Markov			Logistic		
	cov	w	err	cov	w	err	cov	w	err	cov	w	err
<i>dependence</i>	$\kappa \in [2.92, 2.92]$			$\kappa \in [1.00, 4.41]$			$\kappa \in [7.55, 7.55]$			$\kappa \in [2.27, 2.32]$		
BGSN (this paper)	0.949	1.00	1.01	0.950	1.00	1.01	0.890	1.02	1.27	0.783	0.84	1.33
with plain batch means	0.939	0.96	1.01	0.944	0.98	1.01	0.856	0.92	1.27	0.768	0.81	1.33
with the i.i.d. plug-in variance	0.742	0.58	1.01	0.832	0.73	1.01	0.460	0.38	1.27	0.581	0.55	1.33
streaming Newton, i.i.d. theory	0.743	0.58	1.01	0.834	0.73	1.01	0.460	0.38	1.27	0.584	0.55	1.33
with our long-run variance	0.949	1.00	1.01	0.951	1.00	1.01	0.890	1.02	1.27	0.781	0.84	1.33
ASGD + our long-run variance	0.051	1.02	9.77	0.271	1.06	50.34	0.239	1.09	7.95	0.714	1.01	1.91
ASGD, i.i.d. plug-in	0.028	0.59	9.77	0.195	1.03	50.34	0.073	0.38	7.95	0.520	0.66	1.91
ASGD + random scaling	0.083	1.26	9.77	0.679	8.51	50.34	0.230	1.73	7.95	0.925	2.16	1.91
AdaGrad + our long-run variance	0.547	1.03	2.99	0.947	1.01	1.03	0.533	1.00	2.76	0.911	1.01	1.15
offline HAC (<i>not streaming</i>)	0.961	0.99	0.95	0.955	0.99	0.98	0.968	0.98	0.85	0.945	1.00	1.03
oracle variance (<i>infeasible</i>)	0.950	1.00	1.01	0.949	1.00	1.01	0.886	1.00	1.27	0.858	1.00	1.33




Figure 2: The long-run covariance estimator, and why the moving iterate does not spoil it. (a) Measured $\mathbb{E}[\hat{S}]/S$ against block length: plain batch means follow the closed-form bias $1 - \Delta/(bS)$ (no fitted constants), while the two-scale form is unbiased at every b , as Theorem 10 predicts under geometric mixing. (b) The algorithm’s block gradients decomposed exactly into the score at θ^* and the drift $\hat{H}_t(\theta_{t-1} - \theta^*)$. The drift’s contribution to $\hat{S}$ is 0.1% of S at $b = 20$ and 0.9% at $b = 320$, growing like $(b + d)/N$ as Theorem 10 predicts, and the estimator built from the algorithm’s iterates is indistinguishable from the one built at θ^* . This is why our block length need only grow like $\log N$ rather than like a power of N (Remark 3).

at all: it divides by a functional of the iterate path whose limit distribution is known and tabulated, Lee et al., 2022); an offline HAC reference (heteroskedasticity- and autocorrelation-consistent: the textbook two-pass sandwich with a Bartlett-kernel long-run covariance, which stores the whole stream and is therefore not a streaming method); and an infeasible oracle-variance interval that uses the true H and S . First-order baselines are given the *true* Hessian for their sandwich — an advantage no practitioner has — and their step-size constants are tuned by grid search on pilot streams disjoint from the evaluation streams (Appendix B.3). Our method has no learning rate to tune: its step size is fixed at $1/t$ by the preconditioner, with the step-length safeguard of (6) the only other free quantity and its constant fixed at $c_D = 1$ throughout.

5.1. The failure is exactly as predicted, and it is not small

Figure 1(a) sweeps the dependence strength and compares the measured coverage of dependence-blind intervals with the parameter-free prediction




Figure 3: Two real hourly streams. (a, b) Coverage of nominal 95% intervals under Protocol A (real covariates, synthetic autoregressive errors, exact target and exact conditional long-run covariance). The air-quality stream’s median estimated variance inflation is 9.9, so dependence-blind intervals fail badly there. (c) Coverage (left axis) against block length b (horizontal, log scale) on the traffic stream, with the free block-adequacy diagnostic $r_j = |u_j^\top (\hat{S}^{\text{FT}} - \hat{S}^{\text{BM}}) u_j| / (u_j^\top \hat{S}^{\text{BM}} u_j)$ — an estimate of the relative block-length bias of plain batch means, defined where it is first used in Section 5 — on the right axis. Three things to read. Our coverage is flat in b over a fortyfold range and stays within a few hundredths of the infeasible oracle-variance curve throughout, so the shortfall from 0.95 on an 8,000-observation segment is the point estimate and not the covariance. The dependence-blind curve sits far below at every b and does not close. And r_j is U-shaped, lowest near $b = 20$ –40: it rises for small b , where plain batch means are biased, and again for large b , where too few blocks remain to estimate anything stably — which is what makes it a usable guide to the block length, though not a calibration guarantee.

576 $2\Phi(z_{0.025}/\sqrt{\kappa}) - 1$. For κ up to 2.9 the largest absolute discrepancy is 0.0063,
577 of the same order as the Monte Carlo standard error: at $\kappa = 2.92$ the
578 prediction is 0.748 and the measurement is 0.742 ± 0.0114 . BGSN stays at
579 nominal over that range (0.949 at $\kappa = 2.92$), including at $\phi = \psi = 0$, where
580 all four intervals agree — the control showing that the machinery is not
581 simply widening intervals.

582 At the strongest dependence we tried, $\kappa = 6.2$, the measurement is 0.560
583 against a predicted 0.569, and the largest discrepancy over the whole sweep
584 is 0.0085. That gap is finite-sample, not a failure of the formula, and the
585 infeasible oracle interval says so: it covers only 0.946 there, so at $N = 200,000$
586 the central limit theorem itself is not yet accurate at that dependence level.
587 BGSN covers 0.945, i.e. it tracks the oracle rather than the formula — which
588 is the right diagnosis of where the remaining error lives.

589 The oracle comparison is worth stating as a general point, because it
590 separates two things that are easy to confuse. BGSN’s coverage tracks the
591 *infeasible oracle-variance* interval — built from the true H and the true
592 S — to within Monte Carlo error at every dependence level in the sweep.
593 Whatever undercoverage remains is therefore attributable to the central limit
594 theorem at finite N , not to our variance estimator. Plain batch means, by
595 contrast, fall 1–3 percentage points *below* the oracle once $\kappa \geq 4$, which is
596 exactly the $O(1/b)$ bias the two-scale correction removes.

597 Table 1 gives the full picture. On AR-hom, the base method’s interval
598 covers 0.743 and our interval covers 0.949 against a nominal 0.95, with the
599 infeasible oracle interval at 0.950 (which is the sanity check that the central
600 limit theorem itself, and our Monte Carlo harness, are right). The price is
601 width: our interval is $1.71 \times$ the dependence-blind one, which is exactly $\sqrt{\kappa}$
602 — there is no free lunch, the dependence-blind interval is simply the wrong
603 length. On the Logistic design, where κ ranges over $[2.27, 2.32]$, the pattern
604 is the same. AR-het is worth a separate look, because its κ_j varies across
605 coordinates by construction: the dependence-blind interval’s *per-coordinate*
606 coverage there ranges from 0.63 to 0.96, while ours ranges from 0.92 to 0.98.
607 No scalar inflation of the interval width could produce the first pattern from
608 the second; the whole matrix has to be estimated. Crucially, replacing *only*
609 the variance estimator repairs the base method (row “with our long-run
610 variance”), and replacing only the point estimator does not: this isolates the
611 inference component as the operative one.

612 The Markov and Logistic columns need a caveat of their own, and it is the
613 sharpest test of whether our variance estimator is really doing the work we
614 claim. Neither design’s point estimate has reached its asymptotic regime at
615 $N = 200,000$, for two different reasons. The Markov design’s regime chain has
616 an integrated autocorrelation time of 66 observations — the sum $1 + 2 \sum_{k \geq 1}$
617 of its autocorrelations, which is the factor by which dependence multiplies
618 the variance of a sample mean, so that N dependent observations carry about
619 as much information as $N/66$ independent ones — so its effective sample
620 size is a small fraction of N and the point estimate is 1.33 times the efficient
621 standard deviation. On the Logistic design the cause is the curvature rather
622 than the dependence: the weight $\alpha = \sigma'(x^\top \theta)$ is largest at $\theta_0 = 0$, where the
623 warm-up evaluates it, so $\bar{H}$ built there overestimates H and the early steps
624 are correspondingly too short; the point estimate is 1.33 times efficient and
625 the fitted standard error is biased low with it. An interval built from the
626 *correct* asymptotic variance must undercover in that situation, no matter who
627 supplied the variance, simply because $\sqrt{N}(\theta_T - \theta^*)$ is not yet close to its limit.
628 So on these two designs the question is not whether coverage is nominal but

629 whether our interval matches the interval that uses the exact $H^{-1}SH^{-1}$ at
 630 the same iterate. Table 2 and Figure 4 answer it: over a sweep in N both rise
 631 towards nominal together while the dependence-blind interval does not. On
 632 the Markov design the agreement is essentially exact: our coverage is within
 633 0.017 of the oracle's at every N , and at the top of the sweep, $N = 500,000$,
 634 both are 0.918 and 0.917 with the point estimate down to 1.12 times efficient,
 635 against 0.494 for the dependence-blind interval. The direction of travel is
 636 the whole point. Across the fivefold increase in N our coverage and the
 637 oracle's move together towards 0.95 and stay within 0.017 of each other,
 638 while the dependence-blind interval creeps up by a few hundredths towards
 639 its own asymptote, which Proposition 8 puts at $2\Phi(z_{0.025}/\sqrt{7.55}) - 1$, far
 640 below nominal. We stop the sweep at $N = 500,000$ rather than continuing
 641 to where coverage is nominal, because what the experiment is for is the gap
 642 between the first two columns, and Table 2 shows that gap is already at
 643 the level of the Monte Carlo standard error there. On the Logistic design
 644 there is a real gap, and it is ours: our coverage is 0.736 against the oracle's
 645 0.852 at $N = 100,000$, a difference of 0.116, falling monotonically to 0.036
 646 at $N = 500,000$, the top of the sweep. The cause is visible in Table 1: our
 647 fitted width there is 0.84 times the oracle's, i.e. the standard error is biased
 648 low, because $\bar{H}$ is an average over the whole stream and the early blocks
 649 contribute $\alpha = \sigma'(0) = 1/4$, its largest possible value, so $\bar{H}$ overestimates
 650 H and $\bar{H}^{-1}S\bar{H}^{-1}$ underestimates the sandwich. That bias decays with the
 651 Cesàro average of $\|\theta_t - \theta^*\|$, which is why the gap closes; it is a finite-sample
 652 property of averaging curvature along the path, it affects the base method
 653 identically, and no choice of long-run covariance estimator would remove
 654 it. That is the signature we want: our error tracks the oracle's at every
 655 sample size, and the dependence-blind error is the one that survives $N \rightarrow \infty$.
 656 The step-length safeguard bound at most 8 times per run anywhere in that
 657 sweep, and the count does not rise with N , which is the empirical content of
 658 Proposition 1.

659 We report every design in Table 1 at the same N rather than tuning N
 660 per design, which is why the Markov column looks worst there; Table 2 is
 661 where that column is resolved.

662 *The step-length safeguard is what makes the hard design run at all..* Table 5(i)
 663 varies the safeguard. With none, the autoregressive designs are barely af-
 664 fected (1.05 against 1.04 in relative variance) but the regime-switching design
 665 diverges: relative variance 1.1×10^{14} against 1.41 with the safeguard, and
 666 coverage 0.228 against 0.907. The safeguard is not doing this by damping
 667 the algorithm into submission: it binds a median of 7 times in 10,000 blocks

on that design and 3 times on AR-hom, and Table 2 shows the count flat as N grows by a factor of five.

The constant c_D behaves the way a safeguard constant should, which is to say it is irrelevant until it is not. On AR-hom, changing it by a factor of four in either direction moves nothing (1.04 at $c_D = 1/4$ and 1.04 at $c_D = 4$, against 1.04 at our default $c_D = 1$). On the design that actually needs the safeguard both directions cost something, and asymmetrically: $c_D = 1/4$ gives 1.59 against 1.41 — binding 21 times instead of 7, it now restrains steps that were fine — while $c_D = 4$ is too loose to prevent the blow-up at all and the relative variance is 4.3×10^7 . So c_D is a tuning constant with a real failure mode on one side, and we say so rather than advertising insensitivity we only observe on the easy designs.

The schedule shift of Appendix C.1 is the alternative the contraction condition suggests more directly, and it is measurably worse: 3.08 against 1.41 on the regime-switching design, because it slows all 10^4 steps to correct a handful. We report it because we implemented it first, and because the comparison is the evidence for the design we kept.

The first-order rows of Table 1 need a caveat, and it cuts against us if we hide it. These designs are ill-conditioned — $\text{cond}(H) = 400$ for AR-hom — which is precisely the regime streaming second-order methods exist for, and there averaged SGD’s *point* estimate has not converged: its root-mean-square error is 9.77 times the efficient value against 1.01 for BGSN. Its intervals are therefore uninformative about the inference question, whatever covariance they use. Two pieces of evidence show that this is about conditioning and not about our construction. First, within Table 1 itself: streaming AdaGrad, which rescales per coordinate, is efficient on AR-het (1.03) and with our long-run variance it covers 0.947 — so the construction transfers to a first-order method as soon as that method’s point estimate converges. Second, Table 3 repeats the AR-hom design with $\Sigma_x = I$, where every point estimator is efficient (1.07 for averaged SGD, 1.03 for BGSN) and only the covariance can decide coverage: averaged SGD with our long-run variance covers 0.931, the same estimator with the i.i.d. plug-in covers 0.721, and BGSN covers 0.942.

Coverage checks a single quantile, while the theorem claims a whole distribution, so we also report the Kolmogorov–Smirnov distance between the studentised error (the error divided by its own estimated standard error, which is what an interval implicitly compares to a normal quantile) $\sqrt{N}(\theta_{T,j} - \theta_j^*)/\{u_j^\top \hat{S}^{\text{FT}} u_j\}^{1/2}$ and the standard normal, taken per coordinate across replicates. On AR-hom the worst of the d coordinates gives 0.089 against a 5% critical value of 0.079: marginally outside, so at this number of replicates

the studentised distribution is distinguishable from normal, and we report that rather than rounding it in our favour. The dependence-blind studentisation gives 0.207, an order of magnitude further out. The comparison, not the test outcome, is the content: our studentisation is at the edge of detectable non-normality while the dependence-blind one is not close. The standard errors themselves are accurate: 0.997 times the truth on average.

5.2. The long-run covariance estimator behaves as the theory says

Figure 2(a) plots $\mathbb{E}[\hat{S}]/S$ against the block length against the *exact* prediction $\sum_{|k|<b}(1 - |k|/b)\Gamma_k/S$, which is closed form for this design (we plot the exact Bartlett-weighted sum rather than its $1/b$ expansion, so the comparison carries no approximation of its own). Plain batch means track it over the whole range, at $b = 20$ measured 0.934 against predicted 0.936, and the two-scale estimator is within Monte Carlo error of 1 for $b \geq 10$ (at $b = 20$: 0.999); the small residual at $b = 5$ is the $O(\rho^b)$ term, which is only $\rho^5 = 0.03$ there. This is Theorem 10(i)–(ii) with no fitted constants.

Supplementary Fig. S1 measures the rate. Panel (b) is the rate for S itself: under the $b \asymp \log N$ schedule the two-scale estimator’s relative error falls with fitted exponent -0.47 against the $-1/2$ of Theorem 10(iv), while plain batch means under their own optimal $b \asymp N^{1/3}$ schedule give -0.36 against $-1/3$.

Panel (a) puts our estimator and an iterate-path estimator on the same object, because an interval uses the whole sandwich $\Sigma = H^{-1}SH^{-1}$ rather than S alone. We compare our composite $\bar{H}^{-1}\hat{S}^{\text{FT}}\bar{H}^{-1}$, whose relative error falls with fitted exponent -0.43 , with overlapping batch means on the *iterate* path at the block length its analysis prescribes — the estimator of Σ that uses no gradients and no curvature.

Two things must be said carefully here, because the comparison is easy to overstate. First, the iterate-path estimator’s relative error is flat over our range: the fitted exponent is 0.00, and we do *not* read that as its rate. Its published rate is $N^{-1/8}$, which over a sixteenfold increase in N predicts a reduction by a factor $16^{-1/8} = 0.71$ — a change small enough that our range cannot separate it from zero, and our fitted exponent should be treated as uninformative about it rather than as a measurement. Second, what our range *does* establish is the level, and the level is the practical statement: over the whole range the iterate-path estimator’s relative error in Σ sits near 0.86, which is not a usable standard error whatever its asymptotic rate, while ours falls from 0.17 to 0.05. We are comparing an estimator that has entered its asymptotic regime with one that has not, and the honest claim is about the

sample sizes a practitioner would actually have, not about which exponent is larger.

Figure 2(b) addresses the objection that a block average must span enough parameter movement to behave like an independent draw — the reason iterate-based batch means need $b \asymp N^{3/4}$ (Samsonov et al., 2025). We decompose the algorithm’s block gradients exactly into the score at θ^* and the drift $\widehat{H}_t(\theta_{t-1} - \theta^*)$ and evaluate the long-run covariance of each. The drift component contributes 0.0010 of S at $b = 20$ and grows like $\tau_{b,N} = (b + d) \log^2 N/N$, confirming Remark 3: because the blocks carry gradients rather than iterates, the block length is set by the data’s mixing time, not the optimiser’s.

5.3. Curvature schedules

Supplementary Fig. S2 compares deterministic gapping with Bernoulli thinning at a matched number of rank-one updates, with the optimisation frozen at θ^* so that only estimation quality is measured. The excess variance over the independent-data benchmark follows the two predicted curves — $O(\rho^m)$ for gapping, $(\kappa_H - 1)/m$ for thinning — over more than an order of magnitude in m . As stated in Section 7, this is a curvature-side result: Table 5(f) shows that the end-to-end effect on coverage is within Monte Carlo error, which is what the theory predicts, since $\bar{H}$ enters the limit only through a rate.

Where gapping does pay is cost, and Table 1 measures it. Compare BGSN with the row “base method with our long-run variance”: the two use the same variance estimator and the same number of rank-one curvature updates per observation ($1/d$), and they cover the same (0.949 against 0.949 on AR-hom). They differ only in *how* those updates are placed — on a deterministic gap of $m = d$ against Bernoulli retention with $p = 1/d$. At these settings the wall-clock difference is $1.04\times$, which is within our timing noise, so we do not claim a speed-up: a gap avoids drawing and testing a random variable per observation, but that saving is small next to the $O(d^2)$ update it guards. The case for gapping is therefore the one Proposition 11 makes and Supplementary Fig. S2 measures — a strictly better curvature estimate at a matched number of updates, $O(\rho^m)$ excess variance against $(\kappa_H - 1)/m$ — together with the fact that it makes the per-block cost deterministic rather than random. It is a better cost knob, not a faster one.

5.4. Ablations

Supplementary Fig. S4 and Table 5 vary one design choice at a time. (a) Coverage rises monotonically to nominal with N , and the gap to the

783 dependence-blind interval does *not* close — it converges to the predicted
 784 0.749. (b) Coverage as a function of b is flat for the two-scale estimator
 785 over $b \in [10, 320]$ and degrades at both ends for plain batch means (bias at
 786 small b , noise at large b): the practical content of the $O(\rho^b)$ -versus- $O(1/b)$
 787 distinction. (c) The curvature warm-up matters, though less than it did before
 788 the step-length safeguard absorbed part of its job: with essentially no warm-
 789 up, coverage on this design is 0.928 against 0.949 at the default, and above
 790 $c_w \approx 25$ the result is insensitive. (c') is where the warm-up is load-bearing:
 791 on the regime-switching design the relative variance is 2.19 at $c_w = 50$ against
 792 1.41 at $c_w = 200$, because there $50d$ consecutive observations span only about
 793 thirty regime visits and leave the curvature estimate rank-starved. That is
 794 what sets our default, and it is a statement about effectively distinct design
 795 points rather than about observations. (d) Results are insensitive to the
 796 ridge exponent q_r over the whole admissible range, and to the ridge constant
 797 c_l from 0 — no ridge at all — up to our default 0.01; at $c_l = 0.1$ coverage
 798 slips by two points and at $c_l = 1$ the ridge dominates the data terms and the
 799 estimator is no longer efficient. So the ridge is safe to leave alone but not safe
 800 to enlarge, and the honest reading is that the default sits at the top of the
 801 flat region rather than in the middle of it. That $c_l = 0$ works as well as the
 802 default is worth stating plainly: the ridge is in the algorithm so that Lemma 2
 803 can hold for *every* t without an assumption on the iterates, not because
 804 the runs need it. (e) Results are insensitive to the initial curvature scale λ_0
 805 over eight orders of magnitude, which matters because that is a parameter a
 806 user would otherwise have to think about. (f) Gap and thinning behave as
 807 Proposition 11 predicts and do not move coverage, which is what the theory
 808 says: $\bar{H}$ enters the limit only through a rate. (g) At strong dependence
 809 ($\phi = \psi = 0.85$, $\kappa = 6.2$) the two estimators separate in exactly the way the
 810 $O(\rho^b)$ -versus- $O(1/b)$ distinction predicts. Plain batch means need a long
 811 block: coverage rises from 0.921 at $b = 20$ to 0.946 at $b = 160$. The two-scale
 812 correction does not: it is already at 0.944 at $b = 20$ and stays there (0.945 at
 813 $b = 160$). The free diagnostic r_j falls from 0.177 to 0.086 over the same range,
 814 which is what makes it usable as a guide to whether plain batch means would
 815 have been adequate. So the correction is what buys the short block, and the
 816 short block is what keeps the estimator cheap. (h) The weighted averaged
 817 variant is slightly better at small N and slightly worse at large N ; both are
 818 reported. The non-positive-definiteness fallback of Remark 4 never triggered
 819 anywhere: the rate is 0 in every simulation and 0.0% of coordinate-replicate
 820 pairs on the real streams, where an earlier version of these experiments did
 821 not record it at all and where we would have expected it if anywhere, since
 822 there the lag-one correction is largest. It never triggered (0.0% of all runs,

823 all panels).

824 5.5. *The correct interval is also the cheaper one*

825 The authoritative statement is the operation count, because it is exact
826 and independent of the machine. The streaming pass grows as $d^{1.00}$ for BGSN
827 and as $d^{1.86}$ once the i.i.d. plug-in variance is added, because the plug-in score
828 covariance costs $O(d^2)$ per *observation* whereas our long-run covariance costs
829 $O(d^2)$ per *block*; evaluating the whole count at $N = 10^8$ gives $d^{1.01}$ against
830 $d^{1.86}$. Those are the design claim, and they are arithmetic rather than a fit.

831 The exponents fitted at the N of Table 6 are *not* the streaming cost,
832 and we say so rather than quoting them as if they were: with $c_w = 200$ and
833 $N = 200,000$ the one-off $O(c_w d^3)$ warm-up dominates the $O(dN)$ pass for
834 $d \geq 40$ (the crossover is $N \gtrsim c_w d^2$), so both measured exponents sit near 2.
835 What is robust at this N is the *level*: the plug-in variant costs more than BGSN
836 in wall-clock at every d we tried, by a factor between 1.4 and 3.4 (median
837 2.7), and in the main configuration BGSN runs in 0.37 times the time of the
838 dependence-blind streaming Newton method — that is, our valid intervals
839 cost less than the invalid ones, not more. The measured wall-clock exponents,
840 on a shared and heavily loaded machine, are $d^{1.95}$ and $d^{1.85}$; they are noisier
841 than the counts and we report them only as corroboration. Supplementary
842 Fig. S3 plots the same timings against d , where the message is the vertical
843 separation rather than either slope.

844 5.6. *Real streams*

845 Table 4 reports two protocols on hourly traffic volume on Interstate 94
846 and hourly PM2.5 at twelve Beijing monitoring sites. After a fixed regression
847 adjustment (weather plus calendar harmonics) the residual lag-one autocor-
848 relation is 0.76 and 0.91 respectively, and the median estimated variance
849 inflation across coordinates is 5.8 for traffic and 9.9 for air quality. Table 4
850 reports, for each stream, the pooled prediction $d^{-1} \sum_j \{2\Phi(z_{0.025}/\sqrt{\kappa_j}) - 1\}$
851 next to the measurement, so the comparison is against the quantity the table
852 actually measures rather than a single representative κ .

853 Protocol A keeps the real (strongly dependent) covariate stream and
854 replaces the response by $x_t^\top \theta^* + u_t$ with u an autoregression whose coefficient
855 equals that of the real residuals. This gives an *exact* target and an exact
856 conditional long-run covariance while retaining real covariate dependence:
857 BGSN covers 0.818 (traffic) and 0.893 (air quality) against nominal 0.95, while
858 the dependence-blind interval covers 0.597 and 0.543. Protocol B uses the real
859 response and takes the full-stream M -estimator as the target; because that
860 target is itself estimated from a superset of each segment, the correct nominal

level for the comparison is $2\Phi(z_{0.025}/\sqrt{1-1/R}) - 1$, which we compute exactly and report in the table header. BGSN covers 0.795 and 0.858. Neither protocol alone would settle the matter—A has an exact target but a synthetic error process, B has a real error process but an estimated target—and they agree.

Figure 3(c) shows why the block length is the one parameter a practitioner must think about on real data, and that the free diagnostic $r_j = |u_j^\top(\hat{S}^{\text{FT}} - \hat{S}^{\text{BM}})u_j|/(u_j^\top \hat{S}^{\text{BM}} u_j)$ flags the problem. The denominator is the batch-means quadratic form, not the two-scale one: $\hat{S}^{\text{BM}}$ is a sum of outer products and so its quadratic forms are non-negative, whereas those of $\hat{S}^{\text{FT}}$ can be negative (Remark 4), and an earlier version of this diagnostic divided by the latter and produced meaningless numbers on the real streams. On the traffic stream r_j is 0.16 at the $b = 40$ we use, and the sweep in Figure 3(c) shows what it is good for: r_j is U-shaped in b , rising both for short blocks, where plain batch means carry the $O(1/b)$ bias, and for long ones, where too few blocks remain for a stable estimate. Our coverage over that fortyfold range in b is flat and stays within a few hundredths of the infeasible oracle-variance curve, which is the evidence that the shortfall from nominal on these segments is the point estimate rather than the covariance. So r_j is a guide to choosing b , not a calibration guarantee, and we do not claim it predicts coverage.

6. Related work

This paper sits at the intersection of three literatures, and the honest statement of what is new requires being precise about each.

Streaming second-order methods with $O(dN)$ cost. Inversion-free stochastic Newton methods maintain a running inverse curvature estimate by rank-one Sherman–Morrison updates (Sherman and Morrison, 1950) and have been developed for logistic regression (Bercu et al., 2020), Gauss–Newton problems (Cénac et al., 2025), the geometric median (Godichon-Baggioni and Lu, 2024), and general conditioned stochastic gradient descent (Leluc and Portier, 2023; Boyer and Godichon-Baggioni, 2023); Chau et al. (2024) survey the inversion-free idea more broadly, and Byrd et al. (2016); Wang et al. (2017); Moritz et al. (2016) give the stochastic quasi-Newton and L-BFGS counterparts. Our starting point is Godichon-Baggioni and Werge (2025), who thin the rank-one updates with a Bernoulli variable and take mini-batches of size d to reach $O(dN)$ total cost with asymptotic efficiency, and Godichon-Baggioni et al. (2026), who reach the same budget by masking Hessian columns. **The $O(dN)$ budget and the idea of skipping curvature updates are**

Table 2: On the two designs whose point estimate has not converged, our interval tracks the infeasible oracle interval at every sample size ($d = 20$, $b = 20$; R decreases from 200 at the smallest N to 150 at the largest, since what is being measured is the gap between the first two columns and a Monte Carlo standard error near 0.01 is ample for it). *ours* is BGSN; *oracle* uses the exact $H^{-1}SH^{-1}$ at the same iterate and is infeasible; *blind* uses the i.i.d. plug-in. *err* is the root-mean-square error divided by the efficient standard error, so $\text{err} = 1$ means the point estimate has reached the asymptotic regime, and *clips* counts activations of the step-length safeguard out of N/b blocks. At $N = 200,000$ neither design is nominal because $\text{err} > 1$: the central limit theorem has not taken hold, which no variance estimator can repair. What our estimator is responsible for is the difference between the *ours* and *oracle* columns, and the largest such difference over both designs and all four sample sizes is 0.116. The *blind* column is the one that does not converge. Note that *clips* does not grow with N , which is the empirical content of Proposition 1.

design	N	ours	oracle	blind	err	clips
Markov	100,000	0.884	0.867	0.466	1.33	8
	200,000	0.881	0.880	0.472	1.28	8
	500,000	0.918	0.917	0.494	1.12	7
Logistic	100,000	0.736	0.852	0.537	1.33	2
	200,000	0.773	0.853	0.574	1.27	2
	500,000	0.857	0.893	0.669	1.19	2

Table 3: A well-conditioned dependent design isolates the inference question ($\Sigma_x = I$, so $\text{cond}(H) = 1.00$; otherwise identical to the AR-hom column of Table 1: $d = 20$, $N = 200,000$, $b = 20$, $\kappa = 2.92$, $R = 300$). Here every point estimator is efficient ($\text{err} \approx 1$), so coverage is decided by the covariance estimate alone. Averaged SGD with our long-run variance is calibrated; the same estimator with the i.i.d. plug-in is not, by the factor $\sqrt{\kappa} = 1.71$ in width that Proposition 8 predicts. On the ill-conditioned designs of Table 1 the first-order baselines cannot be read this way, because there their point estimates have not converged.

method	cov	w	err
BGSN (this paper)	0.942	1.00	1.03
with plain batch means	0.933	0.97	1.03
with the i.i.d. plug-in variance	0.743	0.59	1.03
streaming Newton, i.i.d. theory	0.740	0.59	1.03
ASGD + our long-run variance	0.931	1.00	1.07
ASGD, i.i.d. plug-in	0.721	0.59	1.07
ASGD + random scaling	0.917	1.28	1.07
AdaGrad + our long-run variance	0.943	1.00	1.03
oracle variance (<i>infeasible</i>)	0.941	1.00	1.03




Figure 4: Separating “the variance estimate is wrong” from “the central limit theorem has not taken hold” on the two designs whose point estimate has not converged at the sample sizes of Table 1. **(a)** The regime-switching design, whose covariate chain has an integrated autocorrelation time of about 66 observations. **(b)** The dependent-logistic design, where instead the curvature weight $\alpha = \sigma'(x^\top \theta)$ is largest at the θ_0 the warm-up evaluates it at. In both, the infeasible curve uses the exact $H^{-1}SH^{-1}$ at the same iterate, so the vertical distance between it and ours is everything our covariance estimator is responsible for: that distance is at most 0.017 on the Markov design and falls from 0.116 to 0.036 on the logistic one. Ours and the infeasible curve rise towards nominal together; the dependence-blind curve does not. **(c)** Why both left panels start low: the root-mean-square error of the point estimate, divided by the efficient standard error, is still above one, so $\sqrt{N}(\theta_T - \theta^*)$ is not yet close to its limit and no choice of variance can give nominal coverage.

898 **theirs, not ours;** [Chen et al. \(2024\)](#) independently thin Hessian *entries*
899 with probability $O(1/d^2)$ in a zeroth-order method. All of this work assumes
900 independent observations. Most of it proves central limit theorems and stops
901 there, but not all: for averaged first-order algorithms under independent data
902 [Godichon-Baggioni \(2019\)](#) gives a recursive online estimator of the asymptotic
903 variance with almost sure and L^2 rates, and [Chen et al. \(2020\)](#); [Zhu et al.](#)
904 [\(2023\)](#); [Godichon-Baggioni and Lu \(2024\)](#) supply plug-in and batch-means
905 constructions. **Online estimation of the *instantaneous* sandwich under**
906 **independent data is therefore not new either;** what is missing is the
907 long-run object under dependence, and the streaming second-order setting.
908 Our contribution on this axis is (i) the analysis of the rank-one recursion when
909 the accumulated terms are dependent, (ii) replacing Bernoulli thinning by a
910 deterministic gap, which is strictly better at matched cost ([Proposition 11](#)),
911 and (iii) turning the central limit theorem into usable intervals.

Table 4: Real hourly streams. Protocol A: real covariates, synthetic autoregressive errors whose coefficient equals the AR(1) coefficient of the real residuals, so the target and the long-run covariance are exact conditionally on the covariate path. Protocol B: fully real response, target the full-stream M -estimator; because that target is estimated from a superset of each segment, the correct nominal level for this comparison is $2\Phi(z_{0.025}/\sqrt{1-1/R}) - 1$, reported in the header. Traffic: $n = 8000$ per segment, $b = 40$. Air quality: $n = 10000$ per segment, $b = 50$, pooled over 12 monitoring sites.

method	traffic (I-94)		air quality (PM2.5)	
	A: exact truth	B: full-stream target	A: exact truth	B: full-stream target
<i>nominal</i>	0.950	0.976	0.950	0.994
BGSN (this paper)	0.818	0.795	0.893	0.858
with plain batch means	0.814	0.795	0.892	0.816
with the i.i.d. plug-in variance	0.597	0.682	0.543	0.361
streaming Newton, i.i.d. theory	0.580	0.682	0.500	0.365
ASGD + our long-run variance	0.564	0.523	0.331	0.330
ASGD, i.i.d. plug-in	0.391	0.455	0.198	0.212
ASGD + random scaling	0.571	0.568	0.422	0.486
oracle variance (<i>infeasible</i>)	0.773	–	0.788	–
<i>stream diagnostics</i>				
residual AR(1) coefficient		0.756	0.909 (median over sites)	
R^2		0.782	0.327	
block-adequacy r_j	0.163	0.271	0.186	0.302

912 *Online inference for stochastic approximation..* Under independent sampling
913 this is a mature area. [Chen et al. \(2020\)](#) give the plug-in and increasing-
914 batch-means variance estimators for averaged stochastic gradient descent;
915 [Zhu et al. \(2023\)](#) make batch means fully recursive with growing overlapping
916 windows and obtain $\|\hat{\Sigma} - \Sigma\| = O(N^{-1/8})$; [Singh et al. \(2025\)](#) show equal
917 batch sizes are preferable and apply the flat-top bias correction; [Ni and Huo](#)
918 [\(2026\)](#) improve the batch-means rate to $N^{-1/6}$ and establish the minimax
919 rate $\Theta(N^{-(1-a)/2})$ for *Hessian-free* covariance estimation from the iterate
920 path; [Wei et al. \(2026\)](#) give a fully online de-biased estimator that needs no
921 second-order information and reaches $N^{(a-1)/2}\sqrt{\log N}$, i.e. that minimax rate
922 up to a logarithm, so the Hessian-free class is close to saturated. [Fang et al.](#)
923 [\(2018\)](#); [Fang \(2019\)](#) bootstrap; [Su and Zhu \(2023\)](#) split the trajectory; [Lee](#)
924 [et al. \(2022\)](#); [Li et al. \(2026\)](#); [Chen et al. \(2024\)](#) studentise by a random scaling
925 matrix and estimate nothing. [Leung and Chan \(2026\)](#) and [Wu \(2009\)](#) design
926 recursive long-run variance estimators directly. **None of these is a claim**

Table 5: Ablations on the homogeneous autoregressive design ($d = 20$, $\kappa = 2.92$, $R = 250$). *cov* columns are coverage of nominal 95% intervals using the two-scale, plain batch-means, and i.i.d. plug-in variance respectively; *eff* is the empirical variance of the point estimate divided by its asymptotic value (1.00 is efficient); *adeq* is the free block-adequacy diagnostic r_j , an estimate of the relative block-length bias of plain batch means. Panel (c') is on the regime-switching design and is the reason the default warm-up is $c_w = 200$ rather than 50: there, $50d$ consecutive observations span only about thirty regime visits, the curvature estimate is still rank-starved, and the $1/t$ schedule diverges. The warm-up must span enough effectively distinct design points, which is not the same as enough observations. Panel (i) varies how the blocked step is safeguarded: *none* is the raw $1/t$ step, *cap* is the step-length safeguard of (6) at $c_D = 1$ (our default) with *cap 0.25* and *cap 4* changing c_D by a factor of four either way, *shift* is the rejected schedule shift of Appendix C.1, and *both* applies the two together; *clips* counts how many of the 10,000 blocks the safeguard acted on. *eff* is in scientific notation where the unsafeguarded step diverges. Panel (j) is on a short stream ($N = 4000$, $d = 11$, the length of a real-data evaluation segment) and varies the cap on the warm-up as a fraction of the stream; $\varpi_w = 1$ means uncapped, which spends 2200 of the 4000 observations on curvature and leaves 45 blocked steps.

panel	setting	eff	cov (2-scale)	cov (BM)	cov (plug-in)
(a) length N	$N = 25,000$	1.23	0.930	0.922	0.710
	$N = 50,000$	1.08	0.937	0.931	0.731
	$N = 100,000$	1.04	0.947	0.938	0.733
	$N = 200,000$	1.04	0.947	0.937	0.736
	$N = 400,000$	1.10	0.941	0.934	0.729
(b) block b	$b = 5$	1.02	0.948	0.905	0.736
	$b = 10$	1.03	0.949	0.926	0.737
	$b = 20$	1.04	0.947	0.937	0.736
	$b = 40$	1.04	0.946	0.943	0.737
	$b = 80$	1.05	0.946	0.944	0.738
	$b = 160$	1.06	0.946	0.944	0.734
	$b = 320$	1.07	0.943	0.944	0.738
	$b = 640$	1.10	0.938	0.942	0.733
(c) warm-up	$c_w = 0.5$	1.13	0.928	0.921	0.724
	$c_w = 5$	1.06	0.941	0.931	0.730
	$c_w = 25$	1.04	0.944	0.936	0.737
	$c_w = 50$	1.04	0.944	0.935	0.741
	$c_w = 100$	1.01	0.949	0.941	0.739
	$c_w = 200$	1.04	0.947	0.938	0.733
	$c_w = 500$	1.12	0.937	0.927	0.713
(c') warm-up, Markov	$c_w = 50$	2.19	0.848	0.812	0.452
	$c_w = 100$	1.76	0.879	0.851	0.467
	$c_w = 200$	1.41	0.907	0.877	0.495
	$c_w = 500$	1.21	0.923	0.900	0.536
(d) ridge	$c_t = 0, q_r = 0.2$	1.04	0.949	0.940	0.741
	$c_t = 0.001, q_r = 0.2$	1.04	0.949	0.940	0.740
	$c_t = 0.01, q_r = 0.2$	1.04	0.947	0.937	0.736
	$c_t = 0.1, q_r = 0.2$	1.04	0.931	0.920	0.703
	$c_t = 1, q_r = 0.2$	3.07	0.605	0.591	0.405
	$c_t = 0.01, q_r = 0.05$	1.03	0.943	0.931	0.729
	$c_t = 0.01, q_r = 0.1$	1.03	0.946	0.934	0.732
	$c_t = 0.01, q_r = 0.24$	1.04	0.948	0.938	0.737

Table 6: Cost. Wall-clock milliseconds for one pass over $N = 200,000$ observations, single-threaded, minimum over 7 interleaved repetitions (interleaved so that a change in the load of this shared machine cannot be mistaken for a difference between methods; the minimum because load can only make a measurement slower). s_{time} and s_{flop} are the exponents in d fitted on the four largest d to the times and to the exact operation counts at this N . Both are dominated by the one-off $O(c_w d^3)$ warm-up here and should *not* be read as the streaming cost: s_{str} is the same exact count with the warm-up excluded and s_{∞} is it at $N = 10^8$, where the warm-up is negligible and (9) reduces to $O(dN)$. Those last two columns are the design claim: the streaming pass is linear in d for BGSN and quadratic once the i.i.d. plug-in variance is added, because the plug-in score covariance costs $O(d^2)$ per *observation* while our long-run covariance costs $O(d^2)$ per *block*. The measured *levels* say the same thing and are robust to load: the plug-in variant costs between 1.4 and 3.4 times more than BGSN, at every d here.

method	$d=5$	$d=10$	$d=20$	$d=40$	$d=80$	$d=160$	s_{time}	s_{flop}	s_{str}	s_{∞}
BGSN (long-run interval)	7.4	6.2	10.3	21.4	115.4	534.9	1.95	2.07	1.00	1.00
with the i.i.d. plug-in variance	10.2	10.9	31.8	73.5	313.9	1392.3	1.85	2.02	1.86	1.86
streaming Newton, Bernoulli $p = 1/d$	9.4	7.7	11.2	27.0	111.2	533.9	1.88	2.07	–	–
streaming Newton, every observation ($p = 1$)	20.6	28.8	78.4	260.7	1135.1	4608.5	1.98	1.97	1.97	1.97
ASGD (first order, random scaling)	4.2	3.5	5.5	8.3	20.2	51.5	1.09	–	–	–
offline HAC, $L = 100$ (<i>not streaming</i>)	151.1	263.2	604.1	2303.0	–	–	–	–	–	–

we duplicate; what we take from them is the batch-means template and the flat-top correction, and what we add is the block-*gradient* construction and its consequence for the rate (Remark 3).

Two recent papers give online covariance estimation for streaming *second-order* methods: Kuang et al. (2026) for a sketched Newton method and Wang et al. (2026) for Nesterov-accelerated sketching. Both are under independent data, both target the instantaneous sandwich, and both require the gradient noise to be a martingale difference with respect to the algorithm’s filtration (Kuang et al., 2026, Assumption 3.2) — exactly the condition that temporal dependence destroys. Kuang et al. (2026) remark that their analysis extends to conditioned stochastic gradient methods generally, which would include an independent-data version of our algorithm; the martingale assumption is what keeps that remark from covering the present setting.

Stochastic approximation on dependent data.. Convergence of stochastic approximation under Markovian or mixing noise is classical (Benveniste et al., 1990; Kushner and Yin, 2003; Borkar, 2008; Andrieu et al., 2005; Fort, 2015; Karimi et al., 2019), and modern rates are available for Markov-chain

944 gradient descent (Sun et al., 2018; Doan et al., 2020; Even, 2023; Duchi
 945 et al., 2012; Nagaraj et al., 2020; Kowshik et al., 2021; Agarwal and Duchi,
 946 2013). Godichon-Baggioni et al. (2023) analyse streaming stochastic gradient
 947 methods with time-dependent and biased gradients, and identify blocked or
 948 growing mini-batches as the device that breaks dependence; they give L^2
 949 rates but no central limit theorem and no inference. **Blocked gradients as
 950 a dependence-breaking device are theirs.**

951 On the inference side, Liu et al. (2023) give a multiplier bootstrap for
 952 ϕ -mixing streams using alternating growing blocks, and establish the long-run
 953 sandwich limit and, qualitatively, the failure of dependence-blind intervals; Li
 954 et al. (2023) give a functional central limit theorem and the semiparametric
 955 efficiency bound $H^{-1}SH^{-1}$ under Markovian sampling, with random scaling
 956 for inference; Samsonov et al. (2025) give Berry–Esseen bounds and an
 957 overlapping-batch-means variance for *linear* stochastic approximation with
 958 Markovian noise; Roy and Balasubramanian (2023) give an online batch-
 959 means estimator of the Markovian long-run sandwich; Li et al. (2026) treat
 960 controlled Markov chains with random scaling. **The long-run sandwich
 961 limit, the efficiency bound, and the existence of online long-run
 962 covariance estimators under dependence are all prior art** — for
 963 first-order methods. Every one of these works is first-order; none maintains
 964 a curvature inverse, and the two that estimate a long-run covariance do so
 965 from the iterate path, at $O(N^{-1/8})$, with block lengths that must grow like
 966 $N^{3/4}$ (Samsonov et al., 2025, Prop. 2, Cor. 2). Remark 3 explains why block
 967 gradients escape that constraint, and Theorem 10 quantifies the gain.

968 *Long-run variance estimation..* The statistical device is heteroskedasticity-
 969 and autocorrelation-consistent covariance estimation (Newey and West, 1987;
 970 Andrews, 1991; Andrews and Monahan, 1992; Newey and West, 1994; White,
 971 1980, 1984), its fixed- b alternative (Kiefer and Vogelsang, 2002, 2005; Sun
 972 et al., 2008), and the batch-means literature from simulation output analysis
 973 and Markov chain Monte Carlo (Glynn and Whitt, 1991, 1992; Damerdj, 1994;
 974 Flegal and Jones, 2010; Jones et al., 2006; Vats et al., 2019). The
 975 two-scale correction we use is the flat-top lag window of Politis and Romano
 976 (1995), refined by Politis (2011) and recast as “lugsail” lag windows by Vats
 977 and Flegal (2022); Singh et al. (2025) bring it to stochastic gradient descent
 978 under independent data, where their Eqs. (15) and (17) coincide with our
 979 (13) and their Corollary 1 shows the convergence *rate* is unchanged. **We do
 980 not claim the correction.** We claim the bias order it attains at block-
 981 gradient resolution under geometric mixing, which is the property that makes
 982 $b \asymp \log N$ admissible, and the streaming implementation. Block bootstrap

and subsampling (Künsch, 1989; Politis et al., 1999; Lahiri, 2003) are the resampling counterparts.

Asymptotics for dependent M -estimation. The population statements we use are standard: sandwich asymptotics for M -estimators and generalised method of moments (Huber, 1967; Hansen, 1982; White, 1982; van der Vaart, 1998), misspecified models with dependent observations (Domowitz and White, 1982), marginal models for correlated responses (Liang and Zeger, 1986; Zeger and Liang, 1986) — the setting our logistic design imitates — and central limit theorems for mixing sequences (Rosenblatt, 1956; Ibragimov, 1962; Herrndorf, 1984; Wooldridge and White, 1988; Doukhan, 1994; Bradley, 2005; Merlevède et al., 2006; Rio, 2017; Dedecker et al., 2007). The technical inputs to our proofs are Davydov’s covariance inequality (Davydov, 1968; Rio, 1993), the Bernstein-type inequality for weakly dependent sequences of Merlevède et al. (2011), the almost sure behaviour of stochastic algorithms (Pelletier, 1998, 2000; Fabian, 1968; Mokkadem and Pelletier, 2011; Robbins and Siegmund, 1971), and the functional central limit theorem of Herrndorf (1984).

7. Limitations

We list the boundaries of the result plainly, because several of them are load-bearing.

Memory is $O(d^2)$, not $O(d)$. The $O(dN)$ claim is about *time*. The method stores a $d \times d$ inverse curvature estimate and two $d \times d$ covariance accumulators, so memory is $O(d^2)$; this is inherited from the independent-data method we extend and is not something we improve. For d in the thousands this is the binding constraint, and a sketched or diagonal-plus-low-rank curvature representation—as in Kuang et al. (2026); Wang et al. (2026); Godichon-Baggioni et al. (2026)—would be the natural remedy. Whether the block-gradient long-run covariance estimator survives such a compression is open.

Geometric mixing is a real restriction, and the block length has to respect it. Theorem 10 buys $b \asymp \log N$ from the geometric decay in Lemma 12. Under algebraic mixing the flat-top bias is polynomial rather than exponential and the admissible b grows accordingly; long-memory streams are outside the theory altogether. This is not merely formal: our air-quality stream has a residual lag-one autocorrelation of 0.91 and a median variance-inflation factor of 9.9, and there the block length needed for calibrated intervals is an order of magnitude larger than for the traffic stream (Table 4). The

1019 free diagnostic $r_j = |u_j^\top (\hat{S}^{\text{FT}} - \hat{S}^{\text{BM}})u_j| / (u_j^\top \hat{S}^{\text{FT}}u_j)$ detects the problem—it
1020 estimates the relative block-length bias of plain batch means—but it does
1021 not remove it. It could be turned into an automatic rule by maintaining
1022 covariance accumulators at several block scales $b, 2b, 4b, \dots$ at once (one extra
1023 $O(d^2)$ update per block per scale, all from the same block gradients) and
1024 reporting from the smallest scale at which r_j falls below a tolerance; we did
1025 not implement this, and we do not claim an automatic rule for b .

1026 $\hat{S}^{\text{FT}}$ need not be positive semidefinite.. This is intrinsic to flat-top and lugsail
1027 estimators (Singh et al., 2025, Remark 5). Marginal intervals need only
1028 positive scalars $u_j^\top \hat{S}^{\text{FT}}u_j$. The documented fallback to $\hat{S}^{\text{BM}}$ never triggered
1029 in any simulation, but it does trigger on the real streams — at a rate of 0.0%
1030 of coordinate-replicate pairs — which is what one should expect, since there
1031 the two designs are ill-conditioned and the lag-one correction is comparable
1032 to the batch-means term. We report the rate rather than claiming it away,
1033 and a user who needs the full matrix (for a joint test, say) must project it.

1034 *The curvature form is required..* Everything rests on $\nabla^2 F(\theta) = \mathbb{E}[\alpha \Psi \Psi^\top]$,
1035 which holds for generalised linear models, ridge and softmax regression,
1036 the geometric median and Gauss–Newton problems, but not for a general
1037 objective. Without it there is no rank-one recursion and no inversion-free
1038 update.

1039 *Warm-up is a real requirement, not an implementation detail..* The $1/t$
1040 schedule with a poor preconditioner leaves a slowly decaying component in
1041 the error. On the autoregressive design the effect is now mild — with a
1042 negligible warm-up, coverage of nominal 95% intervals is 0.928 rather than
1043 0.949 — because the step-length safeguard absorbs most of what a bad
1044 preconditioner does. On the regime-switching design it is not mild: the
1045 relative variance is 2.19 at $c_w = 50$ against 1.41 at $c_w = 200$ (Table 5(c')).
1046 We give a scaling rule ($c_w d$ curvature updates, with $c_w = 200$ throughout)
1047 and show insensitivity above it, but this is a tuning constant, and the
1048 asymptotic theory is silent about it. What sets its size is not the number of
1049 observations but the number of effectively distinct design points they span: on
1050 a persistent regime-switching stream, $50d$ consecutive observations cover only
1051 about thirty regime visits, and the curvature estimate is still rank-starved
1052 (Table 5(c')). The warm-up also has a cost consequence we state rather than
1053 bury: it performs $c_w d$ rank-one updates at $O(d^2)$ each, so the total time is
1054 $O(dN + c_w d^3)$ and the advertised $O(dN)$ only bites once $N \gtrsim c_w d^2$. With
1055 $c_w = 200$ and $d = 40$ that means $N \gtrsim 3 \times 10^5$, so at the $N = 2 \times 10^5$ of

1056 Table 6 the warm-up dominates the streaming pass for every $d \geq 40$ we
 1057 report, which is why that table separates the streaming count from the total.

1058 *Blocking creates a stability condition, and our fix is not fully proved..* The
 1059 contraction condition $\gamma_t \|\bar{H}^{-1} \hat{H}_t\|_{\text{op}} < 2$ is a consequence of blocking the
 1060 gradient, so it is a cost of our construction and not of the method we build
 1061 on. The step-length safeguard in (6) removes the failure it causes, and
 1062 Proposition 1 shows that on the event where the safeguard is eventually
 1063 inactive it changes no asymptotic statement. What we do *not* prove is
 1064 that this event has probability one: the implication we establish runs from
 1065 convergence to eventual inactivity, not back, because the safeguard factor
 1066 depends on the current block’s gradient and so is not predictable, which
 1067 breaks the martingale structure the proofs use. Appendix C says exactly
 1068 what would close the gap — an expanding-truncation construction (Kushner
 1069 and Yin, 2003; Andrieu et al., 2005), or a predictable variant that caps using
 1070 the previous block’s amplification at an extra $O(d^2)$ per block — and we
 1071 did neither. Instead we report the activation counts, which are between 2
 1072 and 7 per run and flat in N across every design and sample size in Tables 2
 1073 and 5. Two further caveats. The condition itself comes from the linearised
 1074 recursion (7), so it is a heuristic for non-quadratic objectives, exact only in
 1075 the linear-model case. And c_D is a tuning constant with a genuine failure
 1076 mode: Table 5(i) shows it is irrelevant over a factor of sixteen on the designs
 1077 that do not need the safeguard, but on the one that does, $c_D = 4$ is too loose
 1078 to prevent the blow-up. The asymptotic theory says nothing about it, and we
 1079 have no rule for choosing it beyond the value that worked across our designs.

1080 *The central limit theorem is proved for fixed b , and we recommend growing*
 1081 *b ..* Theorem 6 is unconditional at each fixed block length, but Theorem 10
 1082 takes $b \asymp \log N$, so the interval we recommend sits in the triangular-array
 1083 regime where our proof needs an extra hypothesis: that the finite random
 1084 constants in Lemma 14 are bounded in probability uniformly in N along
 1085 the sequence b_N . We state that hypothesis in Theorem 6 rather than hiding
 1086 it, and we do not verify it. Closing the gap means redoing Lemma 14 with
 1087 constants explicit in b ; the naive route — bounding $\mathbb{E}\|H^{-1} - P_{t-1}\|_{\text{op}}\|\varepsilon_t\|$
 1088 directly — is the one the appendix shows loses a factor $T^{1/2}$, which is why
 1089 Gordin’s decomposition is there in the first place. Two things make us believe
 1090 the composition is sound in practice rather than merely convenient: growing
 1091 b makes the block scores less dependent, not more, so the mechanism the
 1092 constants come from is working in our favour; and coverage is flat in b over
 1093 $b \in [10, 320]$ at $N = 200,000$ (Table 5(b)), which is the statement the two

1094 theorems need to compose.

1095 *Rates and what we do not prove..* We give an almost sure rate and a central
1096 limit theorem, not a Berry–Esseen bound; the first-order Markovian literature
1097 has the latter for linear stochastic approximation (Samsonov et al., 2025),
1098 and we do not match it. We attain the semiparametric efficiency bound of Li
1099 et al. (2023) rather than establishing it. Our $\tilde{O}(N^{-1/2})$ rate for $\hat{S}^{\text{FT}}$ is not
1100 comparable to the minimax rate $\Theta(N^{-(1-a)/2})$ for Hessian-free covariance
1101 estimation from an iterate path (Ni and Huo, 2026), because we use more
1102 information (block gradients and a maintained curvature estimate) and
1103 because our step exponent is $a = 1$, outside the range that bound is stated
1104 for. Within the Hessian-free class the bound is close to attained (Wei et al.,
1105 2026), so the gap we report is about the information used, not about slack in
1106 their analysis. We make no claim of optimality within our own information
1107 set: we do not know the minimax rate for estimators that may use block
1108 gradients and a running curvature estimate.

1109 *Scope of the empirical evidence..* All designs are convex M -estimation with
1110 $d \leq 160$ and $N \leq 4 \times 10^5$; there is no deep-learning experiment and the
1111 theory would not support one. The dependence in the simulated designs is
1112 stationary and geometrically decaying by construction. On the real streams
1113 we do not know the truth, which is why we run two protocols; the semi-
1114 synthetic protocol has an exact target but a synthetic (autoregressive) error
1115 process, and the fully real protocol has a real error process but a target
1116 that is itself estimated—we correct the nominal level for that exactly, but
1117 the correction assumes the segments are approximately exchangeable, which
1118 non-stationarity would violate. Neither protocol alone would be convincing;
1119 together they bracket the honest answer.

1120 *One thing we deliberately do not claim..* Gapping improves the curvature
1121 estimate at matched cost (Proposition 11, Supplementary Fig. S2), but the
1122 curvature estimate enters the limit distribution only through a rate, so the
1123 end-to-end effect on the point estimate and on coverage is second order and
1124 we could not measure it reliably. The case for gapping is that it is free,
1125 deterministic, and provably better than the randomised alternative—not that
1126 it changes the headline numbers.

1127 8. Conclusions

1128 On a dependent stream, a streaming quasi-Newton method can get the
1129 estimate right and the uncertainty wrong. The error is not small and does

not go away: the reported variance omits the autocovariances of the score, so a nominal 95% interval converges to a coverage below 95% and stays there. Proposition 8 says exactly where: at $2\Phi(z_{\alpha/2}/\sqrt{\kappa}) - 1$, a number computable from one variance-inflation factor κ . At the inflation we measure on hourly air quality this is 0.466, so more than half of the intervals a practitioner would report on that stream miss the value they are meant to cover.

The remedy is a change of reading order rather than an extra estimation stage. Consume the stream in contiguous blocks and take one preconditioned step per block. The block-averaged gradients are then, up to a factor, the non-overlapping batch means of the score process, and batch means are a standard estimator of exactly the long-run variance the interval needs. So the object the statistics requires and the object the computation already produces are the same, and the resulting interval is *cheaper* than the invalid one: an instantaneous plug-in covariance costs $O(d^2)$ per observation, while the blocked long-run covariance costs $O(d^2)$ per block. Gapping the rank-one curvature updates is the second change, and it is a cost knob rather than an inference device — it dominates the randomised thinning the independent-data method uses, at the same number of updates, and makes the per-block cost deterministic.

Three practical consequences are worth stating in plain terms. First, block length is the one parameter that needs thought, and it needs less than one might fear: because the blocks carry gradients rather than parameter values, their length is set by how quickly the *data* forget, not by how quickly the optimiser settles, so b of order $\log N$ suffices where iterate-based methods need a power of N . A diagnostic that costs nothing, r_j , indicates when b is too short. Second, blocking has a price we pay and report: it makes a step-size condition non-trivial, and without the $O(d)$ safeguard of (6) our own method diverges on a persistent regime-switching design. Anyone who blocks a preconditioned step on dependent data will meet this, and Section 2 explains why and what to do. Third, honest intervals are wider. On our designs they are 1.37–2.72 times the dependence-blind width, which is $\sqrt{\kappa}$ and not a tuning choice. The narrower interval was never correct; it was the wrong length.

What we have not done bounds the claims. The central limit theorem is proved unconditionally for each fixed block length, while the estimator we recommend lets b grow with N ; the safeguard is shown to be asymptotically invisible on the event that it stops acting, not unconditionally; memory remains $O(d^2)$, so very high-dimensional streams need a sketched curvature representation whose interaction with block-gradient variance estimation is open; and geometric mixing excludes long-memory data. Section 7 states

each of these precisely. The construction itself is general: any streaming method that forms block gradients and maintains a curvature estimate can report a long-run variance at no extra order of cost, and we would expect the same argument to carry to sketched and diagonal curvature approximations.

Declaration of competing interest

The author declares that he has no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Funding

This research did not receive any specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

CRedit authorship contribution statement

Ihor Kendiukhov: Conceptualization, Methodology, Formal analysis, Software, Investigation, Validation, Visualization, Writing – original draft, Writing – review and editing.

Data availability

All code, saved results and derived data needed to reproduce every number, figure and table in this article are archived and citable (Kendiukhov, 2026). The release used to produce the results reported here is <https://doi.org/10.5281/zenodo.21891313>; the concept identifier 10.5281/zenodo.21891312 always resolves to the most recent release, and development continues at <https://github.com/Kendiukhov/streaming-quasi-newton-dependent-data>. The archive contains the estimator, the experiment scripts, the result files from which every reported number is generated, and the tests that validate every closed-form population quantity against long simulation.

The two real data sets analysed are third-party and are not redistributed. They are downloaded from the UCI Machine Learning Repository by `experiments/fetch_data.py` and are cited in the reference list: Metro Interstate Traffic Volume (Hogue, 2019), 10.24432/C5X60B; and Beijing Multi-Site Air Quality (Chen, 2017), 10.24432/C5RK5G.

1200 **Declaration of generative AI and AI-assisted technologies in the** 1201 **manuscript preparation process**

1202 During the preparation of this work the author used a large language
1203 model (Anthropic Claude) to assist with implementing the software, running
1204 the experiments, and drafting and revising the manuscript, including its proofs.
1205 The author reviewed and edited all output, verified every reference against
1206 its published source, checked the derivations, and takes full responsibility for
1207 the content of this article.

1208 **Acknowledgements**

1209 The author thanks the maintainers of the UCI Machine Learning Reposi-
1210 tory and the depositors of the Metro Interstate Traffic Volume and Beijing
1211 Multi-Site Air Quality data sets for making these streams publicly available.

1212 **References**

- 1213 Agarwal, A., Duchi, J.C., 2013. The generalization ability of online algorithms
1214 for dependent data. *IEEE Transactions on Information Theory* 59, 573–587.
1215 doi:[10.1109/TIT.2012.2212414](https://doi.org/10.1109/TIT.2012.2212414). arXiv:1110.2529.
- 1216 Andrews, D.W.K., 1991. Heteroskedasticity and autocorrelation consistent
1217 covariance matrix estimation. *Econometrica* 59, 817–858. doi:[10.2307/](https://doi.org/10.2307/2938229)
1218 [2938229](https://doi.org/10.2307/2938229).
- 1219 Andrews, D.W.K., Monahan, J.C., 1992. An improved heteroskedasticity
1220 and autocorrelation consistent covariance matrix estimator. *Econometrica*
1221 60, 953–966. doi:[10.2307/2951574](https://doi.org/10.2307/2951574).
- 1222 Andrieu, C., Moulines, É., Priouret, P., 2005. Stability of stochastic ap-
1223 proximation under verifiable conditions. *SIAM Journal on Control and*
1224 *Optimization* 44, 283–312. doi:[10.1137/S0363012902417267](https://doi.org/10.1137/S0363012902417267).
- 1225 Benveniste, A., Métivier, M., Priouret, P., 1990. Adaptive Algorithms and
1226 Stochastic Approximations. volume 22 of *Applications of Mathematics*.
1227 Springer-Verlag, Berlin, Heidelberg. doi:[10.1007/978-3-642-75894-2](https://doi.org/10.1007/978-3-642-75894-2).
- 1228 Bercu, B., Godichon-Baggioni, A., Portier, B., 2020. An efficient stochas-
1229 tic Newton algorithm for parameter estimation in logistic regressions.
1230 *SIAM Journal on Control and Optimization* 58, 348–367. doi:[10.1137/](https://doi.org/10.1137/19M1261717)
1231 [19M1261717](https://doi.org/10.1137/19M1261717). arXiv:1904.07908.

- 1232 Borkar, V.S., 2008. Stochastic Approximation: A Dynamical Systems View-
1233 point. volume 48 of *Texts and Readings in Mathematics*. Hindustan Book
1234 Agency, Gurgaon, India. doi:[10.1007/978-93-86279-38-5](https://doi.org/10.1007/978-93-86279-38-5).
- 1235 Bottou, L., Curtis, F.E., Nocedal, J., 2018. Optimization methods for
1236 large-scale machine learning. *SIAM Review* 60, 223–311. doi:[10.1137/
1237 16M1080173](https://doi.org/10.1137/16M1080173).
- 1238 Boyer, C., Godichon-Baggioni, A., 2023. On the asymptotic rate of con-
1239 vergence of Stochastic Newton algorithms and their Weighted Averaged
1240 versions. *Computational Optimization and Applications* 84, 921–972.
1241 doi:[10.1007/s10589-022-00442-3](https://doi.org/10.1007/s10589-022-00442-3). arXiv:2011.09706. Correction: Com-
1242 put. Optim. Appl. 87(2):705–706, 2024, doi:[10.1007/s10589-023-00511-1](https://doi.org/10.1007/s10589-023-00511-1).
- 1243 Bradley, R.C., 2005. Basic properties of strong mixing conditions. A survey
1244 and some open questions. *Probability Surveys* 2, 107–144. doi:[10.1214/
1245 154957805100000104](https://doi.org/10.1214/154957805100000104).
- 1246 Byrd, R.H., Hansen, S.L., Nocedal, J., Singer, Y., 2016. A stochastic Quasi-
1247 Newton method for large-scale optimization. *SIAM Journal on Optimization*
1248 26, 1008–1031. doi:[10.1137/140954362](https://doi.org/10.1137/140954362).
- 1249 Cénac, P., Godichon-Baggioni, A., Portier, B., 2025. An efficient averaged
1250 stochastic Gauss-Newton algorithm for estimating parameters of nonlin-
1251 ear regressions models. *Bernoulli* 31, 1–29. doi:[10.3150/23-BEJ1637](https://doi.org/10.3150/23-BEJ1637).
1252 arXiv:2006.12920.
- 1253 Chau, H.N., Kirkby, J.L., Nguyen, D.H., Nguyen, D., Nguyen, N.N., Nguyen,
1254 T., 2024. On the inversion-free Newton’s method and its applications.
1255 *International Statistical Review* 92, 284–321. doi:[10.1111/insr.12563](https://doi.org/10.1111/insr.12563).
- 1256 Chen, S., 2017. Beijing Multi-Site Air Quality [dataset]. UCI Machine
1257 Learning Repository. doi:[10.24432/C5RK5G](https://doi.org/10.24432/C5RK5G). donated 19 September 2019;
1258 licensed CC BY 4.0.
- 1259 Chen, X., Lai, Z., Li, H., Zhang, Y., 2024. Online statistical inference for
1260 stochastic optimization via Kiefer-Wolfowitz methods. *Journal of the*
1261 *American Statistical Association* 119, 2972–2982. doi:[10.1080/01621459.
1262 2023.2296703](https://doi.org/10.1080/01621459.2023.2296703).
- 1263 Chen, X., Lee, J.D., Tong, X.T., Zhang, Y., 2020. Statistical inference for
1264 model parameters in stochastic gradient descent. *The Annals of Statistics*
1265 48, 251–273. doi:[10.1214/18-AOS1801](https://doi.org/10.1214/18-AOS1801).

- 1266 Damerdj, H., 1994. Strong consistency of the variance estimator in steady-
1267 state simulation output analysis. *Mathematics of Operations Research* 19,
1268 494–512. doi:[10.1287/moor.19.2.494](https://doi.org/10.1287/moor.19.2.494).
- 1269 Davydov, Y.A., 1968. Convergence of distributions generated by stationary
1270 stochastic processes. *Theory of Probability & Its Applications* 13, 691–696.
1271 doi:[10.1137/1113086](https://doi.org/10.1137/1113086).
- 1272 Dedecker, J., Doukhan, P., Lang, G., León, J.R., Louhichi, S., Prieur, C.,
1273 2007. Weak Dependence: With Examples and Applications. volume 190
1274 of *Lecture Notes in Statistics*. Springer, New York, NY. doi:[10.1007/
1275 978-0-387-69952-3](https://doi.org/10.1007/978-0-387-69952-3).
- 1276 Doan, T.T., Nguyen, L.M., Pham, N.H., Romberg, J., 2020. Convergence rates
1277 of accelerated Markov gradient descent with applications in reinforcement
1278 learning. [arXiv:2002.02873](https://arxiv.org/abs/2002.02873). arXiv:2002.02873; v3, 19 October 2020.
- 1279 Domowitz, I., White, H., 1982. Misspecified models with dependent obser-
1280 vations. *Journal of Econometrics* 20, 35–58. doi:[10.1016/0304-4076\(82\)
1281 90102-6](https://doi.org/10.1016/0304-4076(82)90102-6).
- 1282 Doukhan, P., 1994. Mixing: Properties and Examples. volume 85 of
1283 *Lecture Notes in Statistics*. Springer-Verlag, New York. doi:[10.1007/
1284 978-1-4612-2642-0](https://doi.org/10.1007/978-1-4612-2642-0).
- 1285 Duchi, J.C., Agarwal, A., Johansson, M., Jordan, M.I., 2012. Ergodic mirror
1286 descent. *SIAM Journal on Optimization* 22, 1549–1578. doi:[10.1137/
1287 110836043](https://doi.org/10.1137/110836043).
- 1288 Even, M., 2023. Stochastic gradient descent under Markovian sampling
1289 schemes, in: *Proceedings of the 40th International Conference on Machine
1290 Learning (ICML)*, PMLR. pp. 9412–9439. ArXiv:2302.14428.
- 1291 Fabian, V., 1968. On asymptotic normality in stochastic approximation.
1292 *The Annals of Mathematical Statistics* 39, 1327–1332. doi:[10.1214/aoms/
1293 1177698258](https://doi.org/10.1214/aoms/1177698258).
- 1294 Fang, Y., 2019. Scalable statistical inference for averaged implicit stochastic
1295 gradient descent. *Scandinavian Journal of Statistics* 46, 987–1002. doi:[10.
1296 1111/sjos.12378](https://doi.org/10.1111/sjos.12378).
- 1297 Fang, Y., Xu, J., Yang, L., 2018. Online bootstrap confidence intervals for
1298 the stochastic gradient descent estimator. *Journal of Machine Learning
1299 Research* 19, 1–21. URL: <http://jmlr.org/papers/v19/17-370.html>.

- 1300 Flegal, J.M., Jones, G.L., 2010. Batch means and spectral variance estimators
1301 in Markov chain Monte Carlo. *The Annals of Statistics* 38, 1034–1070.
1302 doi:[10.1214/09-AOS735](https://doi.org/10.1214/09-AOS735).
- 1303 Fort, G., 2015. Central limit theorems for stochastic approximation with
1304 controlled Markov chain dynamics. *ESAIM: Probability and Statistics* 19,
1305 60–80. doi:[10.1051/ps/2014013](https://doi.org/10.1051/ps/2014013). arXiv:1309.3116.
- 1306 Glynn, P.W., Whitt, W., 1991. Estimating the asymptotic variance with
1307 batch means. *Operations Research Letters* 10, 431–435. doi:[10.1016/
1308 0167-6377\(91\)90019-L](https://doi.org/10.1016/0167-6377(91)90019-L).
- 1309 Glynn, P.W., Whitt, W., 1992. The asymptotic validity of sequential stopping
1310 rules for stochastic simulations. *The Annals of Applied Probability* 2, 180–
1311 198. doi:[10.1214/aoap/1177005777](https://doi.org/10.1214/aoap/1177005777).
- 1312 Godichon-Baggioni, A., 2019. Online estimation of the asymptotic variance for
1313 averaged stochastic gradient algorithms. *Journal of Statistical Planning and
1314 Inference* 203, 1–19. doi:[10.1016/j.jspi.2019.01.001](https://doi.org/10.1016/j.jspi.2019.01.001). arXiv:1702.00931.
- 1315 Godichon-Baggioni, A., Lu, W., 2024. Online stochastic Newton methods for
1316 estimating the geometric median and applications. *Journal of Multivariate
1317 Analysis* 202, 105313. doi:[10.1016/j.jmva.2024.105313](https://doi.org/10.1016/j.jmva.2024.105313).
- 1318 Godichon-Baggioni, A., Portier, B., Sallé, G., 2026. Complexity reduction
1319 in online stochastic Newton methods with potential $O(Nd)$ total cost.
1320 *Statistics and Computing* 36, 194. doi:[10.1007/s11222-026-10951-4](https://doi.org/10.1007/s11222-026-10951-4).
1321 article 194; preprint arXiv:2602.05460.
- 1322 Godichon-Baggioni, A., Werge, N., 2025. On adaptive stochastic optimization
1323 for streaming data: A Newton’s method with $o(dn)$ operations. *Journal
1324 of Machine Learning Research* 26, 1–49. URL: [http://jmlr.org/papers/
1325 v26/23-1565.html](http://jmlr.org/papers/v26/23-1565.html).
- 1326 Godichon-Baggioni, A., Werge, N., Wintenberger, O., 2023. Learning from
1327 time-dependent streaming data with online stochastic algorithms. *Trans-
1328 actions on Machine Learning Research* URL: [https://openreview.net/
1329 forum?id=kdfiEu1ul6](https://openreview.net/forum?id=kdfiEu1ul6). arXiv:2205.12549.
- 1330 Hall, P., Heyde, C.C., 1980. *Martingale Limit Theory and Its Application*.
1331 Academic Press, New York.
- 1332 Hansen, L.P., 1982. Large sample properties of generalized method of moments
1333 estimators. *Econometrica* 50, 1029–1054. doi:[10.2307/1912775](https://doi.org/10.2307/1912775).

- 1334 Herrndorf, N., 1984. A functional central limit theorem for weakly dependent
1335 sequences of random variables. *The Annals of Probability* 12, 141–153.
1336 doi:[10.1214/aop/1176993379](https://doi.org/10.1214/aop/1176993379).
- 1337 Hogue, J., 2019. Metro interstate traffic volume. UCI Machine Learning
1338 Repository. doi:[10.24432/C5X60B](https://doi.org/10.24432/C5X60B). dataset; donated 6 May 2019; licensed
1339 CC BY 4.0.
- 1340 Huber, P.J., 1967. The behavior of maximum likelihood estimates under
1341 nonstandard conditions, in: *Proceedings of the Fifth Berkeley Symposium*
1342 *on Mathematical Statistics and Probability, Volume 1: Statistics*, University
1343 of California Press, Berkeley. pp. 221–233.
- 1344 Ibragimov, I.A., 1962. Some limit theorems for stationary processes. *Theory*
1345 *of Probability and Its Applications* 7, 349–382. doi:[10.1137/1107036](https://doi.org/10.1137/1107036).
1346 russian original: *Teor. Veroyatnost. i Primenen.* 7(4), 361–392.
- 1347 Jones, G.L., Haran, M., Caffo, B.S., Neath, R., 2006. Fixed-width output
1348 analysis for Markov Chain Monte Carlo. *Journal of the American Statistical*
1349 *Association* 101, 1537–1547. doi:[10.1198/016214506000000492](https://doi.org/10.1198/016214506000000492).
- 1350 Karimi, B., Miasojedow, B., Moulines, E., Wai, H.T., 2019. Non-asymptotic
1351 analysis of biased stochastic approximation scheme, in: *Beygelzimer, A.,*
1352 *Hsu, D. (Eds.), Proceedings of the Thirty-Second Conference on Learning*
1353 *Theory (COLT)*, PMLR. pp. 1944–1974. ArXiv:1902.00629.
- 1354 Kelly, M., Longjohn, R., Nottingham, K., . The UCI machine learning repos-
1355 itory. <https://archive.ics.uci.edu>. University of California, Irvine.
1356 Accessed 10 August 2026.
- 1357 Kendiukhov, I., 2026. Streaming inversion-free quasi-Newton estimation and
1358 inference on dependent data. doi:[10.5281/zenodo.21891313](https://doi.org/10.5281/zenodo.21891313). software,
1359 version release1, MIT licence. Concept DOI [10.5281/zenodo.21891312](https://doi.org/10.5281/zenodo.21891312)
1360 resolves to the latest release.
- 1361 Kiefer, N.M., Vogelsang, T.J., 2002. Heteroskedasticity-autocorrelation robust
1362 standard errors using the Bartlett kernel without truncation. *Econometrica*
1363 70, 2093–2095. doi:[10.1111/1468-0262.00366](https://doi.org/10.1111/1468-0262.00366).
- 1364 Kiefer, N.M., Vogelsang, T.J., 2005. A new asymptotic theory for
1365 heteroskedasticity-autocorrelation robust tests. *Econometric Theory* 21,
1366 1130–1164. doi:[10.1017/S0266466605050565](https://doi.org/10.1017/S0266466605050565).

- 1367 Kowshik, S., Nagaraj, D., Jain, P., Netrapalli, P., 2021. Streaming lin-
 1368 ear system identification with reverse experience replay, in: Ranzato, M.,
 1369 Beygelzimer, A., Dauphin, Y., Liang, P.S., Vaughan, J.W. (Eds.), Advances
 1370 in Neural Information Processing Systems, Curran Associates, Inc.. pp.
 1371 30140–30152. URL: [https://proceedings.neurips.cc/paper_files/
 1372 paper/2021/file/fd2c5e4680d9a01dba3aada5ece22270-Paper.pdf](https://proceedings.neurips.cc/paper_files/paper/2021/file/fd2c5e4680d9a01dba3aada5ece22270-Paper.pdf).
- 1373 Kuang, W., Anitescu, M., Na, S., 2026. Online covariance matrix estimation
 1374 in sketched Newton methods. Information and Inference: A Journal of the
 1375 IMA 15, iaag012. doi:[10.1093/imaiai/iaag012](https://doi.org/10.1093/imaiai/iaag012).
- 1376 Künsch, H.R., 1989. The jackknife and the bootstrap for general stationary
 1377 observations. The Annals of Statistics 17, 1217–1241. doi:[10.1214/aos/
 1378 1176347265](https://doi.org/10.1214/aos/1176347265).
- 1379 Kushner, H.J., Yin, G.G., 2003. Stochastic Approximation and Recursive
 1380 Algorithms and Applications. volume 35 of *Stochastic Modelling and Applied
 1381 Probability*. 2nd ed., Springer, New York. doi:[10.1007/b97441](https://doi.org/10.1007/b97441).
- 1382 Lahiri, S.N., 2003. Resampling Methods for Dependent Data. Springer Series
 1383 in Statistics, Springer, New York, NY. doi:[10.1007/978-1-4757-3803-2](https://doi.org/10.1007/978-1-4757-3803-2).
- 1384 Lee, S., Liao, Y., Seo, M.H., Shin, Y., 2022. Fast and robust online inference
 1385 with stochastic gradient descent via random scaling, in: Proceedings of the
 1386 AAAI Conference on Artificial Intelligence, pp. 7381–7389. doi:[10.1609/
 1387 aaai.v36i7.20701](https://doi.org/10.1609/aaai.v36i7.20701).
- 1388 Leluc, R., Portier, F., 2023. Asymptotic analysis of conditioned stochas-
 1389 tic gradient descent. Transactions on Machine Learning Research
 1390 ArXiv:2006.02745.
- 1391 Leung, M.F., Chan, K.W., 2026. Principles of statistical inference in
 1392 online problems. Bernoulli 32, 1122–1141. doi:[10.3150/25-BEJ1894](https://doi.org/10.3150/25-BEJ1894).
 1393 arXiv:2209.05399.
- 1394 Li, X., Liang, J., Chen, X., Zhang, Z., 2026. Convergence and inference of
 1395 stream stochastic gradient descent, with applications to queueing systems
 1396 and inventory control. Operations Research 74, 2034–2052. doi:[10.1287/
 1397 opre.2025.1662](https://doi.org/10.1287/opre.2025.1662).
- 1398 Li, X., Liang, J., Zhang, Z., 2023. Online statistical inference for nonlinear
 1399 stochastic approximation with Markovian data. ArXiv:2302.07690.

- 1400 Liang, K.Y., Zeger, S.L., 1986. Longitudinal data analysis using generalized
1401 linear models. *Biometrika* 73, 13–22. doi:[10.1093/biomet/73.1.13](https://doi.org/10.1093/biomet/73.1.13).
- 1402 Liu, R., Chen, X., Shang, Z., 2023. Statistical inference with stochastic gradi-
1403 ent methods under ϕ -mixing data. [arXiv:2302.12717](https://arxiv.org/abs/2302.12717). arXiv:2302.12717.
- 1404 Merlevède, F., Peligrad, M., Rio, E., 2011. A Bernstein type inequality and
1405 moderate deviations for weakly dependent sequences. *Probability Theory
1406 and Related Fields* 151, 435–474. doi:[10.1007/s00440-010-0304-9](https://doi.org/10.1007/s00440-010-0304-9).
- 1407 Merlevède, F., Peligrad, M., Utev, S., 2006. Recent advances in invariance
1408 principles for stationary sequences. *Probability Surveys* 3, 1–36. doi:[10.1214/154957806100000202](https://doi.org/10.1214/154957806100000202).
- 1410 Meyn, S., Tweedie, R.L., 2009. *Markov Chains and Stochastic Stability*.
1411 Cambridge Mathematical Library. 2nd ed., Cambridge University Press,
1412 Cambridge. doi:[10.1017/CB09780511626630](https://doi.org/10.1017/CB09780511626630). with a prologue by Peter W.
1413 Glynn.
- 1414 Mokkadem, A., Pelletier, M., 2011. A generalization of the averaging proce-
1415 dure: The use of two-time-scale algorithms. *SIAM Journal on Control and
1416 Optimization* 49, 1523–1543. doi:[10.1137/080718395](https://doi.org/10.1137/080718395).
- 1417 Moritz, P., Nishihara, R., Jordan, M.I., 2016. A linearly-convergent stochastic
1418 L-BFGS algorithm, in: Gretton, A., Robert, C.C. (Eds.), *Proceedings of
1419 the 19th International Conference on Artificial Intelligence and Statistics
1420 (AISTATS)*, PMLR, Cadiz, Spain. pp. 249–258.
- 1421 Nagaraj, D., Wu, X., Bresler, G., Jain, P., Netrapalli, P., 2020. Least squares
1422 regression with Markovian data: Fundamental limits and algorithms, in:
1423 Larochelle, H., Ranzato, M., Hadsell, R., Balcan, M.F., Lin, H. (Eds.),
1424 *Advances in Neural Information Processing Systems 33 (NeurIPS 2020)*,
1425 Curran Associates, Inc.. pp. 16666–16676.
- 1426 Newey, W.K., West, K.D., 1987. A simple, positive semi-definite, heteroskedas-
1427 ticity and autocorrelation consistent covariance matrix. *Econometrica* 55,
1428 703–708. doi:[10.2307/1913610](https://doi.org/10.2307/1913610).
- 1429 Newey, W.K., West, K.D., 1994. Automatic lag selection in covariance matrix
1430 estimation. *The Review of Economic Studies* 61, 631–653. doi:[10.2307/
1431 2297912](https://doi.org/10.2307/2297912).

- 1432 Ni, Y., Huo, X., 2026. Online covariance estimation in averaged SGD: Im-
 1433 proved batch-mean rates and minimax optimality via trajectory regression.
 1434 [arXiv:2604.10814](#). arXiv:2604.10814.
- 1435 Pelletier, M., 1998. On the almost sure asymptotic behaviour of stochastic
 1436 algorithms. *Stochastic Processes and their Applications* 78, 217–244. doi:[10.1016/S0304-4149\(98\)00029-5](#).
- 1437
 1438 Pelletier, M., 2000. Asymptotic almost sure efficiency of averaged stochastic
 1439 algorithms. *SIAM Journal on Control and Optimization* 39, 49–72. doi:[10.1137/S0363012998308169](#).
- 1440
 1441 Politis, D.N., 2011. Higher-order accurate, positive semidefinite estimation of
 1442 large-sample covariance and spectral density matrices. *Econometric Theory*
 1443 27, 703–744. doi:[10.1017/S0266466610000484](#).
- 1444 Politis, D.N., Romano, J.P., 1995. Bias-corrected nonparametric spectral
 1445 estimation. *Journal of Time Series Analysis* 16, 67–103. doi:[10.1111/j.1467-9892.1995.tb00223.x](#).
- 1446
 1447 Politis, D.N., Romano, J.P., Wolf, M., 1999. Subsampling. Springer Series in
 1448 Statistics, Springer, New York. doi:[10.1007/978-1-4612-1554-7](#).
- 1449 Polyak, B.T., Juditsky, A.B., 1992. Acceleration of stochastic approximation
 1450 by averaging. *SIAM Journal on Control and Optimization* 30, 838–855.
 1451 doi:[10.1137/0330046](#).
- 1452 Rio, E., 1993. Covariance inequalities for strongly mixing processes. *Annales*
 1453 *de l’Institut Henri Poincaré, Probabilités et Statistiques* 29, 587–597.
- 1454 Rio, E., 2017. Asymptotic Theory of Weakly Dependent Random Processes.
 1455 volume 80 of *Probability Theory and Stochastic Modelling*. Springer, Berlin.
 1456 doi:[10.1007/978-3-662-54323-8](#).
- 1457 Robbins, H., Monro, S., 1951. A stochastic approximation method. *The*
 1458 *Annals of Mathematical Statistics* 22, 400–407. doi:[10.1214/aoms/1177729586](#).
- 1459
 1460 Robbins, H., Siegmund, D., 1971. A convergence theorem for non negative
 1461 almost supermartingales and some applications, in: Rustagi, J.S. (Ed.),
 1462 *Optimizing Methods in Statistics*. Academic Press, pp. 233–257. doi:[10.1016/B978-0-12-604550-5.50015-8](#).
- 1463

- 1464 Rosenblatt, M., 1956. A central limit theorem and a strong mixing condition.
1465 Proceedings of the National Academy of Sciences of the United States of
1466 America 42, 43–47. doi:[10.1073/pnas.42.1.43](https://doi.org/10.1073/pnas.42.1.43).
- 1467 Roy, A., Balasubramanian, K., 2023. Online covariance estimation for
1468 stochastic gradient descent under Markovian sampling. [arXiv:2308.01481](https://arxiv.org/abs/2308.01481).
1469 arXiv:2308.01481.
- 1470 Ruppert, D., 1988. Efficient Estimations from a Slowly Convergent Robbins-
1471 Monro Process. Technical Report 781. Cornell University, Operations
1472 Research and Industrial Engineering. [https://hdl.handle.net/1813/](https://hdl.handle.net/1813/8664)
1473 8664.
- 1474 Samsonov, S., Sheshukova, M., Moulines, E., Naumov, A., 2025. Statistical
1475 inference for linear stochastic approximation with Markovian noise, in:
1476 Advances in Neural Information Processing Systems 38 (NeurIPS 2025),
1477 pp. 174565–174626. doi:[10.52202/085713-5807](https://doi.org/10.52202/085713-5807). arXiv:2505.19102.
- 1478 Sherman, J., Morrison, W.J., 1950. Adjustment of an inverse matrix corre-
1479 sponding to a change in one element of a given matrix. The Annals of
1480 Mathematical Statistics 21, 124–127. doi:[10.1214/aoms/1177729893](https://doi.org/10.1214/aoms/1177729893).
- 1481 Singh, R., Shukla, A., Vats, D., 2025. On the utility of equal batch sizes
1482 for inference in stochastic gradient descent. Journal of Machine Learning
1483 Research 26, 1–41. URL: <http://jmlr.org/papers/v26/24-0094.html>.
- 1484 Su, W.J., Zhu, Y., 2023. HiGrad: Uncertainty quantification for online
1485 learning and stochastic approximation. Journal of Machine Learning
1486 Research 24, 1–53. ArXiv:1802.04876.
- 1487 Sun, T., Sun, Y., Yin, W., 2018. On Markov chain gradient descent, in:
1488 Advances in Neural Information Processing Systems 31 (NeurIPS 2018),
1489 Curran Associates, Inc.. pp. 9918–9927. ArXiv:1809.04216.
- 1490 Sun, Y., Phillips, P.C.B., Jin, S., 2008. Optimal bandwidth selection in
1491 heteroskedasticity–autocorrelation robust testing. Econometrica 76, 175–
1492 194. doi:[10.1111/j.0012-9682.2008.00822.x](https://doi.org/10.1111/j.0012-9682.2008.00822.x).
- 1493 van der Vaart, A.W., 1998. Asymptotic Statistics. Number 3 in Cambridge
1494 Series in Statistical and Probabilistic Mathematics, Cambridge University
1495 Press, Cambridge. doi:[10.1017/CB09780511802256](https://doi.org/10.1017/CB09780511802256).

- 1496 Vats, D., Flegal, J.M., 2022. Lugsail lag windows for estimating time-
1497 average covariance matrices. *Biometrika* 109, 735–750. doi:[10.1093/
1498 biomet/asab049](https://doi.org/10.1093/biomet/asab049).
- 1499 Vats, D., Flegal, J.M., Jones, G.L., 2019. Multivariate output analysis
1500 for Markov chain Monte Carlo. *Biometrika* 106, 321–337. doi:[10.1093/
1501 biomet/asz002](https://doi.org/10.1093/biomet/asz002).
- 1502 Wang, H., Du, X., Na, S., 2026. Inference of online Newton methods with
1503 Nesterov’s accelerated sketching, in: *Proceedings of the 43rd International
1504 Conference on Machine Learning (ICML)*. ArXiv:2604.23436.
- 1505 Wang, X., Ma, S., Goldfarb, D., Liu, W., 2017. Stochastic quasi-Newton meth-
1506 ods for nonconvex stochastic optimization. *SIAM Journal on Optimization*
1507 27, 927–956. doi:[10.1137/15M1053141](https://doi.org/10.1137/15M1053141).
- 1508 Wei, Z., Zhu, W., Lyu, J., Wu, W.B., 2026. Refining covariance ma-
1509 trix estimation in stochastic gradient descent through bias reduction.
1510 [arXiv:2604.21203](https://arxiv.org/abs/2604.21203). arXiv:2604.21203.
- 1511 White, H., 1980. A heteroskedasticity-consistent covariance matrix estimator
1512 and a direct test for heteroskedasticity. *Econometrica* 48, 817–838. doi:[10.
1513 2307/1912934](https://doi.org/10.2307/1912934).
- 1514 White, H., 1982. Maximum likelihood estimation of misspecified models.
1515 *Econometrica* 50, 1–25. doi:[10.2307/1912526](https://doi.org/10.2307/1912526).
- 1516 White, H., 1984. *Asymptotic Theory for Econometricians*. Economic Theory,
1517 Econometrics, and Mathematical Economics, Academic Press.
- 1518 Wooldridge, J.M., White, H., 1988. Some invariance principles and central
1519 limit theorems for dependent heterogeneous processes. *Econometric Theory*
1520 4, 210–230. doi:[10.1017/S0266466600012032](https://doi.org/10.1017/S0266466600012032).
- 1521 Wu, W.B., 2005. Nonlinear system theory: Another look at dependence.
1522 *Proceedings of the National Academy of Sciences of the United States of
1523 America* 102, 14150–14154. doi:[10.1073/pnas.0506715102](https://doi.org/10.1073/pnas.0506715102).
- 1524 Wu, W.B., 2009. Recursive estimation of time-average variance constants.
1525 *The Annals of Applied Probability* 19, 1529–1552. doi:[10.1214/08-AAP587](https://doi.org/10.1214/08-AAP587).
- 1526 Wu, W.B., 2011. Asymptotic theory for stationary processes. *Statistics and
1527 Its Interface* 4, 207–226. doi:[10.4310/SII.2011.v4.n2.a15](https://doi.org/10.4310/SII.2011.v4.n2.a15).

- 1528 Zeger, S.L., Liang, K.Y., 1986. Longitudinal data analysis for discrete and
1529 continuous outcomes. *Biometrics* 42, 121–130. doi:[10.2307/2531248](https://doi.org/10.2307/2531248).
- 1530 Zhu, W., Chen, X., Wu, W.B., 2023. Online covariance matrix estimation in
1531 stochastic gradient descent. *Journal of the American Statistical Association*
1532 118, 393–404. doi:[10.1080/01621459.2021.1933498](https://doi.org/10.1080/01621459.2021.1933498).

1533 Appendix A. Proofs

1534 Throughout, C denotes a finite constant that may change from line to
1535 line, $\delta_t = \theta_t - \theta^*$, $\varepsilon_t = \bar{g}_t(\theta^*)$, $P_t = \bar{H}_t^{-1}$, and $\Lambda_k = \mathbb{E}[\varepsilon_0 \varepsilon_k^\top]$. We take $\gamma_t = 1/t$
1536 throughout and ignore the step-length safeguard Π_t of (6); Proposition 1,
1537 proved in Appendix C, shows that on the event where Π_t is eventually
1538 inactive the safeguarded iterates coincide with these from a finite time on,
1539 and Appendix C.1 treats the rejected step-size-shift variant $\gamma_t = (t + t_0)^{-1}$
1540 separately. We write n_t for the number of rank-one curvature updates
1541 completed by the end of block t ; with gap m and no Bernoulli thinning,
1542 $n_t = \lfloor tb/m \rfloor$, so $n_t \asymp t$.

1543 *Order of the argument.* The results are proved in the order of Remark 2:
1544 Lemma 2 (from the ridge alone, no assumption on the iterates); Theorem 4
1545 (uses Lemma 2 only); the preliminary rate $\|\delta_t\| = O(t^{-1/2} \log^c t)$ (uses Theo-
1546 rem 4 and consistency of $\bar{H}$, not its rate); Theorem 3; and then Theorems 5, 6,
1547 9 and 10. Every remainder bound below is uniform over $b \leq C \log N$, which
1548 is the only regime in which Theorem 10 uses them, and the b -dependence is
1549 displayed wherever it is not $O(1)$.

1550 Appendix A.1. Preliminaries on dependence

1551 **Lemma 12** (Geometric decay of the score autocovariances). *Under As-*
1552 *sumption 5, $\|\Gamma_k\|_{\text{op}} \leq C\rho^{|k|}$ with $\rho = c_{\text{mix}}^{\nu/(2+\nu)}$, and $\|\Lambda_k\|_{\text{op}} \leq Cb^{-1}\rho^{(|k|-1)b}$*
1553 *for $|k| \geq 1$ while $\|\Lambda_0\|_{\text{op}} \leq C/b$. Consequently $\sum_{k \in \mathbb{Z}} \|\Lambda_k\|_{\text{op}} \leq C/b$ and*
1554 *$S = \sum_k \Gamma_k$ converges absolutely.*

1555 *Proof.* Davydov’s inequality (Davydov, 1968), in the form of Rio (1993,
1556 Cor. 1.1), gives for random variables $U \in \sigma(\xi_i, i \leq 0)$, $V \in \sigma(\xi_i, i \geq k)$ with
1557 $2 + \nu$ moments, $|\text{Cov}(U, V)| \leq 8\alpha_{\text{mix}}(k)^{\nu/(2+\nu)}\|U\|_{2+\nu}\|V\|_{2+\nu}$. Applying it
1558 coordinatewise to g_0 and g_k and using Assumption 5 gives the first claim.
1559 For the second, $\Lambda_k = b^{-2} \sum_{i \in B_0} \sum_{j \in B_k} \Gamma_{j-i}$; for $k \geq 1$ every lag $j - i$
1560 is at least $(k - 1)b + 1$, so $\|\Lambda_k\|_{\text{op}} \leq b^{-2} \cdot b^2 \cdot C\rho^{(k-1)b}$, and for $k = 0$,
1561 $\|\Lambda_0\|_{\text{op}} \leq b^{-2} \sum_{|l| < b} (b - |l|)C\rho^{|l|} \leq Cb^{-1}$. Summing the geometric series
1562 gives $\sum_k \|\Lambda_k\|_{\text{op}} \leq C/b$. $\square$

1563 **Lemma 13** (Martingale–coboundary decomposition). *Let $(\zeta_t)_{t \geq 1}$ be sta-*
1564 *tionary, centred, $\mathbb{R}^q$ -valued, $\mathcal{F}_t$ -measurable with $\mathcal{F}_t = \sigma(\xi_i : i \leq t)$, and let*
1565 *$\mathbb{E}\|\zeta_1\|^{2+\nu} < \infty$ under Assumption 5. Then $\sum_{n \geq 0} \|\mathbb{E}[\zeta_n \mid \mathcal{F}_0]\|_2 < \infty$, the*
1566 *series $z_t = \sum_{n \geq 0} \mathbb{E}[\zeta_{t+n} \mid \mathcal{F}_t]$ converges in L^2 , and*

$$\zeta_t = m_{t+1} + z_t - z_{t+1}, \quad m_{t+1} = z_{t+1} - \mathbb{E}[z_{t+1} \mid \mathcal{F}_t], \quad (\text{A.1})$$

1567 *where (m_t) is a stationary martingale difference sequence with respect to $(\mathcal{F}_t)$*
1568 *and $\mathbb{E}\|m_1\|^2 < \infty$, $\mathbb{E}\|z_1\|^2 < \infty$. Moreover $\mathbb{E}[m_1 m_1^\top] = b^{-1}S$ when $\zeta_t = \varepsilon_t$.*

1569 *Proof.* For U that is $\sigma(\xi_i : i \geq n)$ -measurable with $2 + \nu$ moments, the stan-
1570 dard consequence of strong mixing (see Rio 2017, Ch. 1) is $\|\mathbb{E}[U \mid \mathcal{F}_0]\|_2 \leq$
1571 $C \alpha_{\text{mix}}(n)^{\nu/(2(2+\nu))} \|U\|_{2+\nu}$, which under Assumption 5 decays geometrically
1572 in n and is therefore summable; this is Gordin’s condition, and (A.1) to-
1573 gether with the stated moment bounds is Gordin’s decomposition; see Hall
1574 and Heyde (1980, Ch. 5). Summing (A.1) telescopes the coboundary, so
1575 $\sum_{t \leq T} \zeta_t = \sum_{t \leq T} m_{t+1} + z_1 - z_{T+1}$; dividing by $\sqrt{T}$ and using $z_{T+1}/\sqrt{T} \rightarrow 0$
1576 in probability identifies $\lim T^{-1} \text{Var}(\sum_{t \leq T} \zeta_t) = \mathbb{E}[m_1 m_1^\top]$, which for $\zeta_t = \varepsilon_t$
1577 is $b^{-1}S$ by (1). $\square$

1578 **Lemma 14** (Weighted sums with adapted weights: convergence and rate).
1579 *Let (ζ_t) be as in Lemma 13, let (A_t) be $\mathcal{F}_{t-1}$ -measurable random matrices, and*
1580 *let (γ_t) be deterministic and non-increasing. Suppose there are deterministic*
1581 *(c_t) , (u_t) and a finite random Λ with, almost surely for all t ,*

$$\|A_t\|_{\text{op}} \leq \Lambda c_t \quad \text{and} \quad \|\gamma_{t+1} A_{t+1} - \gamma_t A_t\|_{\text{op}} \leq \Lambda u_t, \quad (\text{A.2})$$

1582 *and write $V_T = \sum_{t \leq T} \gamma_t^2 c_t^2$ and $U_T = \sum_{t \leq T} u_t$. Then for every $\varpi > 0$, almost*
1583 *surely,*

$$\left\| \sum_{t \leq T} \gamma_t A_t \zeta_t \right\| = O(V_T^{1/2} \log^{1/2+\varpi} T) + O(U_T \log^{1+\varpi} T) + O(\gamma_T c_T T^{1/2+\varpi}). \quad (\text{A.3})$$

1584 *If moreover $V_\infty < \infty$, $U_\infty < \infty$ and $\gamma_T c_T T^{1/2+\varpi} \rightarrow 0$ for some $\varpi > 0$, then*
1585 *$\sum_t \gamma_t A_t \zeta_t$ converges almost surely.*

1586 *Proof. Localisation.* For $K \in \mathbb{N}$ replace A_t by $A_t \mathbf{1}\{\sup_{s \leq t} \max(\|A_s\|_{\text{op}}/c_s, \|\gamma_{s+1} A_{s+1} -$
1587 $\gamma_s A_s\|_{\text{op}}/u_s) \leq K\}$, which is still $\mathcal{F}_{t-1}$ -measurable (the indicator is decreasing
1588 in t and $\mathcal{F}_{t-1}$ -measurable at each t because it involves A_s only for $s \leq t$,
1589 and A_{s+1} enters only through $s+1 \leq t$). On $\{\Lambda \leq K\}$ the two processes
1590 agree and $\mathbb{P}(\bigcup_K \{\Lambda \leq K\}) = 1$, so it suffices to prove the claims under the

1591 *deterministic* envelopes $\|A_t\|_{\text{op}} \leq Kc_t$, $\|\gamma_{t+1}A_{t+1} - \gamma_t A_t\|_{\text{op}} \leq Ku_t$. Fix K
 1592 and drop it.

1593 *Martingale part.* Insert (A.1): $M_T = \sum_{t \leq T} \gamma_t A_t m_{t+1}$ is a martingale,
 1594 because A_t is $\mathcal{F}_t$ -measurable and $\mathbb{E}[m_{t+1} \mid \mathcal{F}_t] = 0$ — this is the exact
 1595 cancellation the decomposition buys, and it holds whatever A_t does. Its
 1596 predictable quadratic variation obeys

$$\langle M \rangle_T = \sum_{t \leq T} \gamma_t^2 \mathbb{E}[\|A_t m_{t+1}\|^2 \mid \mathcal{F}_t] \leq K^2 \sum_{t \leq T} \gamma_t^2 c_t^2 v_t, \quad v_t := \mathbb{E}[\|m_{t+1}\|^2 \mid \mathcal{F}_t],$$

1597 with (v_t) stationary, non-negative and $\mathbb{E}v_t = \mathbb{E}\|m_1\|^2 < \infty$. By Tonelli
 1598 $\mathbb{E}[\sum_{t \leq T} \gamma_t^2 c_t^2 v_t] = \mathbb{E}\|m_1\|^2 V_T$, so Markov's inequality along $T = 2^j$ and
 1599 Borel–Cantelli give $\langle M \rangle_T = O(V_T \log^\varpi T)$ a.s. (and $\langle M \rangle_\infty < \infty$ a.s. if
 1600 $V_\infty < \infty$). The martingale law of the iterated logarithm then gives $\|M_T\| =$
 1601 $O(V_T^{1/2} \log^{1/2+\varpi} T)$ a.s., and if $\langle M \rangle_\infty < \infty$ a locally square-integrable mar-
 1602 tingale converges a.s.

1603 *Coboundary part.* Abel summation gives

$$\sum_{t \leq T} \gamma_t A_t (z_t - z_{t+1}) = \gamma_1 A_1 z_1 - \gamma_T A_T z_{T+1} + \sum_{t < T} (\gamma_{t+1} A_{t+1} - \gamma_t A_t) z_{t+1}.$$

1604 The last sum is at most $K \sum_{t < T} u_t \|z_{t+1}\|$, a non-negative sum with expecta-
 1605 tion $K\mathbb{E}\|z_1\|U_T$, hence $O(U_T \log^{1+\varpi} T)$ a.s. by Markov and Borel–Cantelli
 1606 along dyadic blocks; if $U_\infty < \infty$ its expectation is finite and it converges
 1607 a.s. Finally $\|\gamma_T A_T z_{T+1}\| \leq K\gamma_T c_T \|z_{T+1}\| = O(\gamma_T c_T T^{1/2+\varpi})$ a.s., because
 1608 $\mathbb{E}\|z_1\|^2 < \infty$ and Borel–Cantelli give $\|z_T\| = o(T^{1/2+\varpi})$. $\square$

1609 **Remark 5** (Why the lemma is stated this way). Three features are load-
 1610 bearing and each corresponds to a place where a naive argument fails. (i)
 1611 The weights are *random and adapted*: they contain $\tilde{H}_{t-1}^{-1}$, so a lemma for
 1612 deterministic weights would not apply. Note that adaptedness is essential and
 1613 not cosmetic: for bounded adapted weights and a mixing, centred, stationary
 1614 ζ the sum $\sum_t a_t \zeta_t$ can grow like T if a_t is allowed to correlate with ζ_t —
 1615 take $\zeta_t = u_t + u_{t-1}$ with u_t i.i.d. signs and $a_t = \zeta_{t-1}$ — which is exactly the
 1616 non-martingale obstruction. The decomposition (A.1) removes it because
 1617 $\mathbb{E}[m_{t+1} \mid \mathcal{F}_t] = 0$ *exactly*, whatever A_t does. (ii) The envelope is deterministic
 1618 only up to a finite random factor, which is all that is available for weights built
 1619 from the iterates; the localisation handles it. (iii) We use a martingale law of
 1620 the iterated logarithm rather than a Bernstein inequality for mixing sums.
 1621 The Bernstein inequality of Merlevède et al. (2011) is stated for uniformly
 1622 *bounded* summands, which Assumption 5 does not provide; after (A.1) the

summands form a martingale difference sequence and no such inequality is needed. (iv) The three terms of (A.3) are needed separately: the applications below use $\gamma_t \equiv 1$, for which U_T can diverge logarithmically, so a hypothesis $U_\infty < \infty$ would be too strong.

Lemmas 13 and 14 are the technical substitute for the martingale argument used in the independent-data analysis. Two features matter. First, the weights may be *random and adapted*, which is what we need because they involve $\bar{H}_{t-1}^{-1}$. Second, the conditions (A.2) are of the form already familiar from the independent-data theory. The vanishing ridge supplies the deterministic envelope: it forces $\lambda_{\min}(A_n) \geq c'n^{1-q_r}$, hence $\|\bar{H}_t^{-1}\|_{\text{op}} = W_{n_t}\|A_{n_t}^{-1}\|_{\text{op}} \leq c^{-1}n_t^{q_r} =: c_t$ with $q_r < 1/2$, so with $\gamma_t = 1/t$ the first half of (A.2) is $\sum_t O(t^{2q_r-2}) < \infty$. The second half holds because a single rank-one Riccati update changes the preconditioner by only $O(1/t)$:

Lemma 15 (Increments of the preconditioner). (a) From the ridge alone. *Under Assumption 3 and the ridge of (4) — so using only Lemma 2, with no assumption on the iterates — $\|\bar{H}_t^{-1} - \bar{H}_{t-1}^{-1}\|_{\text{op}} = O(t^{2q_r-1})$ a.s.; hence with $\gamma_t = 1/t$ and $A_t = \bar{H}_{t-1}^{-1}$, $\|\gamma_{t+1}A_{t+1} - \gamma_t A_t\|_{\text{op}} = O(t^{2q_r-2})$ a.s., which is summable for $q_r < 1/2$, so the second half of (A.2) holds with $U_\infty < \infty$.* (b) Sharper, once the curvature has converged. *Under the assumptions of Theorem 3, $\|\bar{H}_t^{-1} - \bar{H}_{t-1}^{-1}\|_{\text{op}} = O(1/t)$ a.s., so with $\gamma_t \equiv 1$ and $A_t = \bar{H}_{t-1}^{-1}$ the increments are $u_t = O(1/t)$ and $U_T = O(\log T)$.*

Proof. $\bar{H}_t^{-1} = W_{n_t}A_{n_t}^{-1}$, and from block $t-1$ to block t the weight W increases by $O(1)$ while A^{-1} changes by $O(1)$ rank-one Sherman–Morrison steps $A^{-1} \mapsto A^{-1} - cA^{-1}uu^\top A^{-1}/(1 + cu^\top A^{-1}u)$. Writing $v = A^{-1}u$, the operator norm of one such step is $c\|v\|^2/(1 + cu^\top v)$, and since $u^\top v = v^\top Av \geq \lambda_{\min}(A)\|v\|^2$ we get the two bounds

$$\|\Delta A^{-1}\|_{\text{op}} \leq \min \left\{ \frac{1}{\lambda_{\min}(A)}, \frac{c\|u\|^2}{\lambda_{\min}(A)^2} \right\}. \quad (\text{A.4})$$

For (a), Lemma 2 gives $\lambda_{\min}(A_n) \geq c'n^{1-q_r}$, so the second bound in (A.4) is $O(n^{2q_r-2})$; multiplying by $W_{n_t} = O(t)$ and using $n_t \asymp t$ gives $\|\bar{H}_t^{-1} - \bar{H}_{t-1}^{-1}\|_{\text{op}} = O(t^{2q_r-1})$, and the change in W contributes $O(1) \cdot \|A^{-1}\|_{\text{op}} = O(t^{q_r-1})$, which is smaller. The stated bound on $\gamma_{t+1}A_{t+1} - \gamma_t A_t$ then follows from $\gamma_t = 1/t$ together with $\gamma_{t+1} - \gamma_t = O(t^{-2})$ and $\|A_t\|_{\text{op}} = O(t^{q_r})$. For (b), Theorem 3 gives $A_{n_t}/n_t \rightarrow H \succ 0$, so $\lambda_{\min}(A_{n_t}) \geq c't$ a.s. for large t , the second bound in (A.4) is $O(t^{-2})$, and multiplying by $W_{n_t} = O(t)$ gives $O(1/t)$. $\square$

1657 *Appendix A.2. The ridge bound (Lemma 2)*

1658 Each curvature update adds $\iota_n e_{k_n} e_{k_n}^\top \succeq 0$ along a cycling basis direction,
 1659 and the data term $\alpha_{i_n} \Psi_{i_n} \Psi_{i_n}^\top$ is positive semidefinite, so it can only help. After
 1660 n updates every basis direction has received at least $\lfloor n/d \rfloor$ ridge increments,
 1661 and dropping the data terms,

$$A_n \succeq \lambda_0 I + \left(\frac{1}{d} \sum_{k \leq n} \iota_k \right) I. \quad (\text{A.5})$$

1662 Now the point of freezing the ridge scale. Because $\iota_k = c_\iota \varsigma_\star k^{-q_r}$ with $\varsigma_\star$ a
 1663 *constant* — measurable with respect to the first curvature-bearing block and,
 1664 by Assumption 3, almost surely finite and strictly positive whenever that
 1665 block’s curvature is not identically zero — we may sum:

$$\sum_{k \leq n} \iota_k = c_\iota \varsigma_\star \sum_{k \leq n} k^{-q_r} \geq \frac{c_\iota \varsigma_\star}{1 - q_r} (n^{1-q_r} - 1),$$

1666 using $\sum_{k \leq n} k^{-q_r} \geq \int_1^{n+1} x^{-q_r} dx$ and $q_r < 1$. Hence $\lambda_{\min}(A_n) \geq c' n^{1-q_r}$ for
 1667 all $n \geq 2$ with $c' = c_\iota \varsigma_\star / \{2d(1 - q_r)\}$, and since $W_{n_t} = 1 + n_t$,

$$\|\bar{H}_t^{-1}\|_{\text{op}} = W_{n_t} \|A_{n_t}^{-1}\|_{\text{op}} \leq \frac{1 + n_t}{c' n_t^{1-q_r}} = O(n_t^{q_r}), \quad \lambda_{\min}(\bar{H}_t) \geq \frac{c' n_t^{1-q_r}}{1 + n_t} \geq c_1 n_t^{-q_r}$$

1668 with $c_1 = c'/2$ for $n_t \geq 1$. Nothing about the iterates is used, and no lower
 1669 bound on the accumulated data terms is needed — which is exactly why the
 1670 lemma may sit at the bottom of the ladder in Remark 2.

1671 **Remark 6** (Why the scale must be frozen). Had we scaled the ridge by
 1672 the *running* harmonic mean $d/\text{tr}(\bar{H}_k^{-1})$ — the natural choice, since it is
 1673 available for free from the maintained inverse — this proof would fail, and
 1674 not for a technical reason. That scale is bounded above by $d\lambda_{\min}(\bar{H}_k)$, so a
 1675 path along which the accumulated data terms contribute negligibly in some
 1676 direction would have its ridge shrink in step with the very eigenvalue the
 1677 ridge is supposed to hold up. Concretely, the only unconditional bound
 1678 then available is $\varsigma_k \geq \lambda_0/W_k = \Theta(1/k)$, giving $\sum_{k \leq n} \iota_k = O(1)$ instead of
 1679 $\Omega(n^{1-q_r})$, and a Gronwall argument in the opposite direction shows this is
 1680 tight: $\lambda_{\min}(A_n)$ can stay $O(1)$, i.e. $\|\bar{H}_n^{-1}\|_{\text{op}} = \Omega(n)$, which would break the
 1681 square-summability $\sum_t \gamma_t^2 \|P_{t-1}\|_{\text{op}}^2 < \infty$ that Theorem 4 needs. The bound
 1682 does hold once $\bar{H}_t \rightarrow H \succ 0$, but that is Theorem 3, three rungs higher, so
 1683 the argument would be circular. Freezing the scale removes the problem at
 1684 no cost: $\varsigma_\star$ serves only to make c_ι dimensionless, and Table 5(d) shows the
 1685 results are unchanged if the scale is dropped altogether.

1686 *Appendix A.3. Consistency (Theorem 4)*

1687 Substituting $\bar{g}_t(\theta_{t-1}) = \nabla F(\theta_{t-1}) + \zeta_t$ into (6), where $\zeta_t = \bar{g}_t(\theta_{t-1}) -$
 1688 $\nabla F(\theta_{t-1})$, and expanding $F(\theta_t) - F(\theta^*)$ to second order using Assumption 1,

$$F(\theta_t) - F(\theta^*) \leq F(\theta_{t-1}) - F(\theta^*) - \gamma_t \nabla F(\theta_{t-1})^\top P_{t-1} \nabla F(\theta_{t-1}) \\ - \gamma_t \nabla F(\theta_{t-1})^\top P_{t-1} \zeta_t + \frac{L}{2} \gamma_t^2 \|P_{t-1}\|_{\text{op}}^2 \|\bar{g}_t(\theta_{t-1})\|^2. \quad (\text{A.6})$$

1689 The last term is summable: $\gamma_t^2 \|P_{t-1}\|_{\text{op}}^2 = O(t^{2q_r-2})$ by Lemma 2 with
 1690 $q_r < 1/2$, and $\mathbb{E}\|\bar{g}_t(\theta_{t-1})\|^2 \leq C + C'(F(\theta_{t-1}) - F(\theta^*))$ by Assumption 2 and
 1691 Jensen. The cross term is the one that would be handled by a martingale
 1692 argument under independence. Instead, decompose $\zeta_t = \varepsilon_t + r_t$ with $r_t =$
 1693 $\bar{g}_t(\theta_{t-1}) - \nabla F(\theta_{t-1}) - \varepsilon_t$, so that

$$\|r_t\| \leq \|\hat{H}_t - H\|_{\text{op}} \|\delta_{t-1}\| + L_\delta \|\delta_{t-1}\|^2 \quad (\text{A.7})$$

1694 where $\hat{H}_t$ is the block curvature; r_t is dominated by $\|\delta_{t-1}\|$ times an L^2 -
 1695 bounded factor, and its contribution to (A.6) is absorbed by the Robbins–
 1696 Siegmund argument (Robbins and Siegmund, 1971). The remaining term
 1697 $-\sum_t \gamma_t \nabla F(\theta_{t-1})^\top P_{t-1} \varepsilon_t$ is handled by Lemma 14 with $A_t = \nabla F(\theta_{t-1})^\top P_{t-1}$,
 1698 which is $\mathcal{F}_{t-1}$ -measurable as required. Its first condition holds because
 1699 $\|P_{t-1}\|_{\text{op}} = O(t^{q_r})$ with $q_r < 1/2$ (Lemma 2) and $\mathbb{E}\|\nabla F(\theta_{t-1})\|^2$ is bounded
 1700 on the events $\{\sup_{s \leq t} \|\theta_s\| \leq K\}$, over which we localise and then let $K \rightarrow \infty$;
 1701 its second condition holds by Lemma 15(a) — which, being a consequence
 1702 of Lemma 2 alone, keeps this rung free of Theorem 3 — together with
 1703 $\|\nabla F(\theta_t) - \nabla F(\theta_{t-1})\|_{\text{op}} \leq L \|\theta_t - \theta_{t-1}\| = O(\gamma_t t^{q_r})$. Robbins–Siegmund
 1704 then gives $F(\theta_t) \rightarrow F(\theta^*)$ a.s. and $\sum_t \gamma_t \nabla F(\theta_{t-1})^\top P_{t-1} \nabla F(\theta_{t-1}) < \infty$ a.s.;
 1705 and since $\gamma_t \lambda_{\min}(P_{t-1}) \geq \gamma_t / \|\bar{H}_{t-1}\|_{\text{op}} \geq c t^{-1}$ (because $\|\bar{H}_t\|_{\text{op}}$ is bounded
 1706 above a.s., $\bar{H}_t$ being an average of terms with $\eta' > 2$ moments) we have
 1707 $\sum_t \gamma_t \lambda_{\min}(P_{t-1}) = \infty$; Assumption 4 then forces $\theta_t \rightarrow \theta^*$ a.s. $\square$

1708 *Appendix A.4. The preliminary rate (rung (iii) of Remark 2)*

1709 This rung is what Theorem 3 needs and Theorem 4 does not give: a rate,
 1710 not just convergence. We prove it using only Lemma 2 and Theorem 4.

1711 **Lemma 16** (Preliminary rate). *Under Assumptions 1 to 5, with $q_r \in (0, 1/4)$,*
 1712 *$\|\theta_T - \theta^*\| = O(T^{-1/2} \log^{1+\varpi} T)$ almost surely, for every $\varpi > 0$.*

1713 *Proof.* Abel summation of (6) exactly as in (A.10) — an algebraic identity
 1714 that uses only $\gamma_t = 1/t$ — gives (A.11): $\delta_T = -T^{-1} H^{-1} \sum_{t \leq T} \varepsilon_t + (\text{A}) +$
 1715 $(\text{B}) - (\text{C})$. We bound the three remainders using only what rungs (i) and (ii)
 1716 provide, namely $\|P_{t-1}\|_{\text{op}} = O(t^{q_r})$ a.s. (Lemma 2), $\theta_t \rightarrow \theta^*$ a.s. (Theorem 4),

and hence, by the local Lipschitz property in Assumption 3 and the continuity in Assumption 2, $\|P_{t-1} - H^{-1}\|_{\text{op}} = o(1)$ and $\|I - P_{t-1}H\|_{\text{op}} = o(1)$ a.s. Write $\rho_t = \|I - P_{t-1}H\|_{\text{op}}$, so $\rho_t \rightarrow 0$ a.s.

Leading term. $T^{-1} \sum_{t \leq T} \varepsilon_t = O(T^{-1/2} \log^{1/2+\varpi} T)$ a.s. by the martingale law of the iterated logarithm applied to (A.1) (the coboundary telescopes to $O(T^{-1})$ after division by T , using $\mathbb{E}\|z_1\|^2 < \infty$ and Borel–Cantelli along $\|z_t\| = o(t^{1/2})$ a.s.).

Term (A). Apply (A.3) with $\gamma_t \equiv 1$ and $A_t = H^{-1} - P_{t-1}$: the envelope may be taken $c_t \equiv 1$ (a finite random constant, since $\|P_{t-1}\|_{\text{op}} = O(t^{q_r})$ makes $\sup_t \|A_t\|_{\text{op}}/c_t$ finite only after we note $A_t \rightarrow 0$; formally take $c_t = 1$ and $\Lambda = \sup_t \|A_t\|_{\text{op}} < \infty$ a.s., which holds because $A_t \rightarrow 0$ and each A_t is finite). By Lemma 15(a) the increments are $u_t = O(t^{2q_r-1})$, so $V_T = O(T)$ and $U_T = O(T^{2q_r})$, and (A.3) gives $\|\sum_{t \leq T} A_t \varepsilon_t\| = O(T^{1/2} \log^{1/2+\varpi} T) + O(T^{2q_r} \log^{1+\varpi} T) + O(T^{1/2+\varpi})$. Dividing by T and using $q_r < 1/4$, which is exactly where that restriction is needed, (A) = $O(T^{-1/2} \log^{1/2+\varpi} T) + O(T^{2q_r-1} \log^{1+\varpi} T) = O(T^{-1/2} \log^{1/2+\varpi} T)$ a.s.

Terms (B) and (C): the bootstrap. (B) = $T^{-1} \sum_{t \leq T} (I - P_{t-1}H) \delta_{t-1}$ so $\|(B)\| \leq (\sup_{t \leq T} \rho_t) \bar{\delta}_T$ with $\bar{\delta}_T = T^{-1} \sum_{t \leq T} \|\delta_{t-1}\|$, and by (A.7) $\|(C)\| \leq CT^{-1} \sum_{t \leq T} \|P_{t-1}\|_{\text{op}} (\|\hat{H}_t - H\|_{\text{op}} \|\delta_{t-1}\| + \|\delta_{t-1}\|^2)$, which is $o(1) \bar{\delta}_T$ a.s. because $\|P_{t-1}\|_{\text{op}} \|\hat{H}_t - H\|_{\text{op}}$ and $\|P_{t-1}\|_{\text{op}} \|\delta_{t-1}\|$ are both $o(1)$ a.s. ($\delta_t \rightarrow 0$ by Theorem 4, and $\|P_{t-1}\|_{\text{op}} \|\hat{H}_t - H\|_{\text{op}} \rightarrow 0$ because $\|P_{t-1}\|_{\text{op}} \rightarrow \|H^{-1}\|_{\text{op}}$ and $\hat{H}_t \rightarrow H$ in the Cesàro sense along the block index). Hence there is a random $T_1 < \infty$ and a sequence $\epsilon_T \downarrow 0$ with

$$D_T := \max_{t \leq T} \|\delta_t\| \leq \epsilon_T D_T + R_T, \quad R_T = O(T^{-1/2} \log^{1/2+\varpi} T) \text{ a.s.}, \quad (\text{A.8})$$

for $T \geq T_1$, where we used $\bar{\delta}_T \leq D_T$ and that the same bound applies to every $t \leq T$ (the right-hand side is non-decreasing in T). For T large enough that $\epsilon_T \leq 1/2$, (A.8) rearranges to $D_T \leq 2R_T$, which is the claim with $\log^{1+\varpi}$ absorbing $\log^{1/2+\varpi}$. $\square$

Two things are worth flagging about this rung. It is the only place the restriction $q_r < 1/4$ (rather than $q_r < 1/2$) is used, and it is used because at this rung the sharper increment bound Lemma 15(b) is not yet available. And the bootstrap in (A.8) is why the rate carries a log factor rather than being clean $T^{-1/2}$; Theorem 5 recovers the sharper statement once Theorem 3 is in hand.

1750 *Appendix A.5. The curvature rate (Theorem 3)*

1751 Write $Y_n = \alpha(\theta^*; \xi_{i_n})\Psi(\theta^*; \xi_{i_n})\Psi(\theta^*; \xi_{i_n})^\top$ and $\tilde{Y}_n = \alpha(\theta_{t(n)-1}; \xi_{i_n})\Psi\Psi^\top$
 1752 for the term actually accumulated, where $t(n)$ is the block containing i_n .
 1753 Then, by (4) and (5),

$$\bar{H}_t - H = \underbrace{\frac{\lambda_0 \varsigma_0 I}{W_{n_t}}}_{(i)} + \underbrace{\frac{1}{W_{n_t}} \sum_{n \leq n_t} \iota_n e_{k_n} e_{k_n}^\top}_{(ii)} + \underbrace{\frac{1}{W_{n_t}} \sum_{n \leq n_t} (Y_n - H)}_{(iii)} + \underbrace{\frac{1}{W_{n_t}} \sum_{n \leq n_t} (\tilde{Y}_n - Y_n)}_{(iv)}. \quad (\text{A.9})$$

1754 Since $W_{n_t} = 1 + n_t \asymp t$, term (i) is $O(1/t)$. For (ii), $\|\cdot\|_{\text{op}} \leq n_t^{-1} \sum_{n \leq n_t} \iota_n \leq$
 1755 $C n_t^{-q_r}$ because $\iota_n \leq C n^{-q_r}$ and $q_r < 1$; the same bound from below along each
 1756 basis direction gives, together with $H \succ 0$, $\lambda_{\min}(\bar{H}_t) \geq c n_t^{-q_r}$ for all t (this is
 1757 the only place the ridge is used, and it is what bounds $\|\bar{H}_t^{-1}\|_{\text{op}}$ for *every* t
 1758 rather than eventually). Term (iv) is bounded, by the local Lipschitz prop-
 1759 erty in Assumption 3 and Cauchy–Schwarz, by $C n_t^{-1} \sum_{n \leq n_t} \|\theta_{t(n)-1} - \theta^*\| \cdot$
 1760 $\|\Psi\Psi^\top\|_{\text{op}}$, which by Lemma 16 is $O(n_t^{-1} \sum_{n \leq n_t} n^{-1/2} \log^{1+\varpi} n) = O(n_t^{-1/2} \log^{1+\varpi} n_t)$
 1761 a.s.

1762 The new content is term (iii). The array $(Y_n)_{n \geq 1}$ is the subsequence of a
 1763 stationary geometrically mixing sequence at spacing m , hence itself stationary
 1764 with mixing coefficients $\alpha_{\text{mix}}(jm) \leq K c_{\text{mix}}^{jm}$, and has $\eta' > 2$ moments by
 1765 Assumption 3. Applying Lemma 14 with $a_n \equiv 1$ (i.e. its maximal-inequality
 1766 half, which does not need square-summability) gives $\|\sum_{n \leq n_t} (Y_n - H)\|_{\text{op}} =$
 1767 $O(n_t^{1/2} \log^{1/2+\varpi} n_t)$ a.s., so (iii) = $O(t^{-1/2} \log^{1/2+\varpi} t)$ a.s. Collecting the four
 1768 terms, $\|\bar{H}_t - H\|_{\text{op}} = O(t^{-\min\{q_r, 1/2\}} \log^{c'} t)$ a.s., which is $O(t^{-\nu_H})$ for every
 1769 $\nu_H < \min\{q_r, 1/2\}$. Finally, on the event $\|\bar{H}_t - H\|_{\text{op}} \leq \lambda_{\min}(H)/2$ (which
 1770 holds for all large t a.s.), $\|\bar{H}_t^{-1} - H^{-1}\|_{\text{op}} \leq \|\bar{H}_t^{-1}\|_{\text{op}} \|\bar{H}_t - H\|_{\text{op}} \|H^{-1}\|_{\text{op}} \leq$
 1771 $2\lambda_{\min}(H)^{-2} \|\bar{H}_t - H\|_{\text{op}}$. $\square$

1772 *Appendix A.6. Proof of Theorem 5 (rate) and Proposition 7 and Theorem 6*
 1773 *(CLT)*

1774 Write (6) as $\delta_t = \delta_{t-1} - \gamma_t P_{t-1} \{H \delta_{t-1} + \varepsilon_t + r_t\}$ with r_t as above. Multi-
 1775 plying by t and using $\gamma_t = 1/t$,

$$t \delta_t = (t-1) \delta_{t-1} + (I - P_{t-1} H) \delta_{t-1} - P_{t-1} \varepsilon_t - P_{t-1} r_t, \quad (\text{A.10})$$

1776 so that, summing from $t = 1$ to T (Abel summation) and dividing by T ,

$$\delta_T = -\frac{1}{T} H^{-1} \sum_{t \leq T} \varepsilon_t + \underbrace{\frac{1}{T} \sum_{t \leq T} (H^{-1} - P_{t-1}) \varepsilon_t}_{(A)} + \underbrace{\frac{1}{T} \sum_{t \leq T} (I - P_{t-1} H) \delta_{t-1}}_{(B)} - \underbrace{\frac{1}{T} \sum_{t \leq T} P_{t-1} r_t}_{(C)}. \quad (\text{A.11})$$

1777 The leading term is *exactly* $-N_T^{-1}H^{-1}\sum_{i\leq N_T}g_i$, because $\sum_{t\leq T}\varepsilon_t = b^{-1}\sum_{i\leq N_T}g_i$
 1778 and $N_T = bT$. This identity is Proposition 7: the leading term does not
 1779 depend on b .

1780 *Term (A).* Apply (A.3) with $\gamma_t \equiv 1$, $\zeta_t = \varepsilon_t$ and the $\mathcal{F}_{t-1}$ -measurable
 1781 weights $A_t = H^{-1} - P_{t-1}$: by Theorem 3 the envelope is $c_t = t^{-\nu_H}$, and by
 1782 Lemma 15(b) the increments are $u_t = O(t^{-1})$. Hence $V_T = O(T^{1-2\nu_H})$ (as
 1783 $\nu_H < 1/2$), $U_T = O(\log T)$, and $\gamma_T c_T T^{1/2+\varpi} = O(T^{1/2-\nu_H+\varpi})$, so

$$\left\| \sum_{t\leq T} (H^{-1} - P_{t-1})\varepsilon_t \right\| = O(T^{1/2-\nu_H} \log^{1/2+\varpi} T) + O(\log^{2+\varpi} T) + O(T^{1/2-\nu_H+\varpi}) = O(T^{1/2-\nu_H+\varpi}) \text{ a.s.}$$

1784 Choosing $\varpi < \nu_H$ gives (A) = $O(T^{-1/2-\nu_H+\varpi}) = o(T^{-1/2})$ a.s. It is worth
 1785 seeing what does the work. The only bound available without (A.1) is
 1786 $\|\mathbb{E}[(H^{-1} - P_{t-1})\varepsilon_t]\| \leq \|H^{-1} - P_{t-1}\|_{\text{op}} \|\varepsilon_t\|_2 = O(t^{-\nu_H} b^{-1/2})$ — Davydov
 1787 gives no decay across a one-block gap, since the weight and the summand are
 1788 only one block apart — and averaging that over $t \leq T$ leaves $O(b^{-1/2} T^{-\nu_H})$,
 1789 which is *not* $o(T^{-1/2})$ for any admissible $\nu_H < 1/2$. The extra $T^{-1/2}$ comes
 1790 entirely from the exact cancellation $\mathbb{E}[A_t m_{t+1} \mid \mathcal{F}_t] = 0$.

1791 *Term (B).* $\|(B)\| \leq T^{-1} \sum_{t\leq T} \|I - P_{t-1}H\|_{\text{op}} \|\delta_{t-1}\| = T^{-1} \sum_{t\leq T} O(t^{-\nu_H}) \|\delta_{t-1}\|$,
 1792 which is a pathwise bound needing no lemma. With the rung-(iii) preliminary
 1793 rate $\|\delta_t\| = O(t^{-1/2} \log^c t)$ this is $O(T^{-1} \cdot T^{1/2-\nu_H} \log^c T) = o(T^{-1/2})$ a.s.

1794 *Term (C).* $r_t = (\hat{H}_t - H)\delta_{t-1} + O(\|\delta_{t-1}\|^2)$, where $\hat{H}_t$ is the block curvature.
 1795 The quadratic part contributes $O(T^{-1} \sum_t \log t/t) = O(\log^2 T/T) = o(T^{-1/2})$.
 1796 For the first part take $\zeta_t = \hat{H}_t - H$, which is stationary, centred and square
 1797 integrable under Assumption 3 (it is the block average of the terms appearing
 1798 in (A.9) evaluated at θ^* , plus a Lipschitz remainder in θ_{t-1} bounded exactly
 1799 as term (iv) there), and apply (A.3) with $\gamma_t \equiv 1$ and the $\mathcal{F}_{t-1}$ -measurable
 1800 weights acting by $A_t \zeta_t = P_{t-1} \zeta_t \delta_{t-1}$. By Lemma 2 and the rung-(iii) rate the
 1801 envelope is $c_t = O(t^{q_r} \cdot t^{-1/2} \log^c t)$ and the increments are $u_t = O(t^{-1} c_t)$, so

$$V_T = O\left(\sum_{t\leq T} t^{2q_r-1} \log^{2c} t\right) = O(T^{2q_r} \log^{2c} T), \quad U_T = O\left(\sum_t t^{q_r-3/2} \log^c t\right) = O(1) \quad (q_r < 1/2),$$

1802 and $\gamma_T c_T T^{1/2+\varpi} = O(T^{q_r+\varpi} \log^c T)$. Hence the sum is $O(T^{q_r+\varpi} \log^{c+1/2+\varpi} T)$
 1803 and (C) = $O(T^{q_r-1+\varpi} \log^{c'} T) = o(T^{-1/2})$ a.s. because $q_r < 1/2$. No inde-
 1804 pendence between $\hat{H}_t$ and δ_{t-1} is used or needed: adjacent blocks are only
 1805 $O(1/b)$ apart in the sense of Lemma 12, not $O(\rho^b)$, and an earlier version of
 1806 this proof wrongly appealed to the latter.

1807 Substituting into (A.11) gives $\|\delta_T\|^2 = O(\log N_T/N_T)$ a.s. (Theorem 5)
 1808 and $\sqrt{N_T} \delta_T = -N_T^{-1/2} H^{-1} \sum_{i\leq N_T} g_i + o_p(1)$. The central limit theorem for

strongly mixing stationary sequences with geometric coefficients and $2 + \nu$ moments (Ibragimov, 1962; Herrndorf, 1984) gives $N^{-1/2} \sum_{i \leq N} g_i \rightsquigarrow \mathcal{N}(0, S)$, and Slutsky completes Theorem 6. The averaged variant follows from the same decomposition with $\gamma_t = ct^{-a}$ and the standard Abel argument for weighted averages (Mokkadem and Pelletier, 2011; Pelletier, 2000). $\square$

Appendix A.7. The L^2 rate (Theorem 9)

Take expectations in (A.11). The leading term has $\mathbb{E} \|N^{-1} H^{-1} \sum_{i \leq N} g_i\|^2 = \text{tr}(H^{-1} S H^{-1})/N + o(1/N)$ by Lemma 12, uniformly in b because it does not involve b at all. For the remainders, replace each almost sure bound above by its L^2 counterpart: the martingale parts are bounded in L^2 by their predictable variations, which were already computed in expectation in the proof of Lemma 14, and the coboundary parts by $\mathbb{E} \|z_1\| \sum_t u_t$. On the descent side, iterating (A.6) in expectation with $\mathbb{E} \|\bar{g}_t(\theta_{t-1})\|^2 \leq C + C' \mathbb{E}[F(\theta_{t-1}) - F(\theta^*)]$ (Assumption 2) and $\gamma_t^2 \|P_{t-1}\|_{\text{op}}^2 = O(t^{2q-2})$ (Lemma 2) gives $\mathbb{E}[F(\theta_t) - F(\theta^*)] = O(\log t/t)$ by the standard Chung-type recursion, and local strong convexity converts this to $\mathbb{E} \|\delta_t\|^2 = O(\log t/t) = O(\log N_t/N_t) \cdot b$; the extra factor b is absorbed because $N_t = bt$ and the leading term is b -free. All constants depend on b only through $\mathbb{E} \|\bar{g}_t\|^2 \leq \mathbb{E} \|g_1\|^2$, which is monotone in b , so the bound is uniform over $b \leq C \log N$. $\square$

Appendix A.8. Coverage of dependence-blind intervals (Proposition 8)

By Theorem 6 and Theorem 3, $\sqrt{N}(\theta_{T,j} - \theta_j^*) \rightsquigarrow \mathcal{N}(0, V_{jj})$ with $V = H^{-1} S H^{-1}$, while $[\bar{H}^{-1} \hat{\Gamma}_0 \bar{H}^{-1}]_{jj} \rightarrow_p [H^{-1} \Gamma_0 H^{-1}]_{jj} = V_{jj}/\kappa_j$. Hence the studentised statistic converges to $\mathcal{N}(0, \kappa_j)$ and $\mathbb{P}(|\mathcal{N}(0, \kappa_j)| \leq z_{\alpha/2}) = 2\Phi(z_{\alpha/2}/\sqrt{\kappa_j}) - 1$. For the homogeneous AR design, $\Gamma_k = \sigma_u^2(\phi\psi)^{|k|}\Sigma_x$ gives $S = \sigma_u^2 \Sigma_x \sum_k (\phi\psi)^{|k|} = \sigma_u^2 \Sigma_x (1 + \phi\psi)/(1 - \phi\psi)$ and $H = \Gamma_0/\sigma_u^2 = \Sigma_x$, so $V = \kappa \sigma_u^2 \Sigma_x^{-1}$ with $\kappa = (1 + \phi\psi)/(1 - \phi\psi)$ and the inflation is the same in every coordinate. $\square$

Appendix A.9. Long-run covariance estimation (Theorem 10)

(i) *Bias of batch means.* For the idealised estimator built from $\varepsilon_t = \bar{g}_t(\theta^*)$, stationarity gives $\mathbb{E}[b \varepsilon_t \varepsilon_t^\top] = b\Lambda_0 = \sum_{|k| < b} (1 - |k|/b) \Gamma_k$. Splitting,

$$b\Lambda_0 = S - \sum_{|k| \geq b} \Gamma_k - \frac{1}{b} \left(\Delta - \sum_{|k| \geq b} |k| \Gamma_k \right), \quad \Delta = \sum_{k \geq 1} k(\Gamma_k + \Gamma_k^\top), \quad (\text{A.12})$$

and by Lemma 12 both remainder sums are $O(b\rho^b)$, giving $\mathbb{E}[\hat{S}^{\text{BM}}] = S - \Delta/b + O(\rho^b)$ up to the drift term treated below.

1841 (ii) *Bias of the two-scale estimator.* With $\Lambda_1 = \mathbb{E}[\varepsilon_0 \varepsilon_1^\top] = b^{-2} \sum_{i \in B_0, j \in B_1} \Gamma_{j-i}$
 1842 and counting multiplicities of each lag (k for $1 \leq k \leq b$, $2b-k$ for $b < k < 2b$),

$$b(\Lambda_1 + \Lambda_1^\top) = \sum_{1 \leq |k| \leq b} \frac{|k|}{b} \Gamma_k + \sum_{b < |k| < 2b} \frac{2b - |k|}{b} \Gamma_k. \quad (\text{A.13})$$

1843 Adding (A.12) in the form $b\Lambda_0 = \sum_{|k| < b} (1 - |k|/b) \Gamma_k$, the coefficient of Γ_k
 1844 becomes exactly 1 for $|k| \leq b$ and $(2b - |k|)/b \in (0, 1)$ for $b < |k| < 2b$, so

$$b\Lambda_0 + b(\Lambda_1 + \Lambda_1^\top) = S - \sum_{|k| \geq 2b} \Gamma_k - \sum_{b < |k| < 2b} \left(1 - \frac{2b - |k|}{b}\right) \Gamma_k = S + O(\rho^b), \quad (\text{A.14})$$

1845 using Lemma 12. This is the flat-top window of Politis and Romano (1995)
 1846 at block resolution; the point is that under geometric mixing its bias is
 1847 exponentially small in b , whereas the batch-means bias in (i) is only $O(1/b)$.
 1848 The algebraic identity $\hat{S}^{\text{FT}} = 2\hat{S}^{\text{BM}}(2b) - \hat{S}^{\text{BM}}(b)$ (using disjoint pairs) is
 1849 immediate from $\hat{S}^{\text{BM}}(2b) = \hat{S}^{\text{BM}}(b) + (b/T) \sum (\bar{g}_{2k-1} \bar{g}_{2k}^\top + \text{transpose})$.

1850 (iii) *Variance.* $\hat{S}_{jk}^{\text{BM}}$ is an average of $T' \asymp N/b$ terms $q_t = b \bar{g}_{tj} \bar{g}_{tk}$.
 1851 These are quadratic in the scores, so bounding their autocovariances by
 1852 Lemma 12 needs $2 + \delta$ moments of q_t , i.e. $4 + 2\delta$ moments of the scores;
 1853 this is exactly why Assumption 5 asks for $\mathbb{E}\|g_1\|^{4+\nu'} < \infty$ with a mar-
 1854 gin $\nu' > 0$ rather than plain fourth moments. Given that, (q_t) is a mea-
 1855 surable function of a geometrically mixing sequence at block resolution,
 1856 hence itself geometrically mixing with $\alpha_{\text{mix}}(k) \leq K\rho^{kb}$, and Davydov gives
 1857 $|\text{Cov}(q_t, q_{t+k})| \leq C\rho^{\delta kb/(2+\delta)}$. Summing the geometric series, $\text{Var}[\hat{S}_{jk}^{\text{BM}}] =$
 1858 $T'^{-2} \sum_{t,t'} \text{Cov}(q_t, q_{t'}) = O(1/T') = O(b/N)$, and the same argument applies
 1859 to the lag-one term $bT'^{-1} \sum \bar{g}_{2k-1,j} \bar{g}_{2k,k}$, hence to $\hat{S}^{\text{FT}}$.

1860 *Drift.* Write $\bar{g}_t(\theta_{t-1}) = \varepsilon_t + \Xi_t$ with $\Xi_t = \hat{H}_t \delta_{t-1} + O(\|\delta_{t-1}\|^2)$. Then
 1861 $\hat{S}^{\text{FT}}$ built from $\bar{g}_t(\theta_{t-1})$ differs from the idealised version by terms bounded
 1862 by $bT'^{-1} \sum_{t \in \mathcal{T}} \{\|\Xi_t\|^2 + 2\|\Xi_t\| \|\varepsilon_t\|\}$. Two facts are needed. First, on the
 1863 retained window $t \asymp T$, Theorem 5 gives $\mathbb{E}\|\delta_{t-1}\|^2 = O(\log t/(bt))$. Second,
 1864 $\hat{H}_t$ is a b -sample estimate of a $d \times d$ matrix, so for δ independent of block
 1865 t , $\mathbb{E}\|\hat{H}_t \delta\|_{\text{op}}^2 = \delta^\top \mathbb{E}[\hat{H}_t^2] \delta \leq (1 + Cd/b) \|H\delta\|_{\text{op}}^2$ under Assumption 3 (the d/b
 1866 term is the fluctuation $\mathbb{E}[(\hat{H}_t - H)M(\hat{H}_t - H)]$, whose trace contributes a
 1867 factor $\text{tr}(HM)/b$). Hence

$$bT'^{-1} \sum_{t \in \mathcal{T}} \mathbb{E}\|\Xi_t\|^2 = O\left((1+d/b) T'^{-1} \sum_t \log t/t\right) = O((1+d/b) b \log^2 N/N) = O((b+d) \log^2 N/N),$$

1868 and the cross term is of smaller order by Cauchy–Schwarz. Writing $\tau_{b,N} =$
 1869 $(b+d) \log^2 N/N$, this is dominated by the $\sqrt{b/N}$ standard deviation in (iii)

whenever $b + d = o(\sqrt{N})$ up to logarithms; in the reported experiments $b + d \leq 340$ and $N \geq 2.5 \times 10^4$. Note that $\tau_{b,N}$ grows as b falls below d , which is visible in Figure 2(b) and is why we do not recommend $b \ll d$.

(iv) *Rate.* The entrywise bounds give $\|\widehat{S}^{\text{FT}} - S\|_{\text{op}} \leq \|\widehat{S}^{\text{FT}} - S\|_F \leq d \max_{jk} |\widehat{S}_{jk}^{\text{FT}} - S_{jk}|$, hence $\|\widehat{S}^{\text{FT}} - S\|_{\text{op}} = O_p(d\sqrt{b/N} + \rho^b + \tau_{b,N})$; the factor d is the price of the conversion and we do not try to remove it (a matrix-Bernstein argument under an explicit sub-exponential condition would give $\sqrt{db/N}$). Choosing $b = \lceil c \log N \rceil$ with $c > 2/\log(1/\rho)$ makes $\rho^b = O(N^{-2})$ and gives $O_p(\sqrt{\log N/N})$. For $\widehat{S}^{\text{BM}}$ the bias term $1/b$ replaces ρ^b and the trade-off $\sqrt{b/N} \asymp 1/b$ gives $b \asymp N^{1/3}$ and $O_p(N^{-1/3})$. Finally, $\bar{H}^{-1} \rightarrow_p H^{-1}$ (Theorem 3) and $\widehat{S}^{\text{FT}} \rightarrow_p S$ give $u_j^\top \widehat{S}^{\text{FT}} u_j \rightarrow_p V_{jj}$, and Theorem 6 with Slutsky gives the studentised limit. $\square$

Appendix A.10. Curvature schedules (Proposition 11)

Let $\widehat{H}^{\text{gap}} = n^{-1} \sum_{j \leq n} Y_{jm}$. By stationarity,

$$n \text{Var}[\widehat{H}^{\text{gap}}] = \Gamma_0^H + 2 \sum_{j \geq 1} \left(1 - \frac{j}{n}\right) \Gamma_{jm}^H = \Gamma_0^H + 2 \sum_{j \geq 1} \Gamma_{jm}^H + o(1) = \Gamma_0^H + O(\rho^m),$$

using $\|\Gamma_k^H\|_{\text{op}} \leq C\rho^k$ (Lemma 12 applied to Y , which has $\eta' > 2$ moments by Assumption 3). For Bernoulli thinning, let $Z_i \stackrel{iid}{\sim} \text{Bernoulli}(p)$ be independent of the data and let

$$\widehat{H}^{\text{Bern}} = \left(\sum_{i \leq N} Z_i \right)^{-1} \sum_{i \leq N} Z_i Y_i$$

be the estimator with the *realised* normaliser — which is what the algorithm computes, since W counts the updates actually made. Using a fixed normaliser pN instead would leave a term $p(1-p) \text{vec}(H) \text{vec}(H)^\top$ coming from $\mathbb{E}[Y_i] = H \neq 0$, which is $\Theta(1)$ and would swamp the $O(1/m)$ effect; the ratio form removes it, because $\widehat{H}^{\text{Bern}} - H = (\sum_i Z_i)^{-1} \sum_i Z_i (Y_i - H)$ *exactly*. Writing $n_Z = \sum_i Z_i$ with $\mathbb{E} n_Z = pN = n$ and $n_Z/n \rightarrow 1$ a.s., the delta method for a ratio gives $n \text{Var}[\widehat{H}^{\text{Bern}}] = n^{-1} \text{Var} \left[\sum_i Z_i (Y_i - H) \right] + o(1)$, and

$$\text{Var} \left[\sum_i Z_i (Y_i - H) \right] = \sum_i p \Gamma_0^H + \sum_{i \neq j} p^2 \Gamma_{i-j}^H = Np \Gamma_0^H + 2p^2 \sum_{k \geq 1} (N-k) \Gamma_k^H.$$

Dividing by $n = pN$ and using $p = 1/m$, $n \text{Var}[\widehat{H}^{\text{Bern}}] = \Gamma_0^H + 2p \sum_{k \geq 1} \Gamma_k^H + o(1) = \Gamma_0^H + m^{-1}(S^H - \Gamma_0^H) + o(1)$. $\square$

Remark 7. The contrast is that random thinning leaves a fixed fraction p of *all* lag pairs in the sum, whereas a deterministic gap removes every lag below m . For Gaussian AR(1) covariates and $\alpha \equiv 1$, Γ_k^H decays like ϕ^{2k} and $\kappa_H = (1 + \phi^2)/(1 - \phi^2)$.

Appendix B. Experimental details

Appendix B.1. Data-generating processes and their population targets

Every simulated design pairs a dependent covariate process with a dependent error or latent process, for the reason given in Remark 1. In all four, $\theta^* = (-1, \dots, 1)$ evenly spaced in $\mathbb{R}^d$, $d = 20$, and the covariate covariance is the ill-conditioned $\Sigma = Q \text{diag}((i/d)^c) Q^\top$ of [Boyer and Godichon-Baggioni \(2023\)](#) with Q a random orthogonal matrix, $c = 2$ (condition number 400) except for the logistic design where $c = 1$ (condition number 20) to keep the linear predictor in a sensible range. The population (H, Γ_0, S) is fixed once per design; Monte Carlo replicates vary only the realised path, which is checked in `tests/test_oracles.py`.

AR-hom (closed form). $x_t = \phi x_{t-1} + \sqrt{1 - \phi^2} \Sigma^{1/2} \eta_t$, $u_t = \psi u_{t-1} + \sqrt{1 - \psi^2} \sigma_u \varepsilon_t$, $y_t = x_t^\top \theta^* + u_t$, with η, ε standard normal and independent, $\phi = \psi = 0.7$, $\sigma_u = 1$. Then $H = \Sigma$, $\Gamma_k = \sigma_u^2 (\phi\psi)^{|k|} \Sigma$ and $S = \sigma_u^2 \Sigma (1 + \phi\psi)/(1 - \phi\psi)$, so $\kappa_j \equiv 2.922$.

AR-het (closed form). $x_t = A_\phi x_{t-1} + \Sigma_\eta^{1/2} \eta_t$ with $A_\phi = \text{diag}(\phi_1, \dots, \phi_d)$, ϕ_j evenly spaced on $[0, 0.9]$. The stationary covariance is the Hadamard product $\Sigma_x = \Sigma_\eta \circ [1/(1 - \phi_i \phi_j)]$, which is positive definite because $[1/(1 - \phi_i \phi_j)] = \sum_{k \geq 0} (\phi \phi^\top)^{\circ k}$ is a sum of Hadamard powers of a rank-one positive semidefinite matrix whose first d terms are linearly independent when the ϕ_j are distinct. With AR(1) errors, $\mathbb{E}[x_t x_{t+k}^\top] = \Sigma_x A_\phi^k$, $\Gamma_k = \sigma_u^2 \psi^k \Sigma_x A_\phi^k$ and $S = \sigma_u^2 (\Sigma_x + \Sigma_x G_\psi + G_\psi \Sigma_x)$ with $G_\psi = \psi A_\phi (I - \psi A_\phi)^{-1}$. Because the ϕ_j differ, κ_j varies across coordinates, so a scalar correction cannot pass this design.

Markov (closed form). s_t is a $K = d + 1$ state Markov chain with $\mathbb{P}(s_t = s_{t-1}) = 0.97$, uniform off-diagonal transitions and uniform stationary law; $x_t = V_{\cdot, s_t} + \Sigma_\nu^{1/2} \nu_t$, so covariates are a Markov-modulated Gaussian mixture—non-Gaussian marginals, discrete latent state, regime (non-linear) dependence. The errors are AR(1) with $\psi = 0.8$. Then $\Sigma_x = V \text{diag}(\pi) V^\top + \Sigma_\nu$, $\mathbb{E}[x_t x_{t+k}^\top] = V \text{diag}(\pi) P^k V^\top$ and $S = \sigma_u^2 (\Sigma_x + R + R^\top)$ with $R = V \text{diag}(\pi) \{\psi P (I - \psi P)^{-1}\} V^\top$.

1931 The choice of centres matters and is worth recording, because the obvious
1932 construction produces a design that *looks* strongly dependent but has no
1933 inferential consequence. With randomly drawn centres and $K \ll d$, the regime
1934 component $V\text{diag}(\pi)V^\top$ has rank K and sits in high-curvature directions of
1935 Σ_x ; the sandwich $H^{-1}\Gamma_k H^{-1}$ is then negligible against $H^{-1}\Gamma_0 H^{-1}$, whose
1936 mass is in the *low*-curvature directions, and κ_j came out in $[1.00, 1.05]$ for
1937 every setting we tried. We therefore take $K = d + 1$ and the centres to
1938 be a regular simplex scaled by the target covariance: $V = \sqrt{K} \text{chol}(\Sigma)U$
1939 with $U \in \mathbb{R}^{d \times K}$ an orthonormal basis of $\mathbf{1}^\perp$ in $\mathbb{R}^K$, which gives $V\pi = 0$
1940 and $V\text{diag}(\pi)V^\top = \Sigma$ exactly. The regime component is then full rank and
1941 proportional to Σ , so $R_k = \lambda_2^k \Sigma$ with λ_2 the chain’s second eigenvalue and
1942 $\kappa_j \rightarrow (1 + \psi\lambda_2)/(1 - \psi\lambda_2)$ as $\Sigma_\nu \rightarrow 0$: at our settings, $\kappa_j = 7.55$ in every
1943 coordinate against the noise-free limit 7.88.

1944 This design is the one that forces both the longer warm-up and the
1945 stability shift, and the same feature causes both. With $\mathbb{P}(s_t = s_{t-1}) = 0.97$
1946 the chain dwells in a regime for about 33 observations and its integrated
1947 autocorrelation time is $(1 + 0.97)/(1 - 0.97) \approx 66$, so a block of $b = 20$
1948 consecutive observations almost always lies inside a single regime and the
1949 block Hessian $\hat{H}_t$ is nearly rank one. Two consequences, which we separated
1950 only after mistaking the second for the first. A warm-up of $50d$ *consecutive*
1951 observations spans only about thirty regime visits, so it leaves $\bar{H}$ rank-starved;
1952 $200d$ fixes that (Table 5(c')). But a well-estimated $\bar{H}$ is not enough, because
1953 the near-rank-one $\hat{H}_t$ makes $\|\bar{H}^{-1}\hat{H}_t\|_{\text{op}}$ large — median 33 here, against 20
1954 on AR-hom, with block maxima 89 and 55 — so the contraction condition
1955 (8) fails at $\gamma_1 = 1$ and the first few steps overshoot by orders of magnitude.
1956 At $c_w = 200$ and no shift, the relative error of the point estimate still exceeds
1957 one on 7 of 10 replicates. Both fixes are needed, and neither substitutes for
1958 the other (Table 5(i)).

1959 *Logistic (simulated oracle)*.. x_t is AR(1) with $\phi = 0.7$; the latent error is
1960 $e_t = \text{logit}(\Phi(\nu_t))$ with ν_t a unit-variance AR(1) with $\psi = 0.8$, so e_t has
1961 an exact standard logistic marginal and $y_t = \mathbf{1}\{x_t^\top \theta^* + e_t > 0\}$ satisfies
1962 $\mathbb{P}(y_t = 1 \mid x_t) = \sigma(x_t^\top \theta^*)$ *exactly*. The working logistic likelihood is therefore
1963 correctly specified marginally— θ^* is the true parameter and $\Gamma_0 = H$ by the
1964 information equality—while the score is serially correlated through ν . This
1965 is the marginal-model situation of [Liang and Zeger \(1986\)](#). H is computed
1966 to machine precision by Gauss–Hermite quadrature using the decomposition
1967 $x = (\Sigma_x \theta^* / \sigma_\ell^2) \ell + x_\perp$ with $\ell = x^\top \theta^*$ the linear predictor and $x_\perp \perp \ell$; S
1968 has no closed form and is obtained from eight independent streams of 10^6
1969 observations with the score evaluated at the known θ^* and a Bartlett long-run

covariance at lag 100 (the score autocovariance decays like $\psi^k = 0.8^k$, so the truncation error is below 10^{-9}). The resulting relative Monte Carlo error of the sandwich diagonal is 7.5×10^{-3} ; it is the only non-exact population quantity in the paper.

Appendix B.2. Real streams: features, segments and Protocol B’s nominal level

Responses. Both responses are log-transformed and then standardised to zero mean and unit variance, exactly as the covariates are. The standardisation is not cosmetic. Without it the mean of $\log(\text{traffic volume})$, about 7.9, sits entirely in the intercept, which is then roughly 400 standard errors from the $\theta_0 = 0$ that a streaming method starts from; a $1/t$ schedule needs of order that many steps merely to cover the distance, so on a 4000-hour segment (which affords only about 100 blocked steps at $b = 40$) the experiment would be measuring the distance to the initialisation rather than anything about inference — the infeasible oracle-variance interval covered 0.14 in that configuration, which is the diagnostic that the point estimate, not the covariance, was at fault. Standardising the response is an affine reparametrisation of a linear model and changes no coverage statement; it is also what any practitioner would do.

Features. Traffic: intercept, standardised temperature and cloud cover, two daily harmonics and one weekly harmonic, a holiday indicator propagated from the labelled midnight row to the whole calendar day, and a weekend indicator ($d = 11$). Rainfall is *excluded*: `rain_1h` is identically zero over stretches of several thousand consecutive hours, so $\log(1 + \text{rain})$ is a constant column on some contiguous segments, which makes the design there exactly rank deficient and the per-segment population targets undefined. Air quality: intercept, standardised temperature, pressure, dew point, $\log(1 + \text{rain})$ and wind speed, plus the same harmonics ($d = 12$). Feature construction was fixed before any response was looked at, and `exp6_real.py` asserts that every evaluation segment has full rank, reporting the worst per-segment condition number.

Segments and hold-out. The leading 15% of each stream is reserved for choosing the first-order baselines’ step-size constants and is excluded from every evaluation segment; the constants are chosen per stream, and the grid and whether the optimum is interior are recorded in `results/exp6_real.json`. Evaluation segments are the disjoint contiguous blocks that follow.

2006 *Protocol B’s nominal level.* Let R be the number of evaluation segments,
 2007 all of the same length, let $\hat{\theta}_k$ be the estimate from segment k and $\hat{\theta}_{\text{full}}$
 2008 the full-stream M -estimator. If (i) the segment estimates are approximately
 2009 uncorrelated, (ii) they have a common variance V , and (iii) $\hat{\theta}_{\text{full}} \approx R^{-1} \sum_k \hat{\theta}_k$,
 2010 then

$$\text{Var}(\hat{\theta}_k - \hat{\theta}_{\text{full}}) = \text{Var}(\hat{\theta}_k) - \text{Var}(\hat{\theta}_{\text{full}}) = V(1 - 1/R),$$

2011 so an interval built to cover θ^* with probability $1 - \alpha$ covers $\hat{\theta}_{\text{full}}$ with proba-
 2012 bility $2\Phi(z_{\alpha/2}/\sqrt{1 - 1/R}) - 1$, which is what the header of Table 4 reports.
 2013 Assumption (iii) requires the full-stream estimator to be (approximately) the
 2014 equally weighted average of the segment estimators, which holds when the
 2015 segments have equal length and similar $\hat{H}$; (i) is the same approximation that
 2016 makes segments usable as replicates at all. The correction is small but not
 2017 negligible—at the R of our traffic stream it raises the target from 0.950 to
 2018 0.976, and at the R of the air-quality stream to 0.994—and both the corrected
 2019 target and the raw coverage appear in Table 4, computed from R rather
 2020 than quoted from memory. The correction is also least trustworthy exactly
 2021 where it is largest: with only a handful of segments per stream, assumption
 2022 (iii) is a crude approximation and the corrected target is close to one, which
 2023 compresses the comparison. That is the main reason we treat Protocol A,
 2024 whose target is exact by construction, as the primary real-data evidence and
 2025 Protocol B as the check that the conclusion survives a real error process.

2026 *Appendix B.3. Hyperparameters*

2027 Unless a sweep says otherwise: block length $b = 20$ ($= d$), gap $m = b$
 2028 (one curvature slot per block), no Bernoulli thinning ($p = 1$), log-weights
 2029 off ($w' = 0$), ridge $c_i = 0.01$, $q_r = 0.4$, initial curvature scale $\lambda_0 = 10^{-6}$
 2030 (relative to the first block’s mean curvature magnitude, so it is dimensionless),
 2031 warm-up $c_w = 200$ observations per parameter capped at $\varpi_w = 10\%$ of the
 2032 stream, retained window dyadic (last half to last three quarters of blocks),
 2033 step-length safeguard constant $c_D = 1$ in (6), and the schedule shift of
 2034 [Appendix C.1](#) switched off (`t0_mult` = 0). Table 5 varies each of these. The
 2035 warm-up constant is 200 rather than the 50 that suffices on the autoregressive
 2036 designs because of the regime-switching design, for the reason given in
 2037 [Appendix B.1](#). The warm-up and the safeguard address different failures and
 2038 neither substitutes for the other: the warm-up makes $\bar{H}$ informative, and the
 2039 safeguard makes the first steps safe to take even when it is. With $c_w = 200$
 2040 and no safeguard the regime-switching design still diverges (Table 5(i)); with
 2041 the safeguard and $c_w = 50$ it no longer diverges but the curvature is still
 2042 rank-starved, and the relative variance is 2.19 against 1.41 at $c_w = 200$

(Table 5(c')). The two devices are complementary in exactly that way: the safeguard turns a divergence into an inefficiency, and the warm-up removes the inefficiency. The 10% cap on the warm-up never binds in the simulations, where $N = 200,000$ and $c_w d \leq 4000$; it binds on every real-data segment, where N is 8,000 or 10,000 (Table 5(j)).

The dependence-blind competitor “streaming Newton, i.i.d. theory” uses the base method’s schedule exactly: every observation is a candidate curvature slot ($m = 1$) with Bernoulli retention $p = 1/d$, and the interval uses the i.i.d. plug-in $\hat{\Gamma}_0 = N^{-1} \sum_i s_i(\theta_{t-1}) s_i(\theta_{t-1})^\top$, accumulated over exactly the observations the blocked estimators use—that is, excluding the warm-up, so that the comparison is not contaminated by scores evaluated at θ_0 .

First-order baselines use step $\gamma_t = c t^{-3/4}$ (averaged SGD) or a constant c with AdaGrad scaling, with c chosen from a 19-point logarithmic grid on $[10^{-3}, 10^{1.5}]$ minimising squared error on eight *pilot* streams whose seeds are disjoint from the evaluation streams; the selected constants are recorded in `results/exp1_main.json`. They are additionally handed the *true* Hessian for their sandwich. Random scaling uses the two-sided 95% quantile 6.811 of the pivot $|W(1)|/(\int_0^1 (W(r) - rW(1))^2 dr)^{1/2}$, taken from Kiefer and Vogelsang (2002) as tabulated by Lee et al. (2022, Table 2). Our own simulation of that quantile, using the exact Karhunen–Loève representation $\int_0^1 B^2 = \sum_{k \geq 1} Z_k^2 / (k\pi)^2$ with Z_k i.i.d. standard normal independent of $W(1)$, gives 6.746 from 4×10^6 draws (the sampler is validated by reproducing $\mathbb{E} \int_0^1 B^2 = 1/6$ and $\text{Var} \int_0^1 B^2 = 1/45$); we use the larger published value, which is generous to that baseline. We also flag a normalisation trap: several papers quote 5.32 as “the 95% critical value”, which is the 95% quantile of the *signed* statistic, i.e. the two-sided 90% value, and using it would make the random-scaling intervals substantially too narrow.

The iterate-path comparator. Supplementary Fig. S1(a) compares against overlapping batch means computed from the iterate path, in the matrix form of the estimator used in the dependent-data literature: with T recorded iterates, window length b_n , window averages $\bar{\theta}_{b_n,t} = b_n^{-1} \sum_{k=t}^{t+b_n-1} \theta_k$ and overall average $\bar{\theta}$,

$$\hat{\Sigma}^{\text{it}} = \frac{N}{T} \cdot \frac{b_n}{T - b_n + 1} \sum_t (\bar{\theta}_{b_n,t} - \bar{\theta})(\bar{\theta}_{b_n,t} - \bar{\theta})^\top, \quad b_n = \lceil T^{3/4} \rceil,$$

which estimates the covariance of $\sqrt{N}(\bar{\theta} - \theta^*)$, i.e. the same Σ our sandwich estimates. It uses no gradients and no curvature estimate. We compute it from our own recorded trajectory, with cumulative sums, so it costs $O(Td)$

time and $O(Td)$ memory; it is a reference, not a streaming method, and we make no claim about the constants of any particular published recursion.

Appendix B.4. Reproduction and compute

Everything in this paper is public: the estimator, the experiment scripts, the saved result files from which every number, figure and table is generated, and the tests that validate every closed-form population quantity against long simulation. The exact release used to produce the results reported here is archived with a persistent identifier ([Kendiukhov, 2026](https://zenodo.org/record/21891313)),

[10.5281/zenodo.21891313](https://zenodo.org/record/21891313) (development version: <https://github.com/Kendiukhov/streaming-quasi-newton-dependent-data>),

and the concept identifier [10.5281/zenodo.21891312](https://zenodo.org/record/21891312) always resolves to the most recent release. The two real data sets are third-party and are not redistributed; `fetch_data.py` downloads them from the UCI Machine Learning Repository ([Hogue, 2019](https://archive.ics.uci.edu); [Chen, 2017](https://archive.ics.uci.edu); [Kelly et al.](https://archive.ics.uci.edu)). `make all` runs everything. All BLAS libraries are pinned to a single thread and parallelism is over Monte Carlo replicates (`experiments/_env.py` explains why: on the machine used, an unpinned tall-skinny $(d, N) \times (N, d)$ product with $N \approx 4 \times 10^4$, $d \approx 12$ took 18s against 5ms pinned). Total compute for every number in the paper is a few CPU-hours. The machine was shared and its load varied by more than an order of magnitude during the runs, so Table 6 reports (i) exact machine-independent operation counts computed from the algorithm itself and (ii) wall-clock times measured with the repetitions *interleaved* across methods and summarised by the minimum over repetitions—round-robin ordering exposes every method to the same load pattern, and external load can only make a measurement slower, so the minimum is the robust summary. The operation counts are the authoritative statement; the times corroborate them.

Appendix B.5. Additional results

Appendix C. Safeguarding the blocked step

Proof of Proposition 1. On $\mathcal{A}$ there are finitely many t at which Π_t is active, so $T_0 = \max\{t : \Pi_t \text{ active}\}$ is finite and for $t > T_0$ the recursion is $\theta_t = \theta_{t-1} - \gamma_t \bar{H}_{t-1}^{-1} \bar{g}_t(\theta_{t-1})$ verbatim. Two sequences that agree from a finite index on have the same limit points, the same almost sure rate and the same limit law, since all three are tail properties; the initial condition at T_0 is a finite random vector, and Theorems 4 to 6 hold for the recursion started from any initial condition (Assumption 4 is global). For the last claim, if $\theta_t \rightarrow \theta^*$

2114 and $\|\gamma_t \bar{H}_{t-1}^{-1} \bar{g}_t(\theta_{t-1})\| \rightarrow 0$ almost surely then for almost every ω there is
 2115 $T_1(\omega)$ beyond which $\|\gamma_t \bar{H}_{t-1}^{-1} \bar{g}_t\| < c_D(1 + \|\theta^*\|)/2 \leq c_D(1 + \|\theta_{t-1}\|)$, so Π_t
 2116 is inactive for $t > T_1$ and $\omega \in \mathcal{A}$. $\square$

2117 We repeat here what Section 2 says, because it is the kind of gap a reader
 2118 is entitled to be told about twice. Proposition 1 is conditional on $\mathcal{A}$ and
 2119 does not establish $\mathbb{P}(\mathcal{A}) = 1$ unconditionally: the implication runs from
 2120 convergence to eventual inactivity, not the other way. The reason we do
 2121 not close the loop is that the safeguard factor $\min\{1, c_D(1 + \|\theta_{t-1}\|)/\|v\|\}$
 2122 depends on $\bar{g}_t$, hence is $\mathcal{F}_t$ - and not $\mathcal{F}_{t-1}$ -measurable, and multiplying the
 2123 noise term $\gamma_t \bar{H}^{-1} \varepsilon_t$ by an $\mathcal{F}_t$ -measurable factor destroys the martingale-
 2124 difference property that Theorems 4 and 6 rest on. Two routes would close it
 2125 and both are outside our scope: the expanding-truncation construction of
 2126 Kushner and Yin (2003) and Andrieu et al. (2005), which replaces the cap by
 2127 a projection onto a ball of radius increasing with the number of activations
 2128 and proves the activation count finite; or an $\mathcal{F}_{t-1}$ -measurable variant that
 2129 caps using the *previous* block’s amplification, $\gamma_t \wedge \tau / \|\bar{H}_{t-2}^{-1} \hat{H}_{t-1}\|_{\text{op}}$, which is
 2130 predictable and so leaves the martingale structure intact, at the cost of an
 2131 extra $O(d^2)$ per block. What we do instead is report the activation counts,
 2132 which are between 2 and 7 per run and do not grow with N .

2133 *Appendix C.1. The step-size shift: the alternative we implemented and re-* 2134 *jected*

2135 The variant replaces the step sequence $\gamma_t = c t^{-a}$ of (6) by $\gamma_t = c(t+t_0)^{-a}$
 2136 with

$$t_0 = \max\left\{0, \frac{c}{2} \max_{t \leq T_w} \|\bar{H}_{T_w}^{-1} \hat{H}_t\|_{\text{op}} - 1\right\}, \quad (\text{C.1})$$

2137 $T_w = \lceil c_w d/b \rceil$ being the last warm-up block, and switches the step-length
 2138 safeguard Π_t off. By construction (8) then holds at $t = 1$ for every block whose
 2139 amplification does not exceed the warm-up maximum. The spectral norms
 2140 are obtained by four power iterations per warm-up block with $\hat{H}_t$ applied in
 2141 factored form, so no $d \times d$ product is formed and the whole computation costs
 2142 $O(c_w(d^2 + db)d/b)$ flops once, $O(1)$ in N ; with $b = m = d$ that is $O(c_w d^2)$,
 2143 some 4–5% of the total at the settings of Table 6. Setting `t0_mult` = 0.5 and
 2144 `step_cap` = 0 in the code selects this variant.

2145 *Correctness of the rejected shift variant.* Write T_w for the last warm-up block.
 2146 For every $t \leq T_w$, $\|\bar{H}_{T_w}^{-1} \hat{H}_t\|_{\text{op}} \leq \|\bar{H}_{T_w}^{-1}\|_{\text{op}} \|\hat{H}_t\|_{\text{op}}$. The first factor is finite
 2147 almost surely by Lemma 2, which bounds it by $c_1^{-1} n_{T_w}^{q_r}$ using only the ridge in
 2148 (4) and no property of the iterates. The second is finite almost surely because

2149 $\|\widehat{H}_t\|_{\text{op}} \leq b^{-1} \sum_{i \in B_t} \alpha_i \|\Psi_i\|^2$ and Assumption 3 gives $\mathbb{E}\|\alpha\Psi\Psi^\top\|_{\text{op}}^{\eta'} \leq C_{\eta'}$
 2150 with $\eta' > 2$, so each summand is finite a.s. and the maximum is over the
 2151 finitely many blocks $t \leq T_w$. Hence $t_0 < \infty$ a.s. Measurability with respect
 2152 to $\mathcal{F}_{T_w} = \sigma(\xi_1, \dots, \xi_{bT_w})$ is immediate: (C.1) is a function of $\bar{H}_{T_w}$ and of
 2153 $\{\widehat{H}_t\}_{t \leq T_w}$, all of which are built from the warm-up window and from θ_0 , and
 2154 the warm-up takes no parameter step.

2155 For the second claim, fix ω in the almost sure event $\{t_0 < \infty\}$ and let
 2156 $K = \lceil t_0 \rceil$. Then $c(t+t_0)^{-a} \geq c(t+K)^{-a}$, and $\sum_{t \geq 1} (t+K)^{-a} = \sum_{t > K} t^{-a} =$
 2157 ∞ for $a \leq 1$, so $\sum_t \gamma_t = \infty$. Likewise $c(t+t_0)^{-a} \leq ct^{-a}$, so $\sum_t \gamma_t^{1+\delta} \leq$
 2158 $c^{1+\delta} \sum_t t^{-a(1+\delta)} < \infty$ whenever $a(1+\delta) > 1$, which is the same condition
 2159 as for the unshifted sequence. Finally $(t+t_0)^{-a}/t^{-a} = (1+t_0/t)^{-a} \rightarrow 1$.
 2160 Conditioning on $\mathcal{F}_{T_w}$ is harmless because $\mathcal{F}_{T_w}$ is generated by finitely many
 2161 observations, so it changes no tail statement and, by Assumption 5, leaves
 2162 the stream after the warm-up geometrically mixing with the same constants
 2163 up to a factor.

2164 Theorem 4 and the preliminary rate use the step sequence only through
 2165 $\sum_t \gamma_t = \infty$ and $\sum_t \gamma_t^2 \|P_{t-1}\|_{\text{op}}^2 < \infty$, both just verified, so they transfer
 2166 immediately. Theorems 5 and 6 and Proposition 7 deserve one extra line,
 2167 because their proof uses the identity $\gamma_t = 1/t$ explicitly in the Abel summation
 2168 (A.10), and we would rather write the shifted identity out than assert that
 2169 nothing changes. With $\gamma_t = (t+t_0)^{-1}$, multiplying $\delta_t = \delta_{t-1} - \gamma_t P_{t-1} \{H\delta_{t-1} +$
 2170 $\varepsilon_t + r_t\}$ by $t+t_0$ gives

$$(t+t_0)\delta_t = (t-1+t_0)\delta_{t-1} + (I - P_{t-1}H)\delta_{t-1} - P_{t-1}\varepsilon_t - P_{t-1}r_t, \quad (\text{C.2})$$

2171 which is (A.10) with t replaced by $t+t_0$ throughout. Summing from $t=1$ to
 2172 T now telescopes to $(T+t_0)\delta_T = t_0\delta_0 + \sum_{t \leq T} \{\cdot\}$ instead of $T\delta_T = \sum_{t \leq T} \{\cdot\}$,
 2173 so dividing by $T+t_0$,

$$\delta_T = \frac{t_0 \delta_0}{T+t_0} - \frac{1}{T+t_0} H^{-1} \sum_{t \leq T} \varepsilon_t + \frac{T}{T+t_0} \{(\text{A}) + (\text{B}) - (\text{C})\},$$

2174 with (A), (B), (C) exactly as in the unshifted proof. Two differences from
 2175 (A.10), and both are negligible. The new first term is $O(t_0/T) = O(1/T) =$
 2176 $o(T^{-1/2})$ because t_0 is a finite constant, so it does not enter the limit. And
 2177 $T/(T+t_0) = 1 + O(1/T)$ and $(T+t_0)^{-1} = T^{-1}(1 + O(1/T))$, so every term
 2178 of the unshifted decomposition is multiplied by $1 + O(1/T)$; since each of
 2179 them is $O_p(T^{-1/2})$ or smaller, the perturbation is $O_p(T^{-3/2})$. The remaining
 2180 estimates — the bounds on (A), (B) and (C), the martingale central limit
 2181 theorem applied to $T^{-1/2} \sum_{t \leq T} \varepsilon_t$, and the Gordin decomposition — do
 2182 not involve γ_t at all. Hence the rate, the block-freeness statement and the

2183 limit law are unchanged. So the shift variant has the cleaner theory of the
 2184 two; it is the empirical comparison, not the theory, that made us prefer the
 2185 safeguard. $\square$

2186 Two remarks on what this variant does and does not do. First, (C.1) does not
 2187 make (8) hold for *every* t : the maximum in (C.1) is over warm-up blocks, so a
 2188 later block with an unusually badly conditioned $\hat{H}_t$ can still violate (8) at that
 2189 single t . What the shift buys is that the violation happens with $\gamma_t \leq (1+t_0)^{-1}$
 2190 rather than $\gamma_1 = 1$, so a single bad block moves θ by a bounded multiple of its
 2191 distance to θ^* instead of by a large one, and the recursion recovers. Second,
 2192 the proposition is about asymptotics; the finite-sample effect of the shift
 2193 is a trade, because a larger t_0 also slows the decay of the initial condition,
 2194 which decays like $t_0/(T+t_0)$ rather than $1/T$. This is why we set t_0 from
 2195 (8) — the smallest value that restores contraction — rather than by taking a
 2196 comfortable margin, and why Table 5 reports the sensitivity.